



CPTS 223 Advanced Data Structure C/C++

Midterm Exam Review


Tree: Red-Black Trees

Overview

- Tree data structure
- Binary search trees
 - Support $O(\lg(n))$ operations
 - **Balanced trees**
- STL **set and map** classes
- B-trees for accessing secondary storage
- Applications of Tree

This time: RB trees (practically used self-balanced BSTs)

Motivation



OS + compiler

OS + compiler + Make

OS + compiler + Make + Cmake

→ **executable** program

→ **automation** of executable program

→ **Cross-platform** automation of
executable program

C++ templated class

```
template <class T>
class MyPair {
    T values [2];
public:
    mypair (T first, T second)
    {
        values[0]=first;
        values[1]=second;
    }
};
```

- Templated class
- Type of 'T' defined at instantiation time
- Can have multiple templated types
- Works well if 'T' follows good OO design

```
MyPair<int> pair1 = new MyPair<int>();
```

Lvalues and Rvalues

- Lvalues
 - Permanent variables or objects
 - Persist beyond immediate use
 - Passed to a function with &
 - `string &rstr = str;`
- Rvalues
 - Temporary values
 - Could be as simple as a number
 - `myfunc(2)`
 - `'2'` is a Rvalue

How about `x+y`? A lvalue or rvalue?

The Big-Five

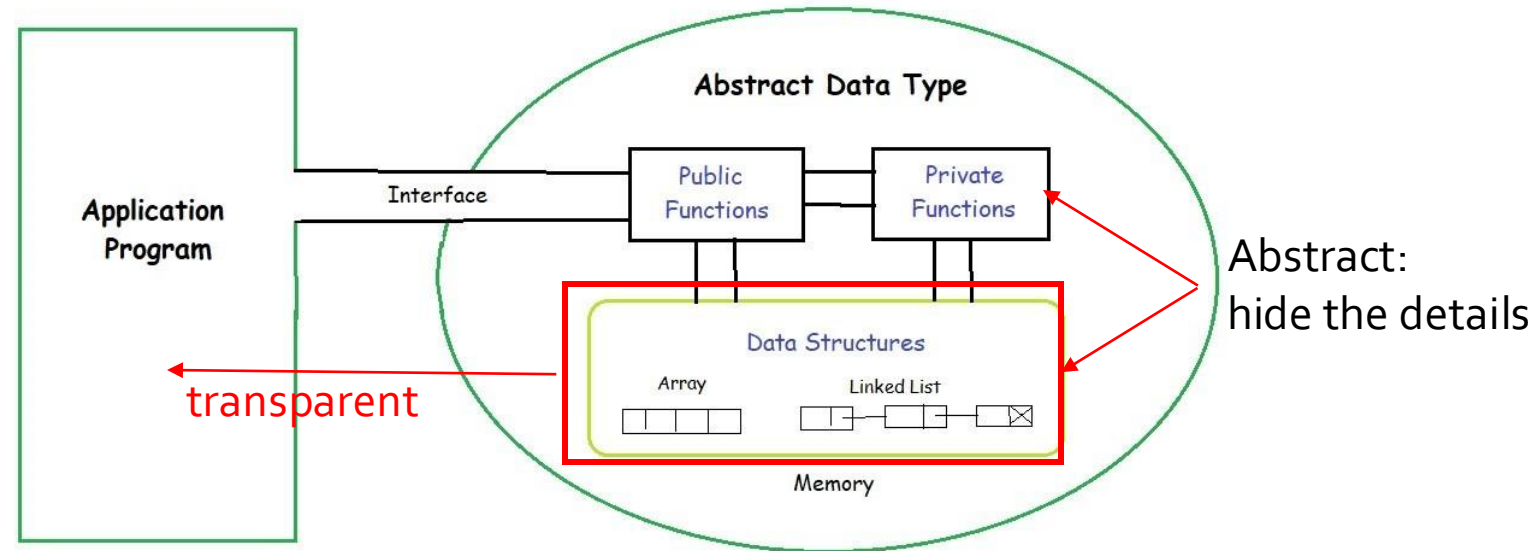
- Big Five: why and when do you need to do them? Resource management
 - Copy constructor
 - Move constructor
 - Copy Assignment operator=
 - Move Assignment operator=
 - Destructor

Interface of the Big-Five

```
~IntCell( );                // Destructor
IntCell( const IntCell & rhs ); // Copy constructor
IntCell( IntCell && rhs );    // Move constructor
IntCell & operator= ( const IntCell & rhs ); // Copy assignment operator
IntCell & operator= ( IntCell && rhs );    // Move assignment operator
```

Abstract Data Type

ADTs



Examples: running time?

// Assume A is an integer array of size n

Algorithm1 (A, n)

max = infinity;

for ($i=1$ to n) {

if ($A[i] > \text{max}$)

max = $A[i]$;

}

Output max;

$T(n) = n + \text{const}$

Algorithm3 ($A, 1, n$)

if ($n < 2$) return;

$x = \text{floor}(n/2)$;

$T(1) = 1$;

Algorithm3 ($A, 1, x$)

Algorithm3 ($A, x+1, n$)

$T(n) = n(1 + \text{const})$

Definition: (1) Let $T(n)$ denote the time take by an algorithm on an input of size n . (2) $T(1) = 1$

Algorithm2 ($A, 1, n$)

if ($n < 2$) return;

mid = floor($n/2$);

if (condition#1)

Algorithm2 ($A, 1,$

mid);

else

Algorithm2 ($A,$

mid+1, n);

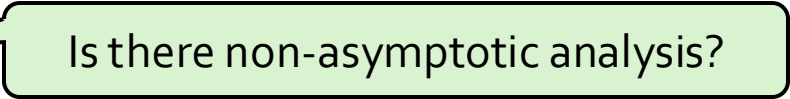
$T(n) = 1 + \text{const} * \log_2 n$

Is running time comparable?

Benchmarks can be sensitive

- Given the list {3, 9, 1, 2, 3, 5} as input
- Find(X): return the index of X
- We expect Find(3) to execute faster than Find(5)
 - Which is more representative of this input list, Find(3) or Find(5)?
 - How is our benchmark affected if we decide to use different lists?
 - e.g. {1, 2, 3, 3, 3, 4, 4, 4, 5, 9}
- Benchmarks are great for telling us how an algorithm will execute under a given set of circumstances
- but we ideally want something more **generalizable**

What is algorithm analysis?

- A mathematical technique for estimating the **rate** at which execution time **grows** relative to the **size of its input** parameters
- A formal name: **asymptotic analysis**  Is there non-asymptotic analysis?
- Asymptotic analysis is a method of estimation that groups algorithms based on their **growth rate**
- Asymptotic analysis is **unable** to tell us for sure how one algorithm will perform exactly in absolute timed execution relative to another
- but it does give us some good **hints**

Algorithm complexity

- $T(n)$ is time to run given an input size of n elements
- $T(n) = O(f(n))$: exist $[c, n_o]$ such that $T(n) \leq cf(n)$ when $n \geq n_o$
 - e.g., $T(n) \leq 2.45 n^2$, where $f(n)=n^2$
- $T(n) = \Omega(g(n))$ when $+[c, n_o]$ such that $T(n) \geq cg(n)$ when $n \geq n_o$
 - e.g., $T(n) \geq 1.03 n^2$, where $g(n)=n^2$
- $T(n) = \Theta(h(n))$ if and only if $T(n) = O(h(n))$ and $T(n) = \Omega(h(n))$
 - e.g., $1.03 n^2 \leq T(n) \leq 2.45 n^2$, where $h(n)=n^2$

Bounds

- $O(f(n))$ is an **UPPER bound** of $T(n)$ -- "Worst case can be no more than"
 $\leftarrow T(n) \leq cf(n)$
- $\Omega(g(n))$ is a **LOWER bound** of $T(n)$ -- "Best case can be no faster than"
 $\leftarrow T(n) \geq cg(n)$
- Only the **order** of the algorithm
- **No details, e.g., constants** (asymptotic analysis)
- $T(n) = \Theta(g(n))$ is when $O(g(n)) = \Omega(g(n))$ -- "It must be exactly"
- You will find **Theta (Θ)** also used as an **average case**

Nested loops

- ```
for(int i = 0; i < num_items; i++) {
 for(int j = 0; j < num_items; j++) {
 swap(items[i], items[j])
 }
}
```
- In this case, we multiply the effect that num\_items has on the growth rate, yielding  $O(n^2)$

# Unrelated loops

---

```
for(int i = 0; i < num_items; i++) {
 cout << "hello";
}
for(int j = 0; j < num_items; j++) {
 cout << "goodbye";
}
```

- $O(n + n)$  or  $O(2n) \rightarrow$  simplify to  $O(n)$

# Reduction of non-constant time

---

- In Big-O analysis, we always drop coefficients:
  - $O(2n) \rightarrow O(n)$
  - $O(40n) \rightarrow O(n)$
  - $O(1000000n) \rightarrow O(n)$
- This is because Big-O cares about placing algorithms into performance groups, not absolute  $T(n)$  calculations

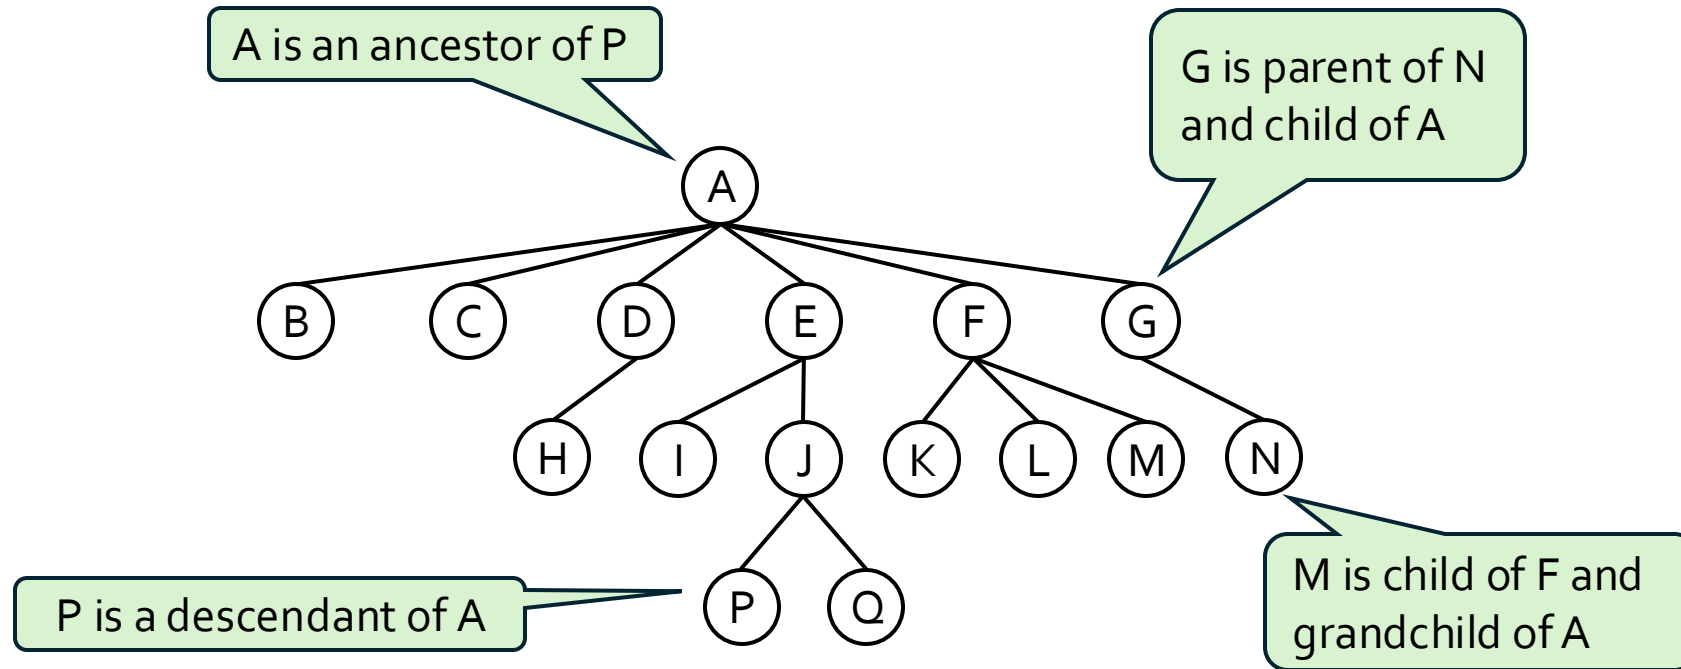


# How to compare algorithms?

| Function    | Name                               |
|-------------|------------------------------------|
| $c$         | Constant                           |
| $\log(n)$   | Logarithmic                        |
| $\log^2(n)$ | Log-squared                        |
| $n$         | Linear                             |
| $n \log(n)$ | (Will see this in sorting *a lot*) |
| $n^2$       | Quadratic                          |
| $n^3$       | Cubic                              |
| $2^n$       | Exponential                        |

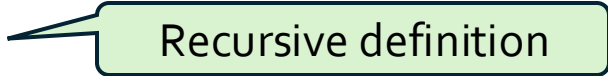
Tree: Binary Search Tree

# Trees



# Basic definitions

---

- A tree  $T$  is a set of nodes that form a **directed acyclic graph (DAG)** such that:
  - Each non-empty tree has a root node and zero or more sub-trees  $T_1, \dots, T_k$
  - Each sub-tree is a tree  Recursive definition
  - An internal node is connected to its children by a directed edge
- Each node in a tree has only one parent
  - Except the root, which has no parent

# Basic definitions

- **Internal** node: nodes with at least one child
- **Leaf** node: nodes with no children
- **Siblings**: nodes with the same parent
- A **path** from node  **$n_1$**  to  **$n_k$**  is a sequence of nodes  $n_1, n_2, \dots, n_k$  such that  $n_i$  is the parent of  $n_{i+1}$  for  $1 \leq i < k$ 
  - The **length** of a path is the number of edges on the path (i.e.,  **$k-1$** )
  - Each node has a path of length 0 to itself
  - There is exactly one path from the root to each node in a tree
  - Nodes  $n_i, \dots, n_k$  are **descendants** of  $n_i$  and **ancestors** of  $n_k$
  - Nodes  $n_{i+1}, \dots, n_k$  are **proper descendants** of  $n_i$
  - Nodes  $n_i, \dots, n_{k-1}$  are **proper ancestors** of  $n_k$

Nodes →  
either a leaf or an  
internal node


# Basic definitions

---

- The **depth** of a node  $n_i$  is the length of the path from the root to  $n_i$ 
  - The **root node has a depth of 0**
  - The **depth of a tree == the depth of its deepest leaf**
- The **height** of a node  $n_i$  is the **length of the longest path** under  $n_i$ 's subtree
  - **All leaves have a height of 0**
- **height of tree = height of root = depth of tree**

Tree: Binary Search Tree

# Height and depth




1. Height of each node?
2. Height of the tree?
3. Depth of each node?
4. Depth of the tree?

e.g.,  $\text{height}(E)=2$ ,  $\text{height}(L)=0$

Tree: Binary Search Tree

# Height and depth




1. Height of each node?
2. Height of the tree?
3. Depth of each node?
4. Depth of the tree?

= 3 (height of longest path from root)

Tree: Binary Search Tree

# Height and depth




1. Height of each node?
2. Height of the tree?
3. Depth of each node?
4. Depth of the tree?

e.g.,  $\text{depth}(E)=1$ ,  $\text{depth}(L)=2$



Tree: Binary Search Tree

# Height and depth




1. Height of each node?
2. Height of the tree?
3. Depth of each node?
4. Depth of the tree?

= 3 (length of the path to the deepest node)

Tree: Binary Search Tree

# Traversals


---



- Pre-order: **root** - left - right
- Post-order: left - right - root
- In-order: left - **root** - right

Tree: Binary Search Tree

# Traversals




- Pre-order: + + a \* b c \* + \* d e f g
- Post-order: a b c \* + d e \* f + g \* +
- In-order: a + b \* c + d \* e + f \* g

Tree: Binary Search Tree



# Inserting into BSTs

- e.g., insert 5



# Removing from BSTs


- There are two cases for removal
- **Case 1:** Node to remove has 0 or 1 child
  - Action: remove it and make appropriate adjustments to retain BST structure
  - e.g., `remove(4)`



Tree: Binary Search Tree

# Removing from BSTs

- **Case 2:** Node to remove has 2 children
  - Action:
    - Replace node element with **successor**
    - Remove the **successor** (case 1)
  - e.g., remove(2)



Tree: AVL Tree

# Balanced BSTs


---

- AVL (Adelson-Velskii and Landis) trees
  - Height of left and right subtrees at every node in BST differ by at most 1
  - Balance forcefully maintained for every update (via rotations)
  - BST depth always  $O(\lg(n))$

Tree: AVL Tree

# AVL trees

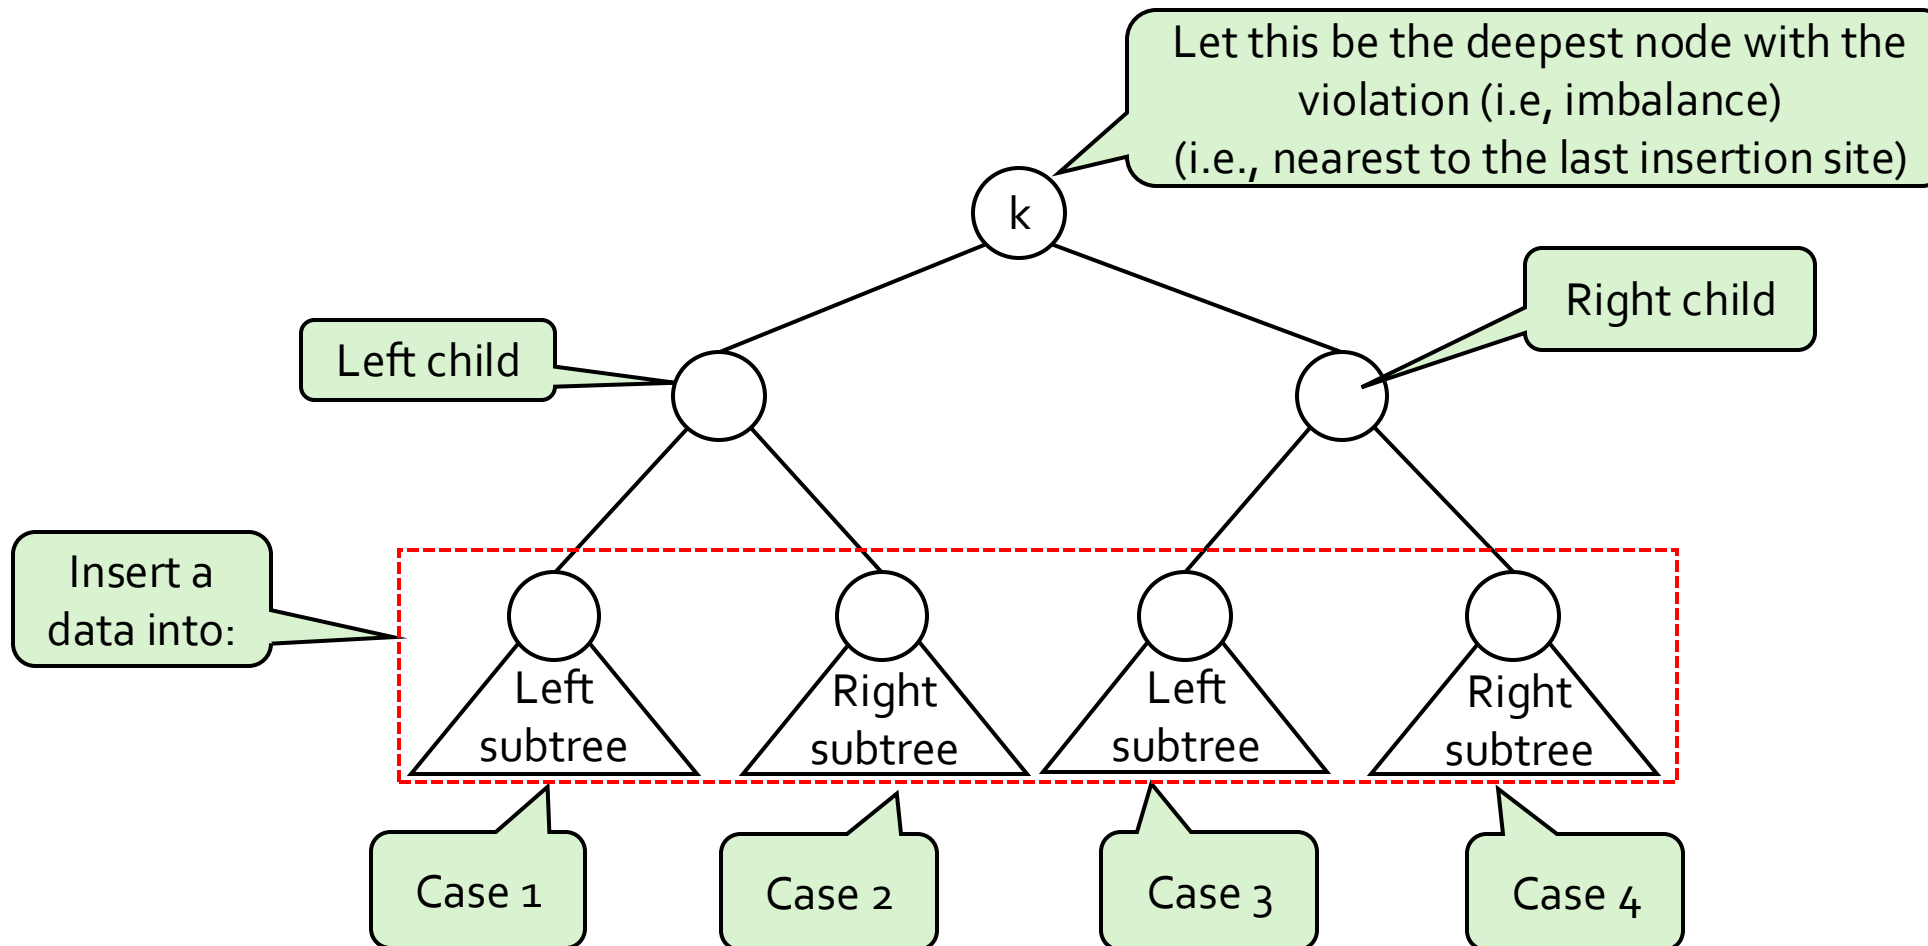
- Which of these is a valid AVL tree?





Tree: AVL Tree


# Identify cases for AVL insert



Tree: AVL Tree

# AVL insert (single rotation)

- Case 1: single rotation right



**Figure 4.34** Single rotation to fix case 1

Tree: AVL Tree


# AVL insert (double rotation)

- Case 2: left-right double rotation

Can be implemented as two successive rotations


height diff=2

Insert a data into:



**Figure 4.38** Left-right double rotation to fix height diff=2

height diff=1




- BST order?
- AVL balance condition?

$A < k_1 < B < k_2$   
 $< C < k_3 < D$

Tree: Splay Tree

# Splaying: zig-zag

- Node X is right-child of parent, which is left-child of grandparent (or vice-versa)
- Perform **double** rotation (**left, right**)




**Figure 4.48** Zig-zag

Tree: Splay Tree

# Splaying: zig-zig

- Node X is left-child of parent, which is left-child of grandparent (or right-right)
- Perform **double** rotation (**right-right**)

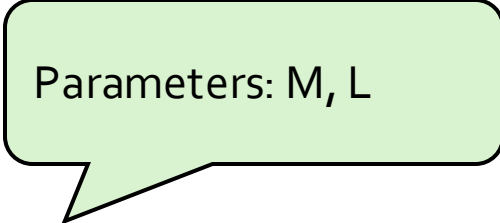


**Figure 4.49** Zig-zig

# B-tree (B+ tree) definition

---

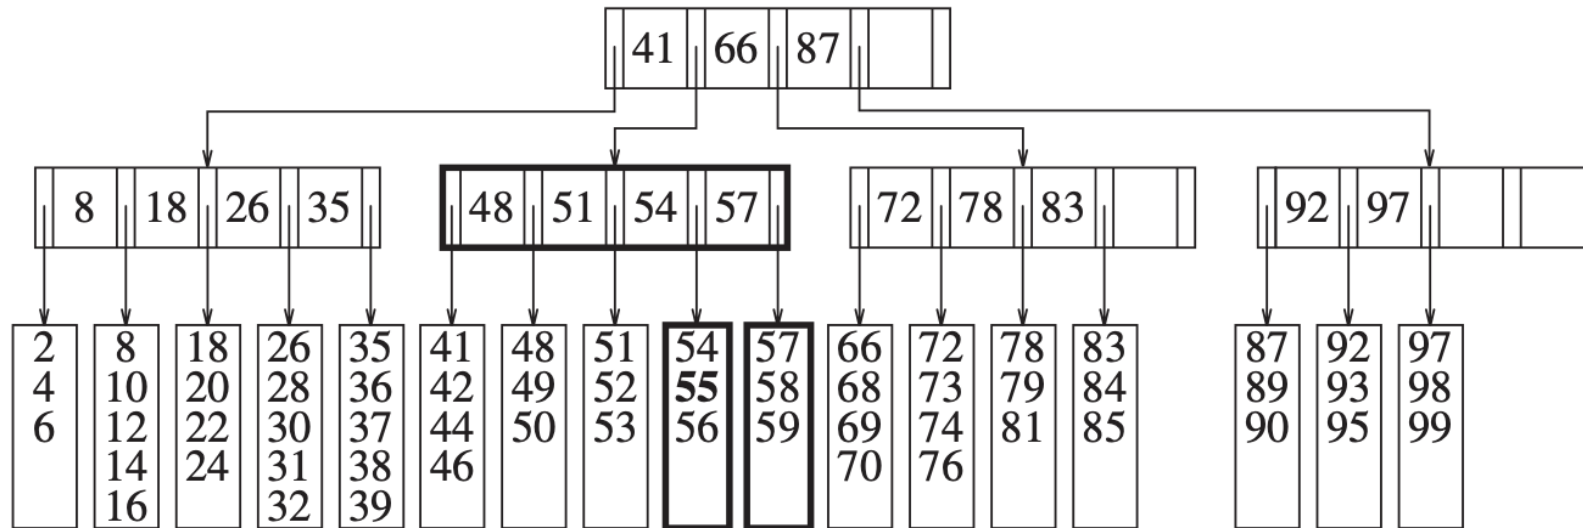
- Leaves store the real data items
- Internal nodes store up to **M-1 keys**
  - s.t., **key i** is the **smallest key** in subtree **i+1**
- Root can have between 2 to **M children**
- Each internal node (except root) has between  $\text{ceil}(M/2)$  to **M children**
- All leaves are at the same depth
- Each **leaf** has between  $\text{ceil}(L/2)$  and **L data items**, for some L



Parameters: M, L

Tree: B-Tree and B+ Tree

# B-tree of order 5

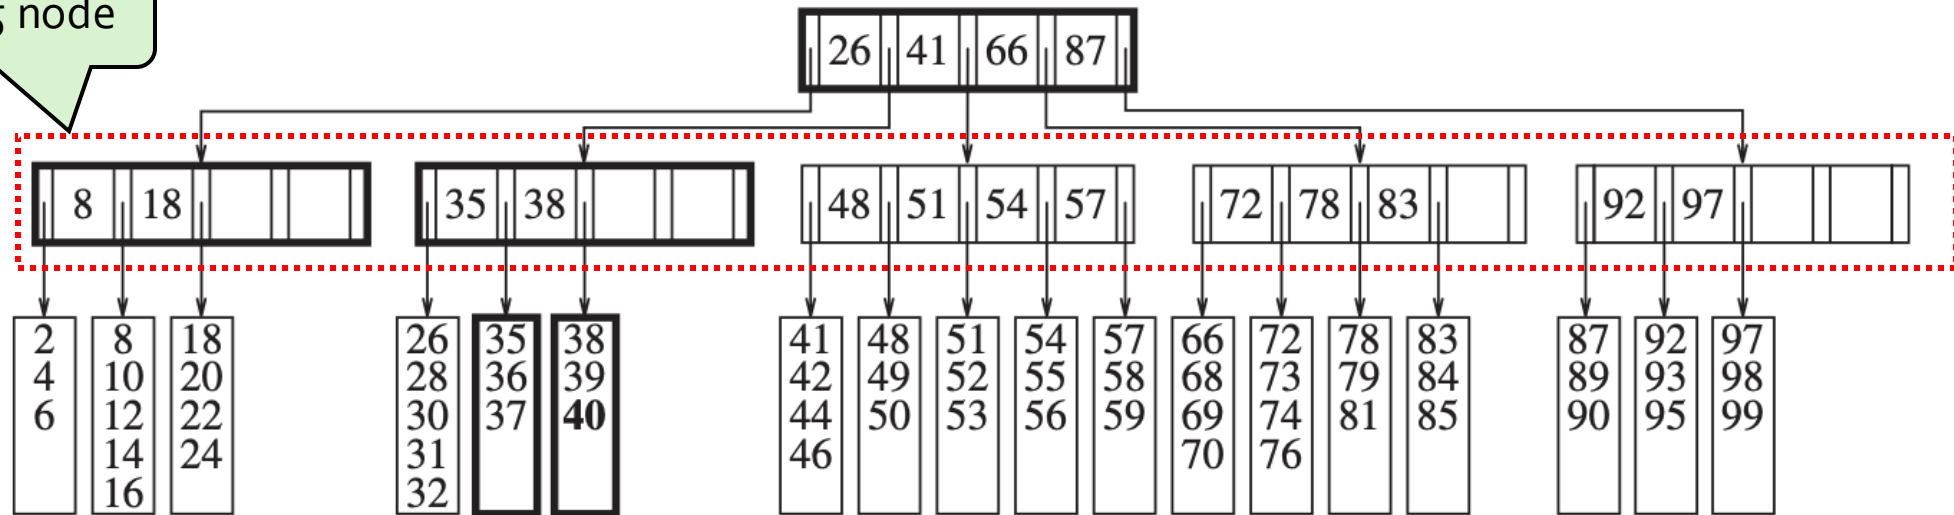


**Figure 4.65** Insertion of 55 into the B-tree in Figure 4.64 causes a split into two leaves

Tree: B-Tree and B+ Tree

# B-tree of order 5: insert(40)

Split into 5 node

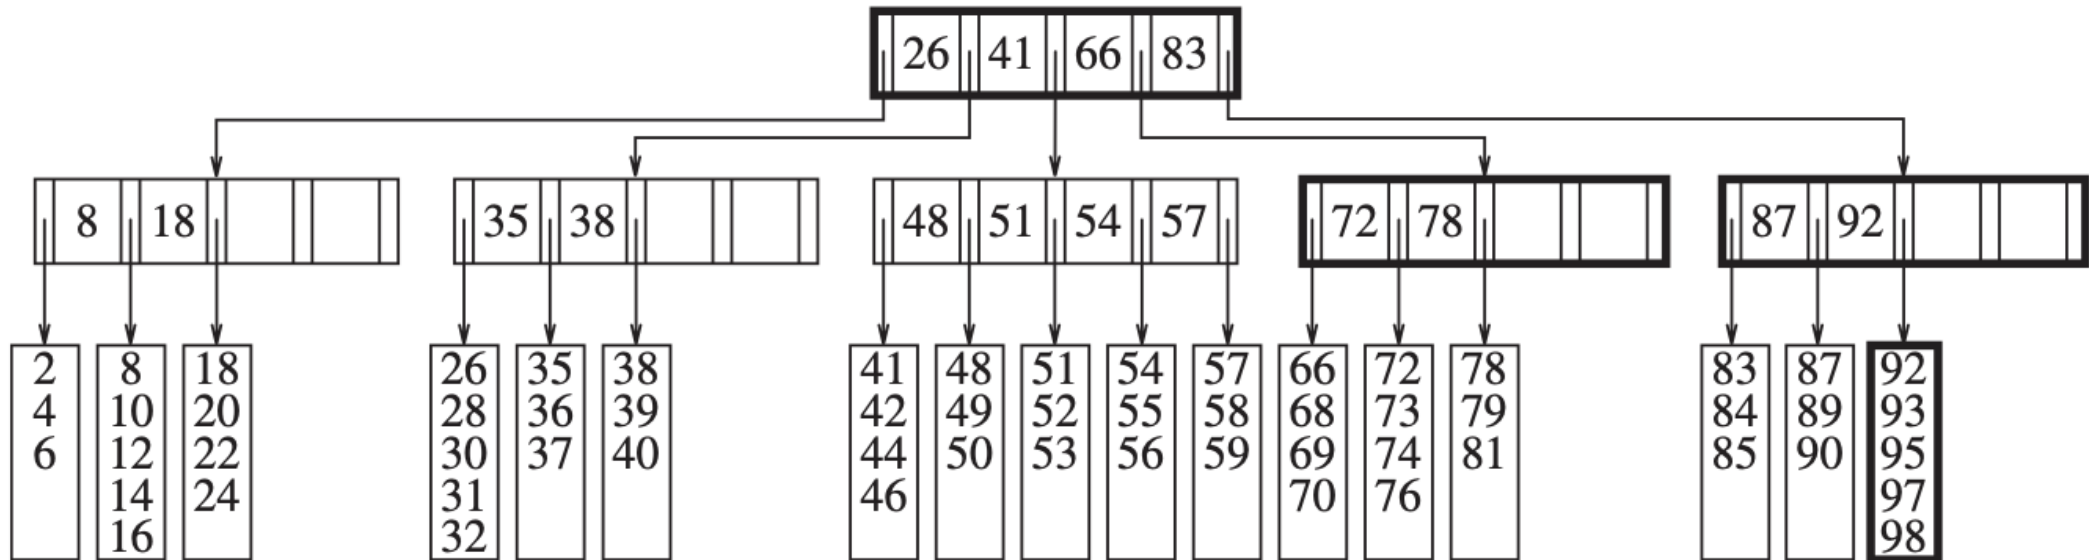


**Figure 4.66** Insertion of 40 into the B-tree in Figure 4.65 causes a split into two leaves and then a split of the parent node



Tree: B-Tree and B+ Tree

# B-tree of order 5: delete(99)



**Figure 4.67** B-tree after the deletion of 99 from the B-tree in Figure 4.66

Tree: Set and Map

# STL set and map classes

---

