

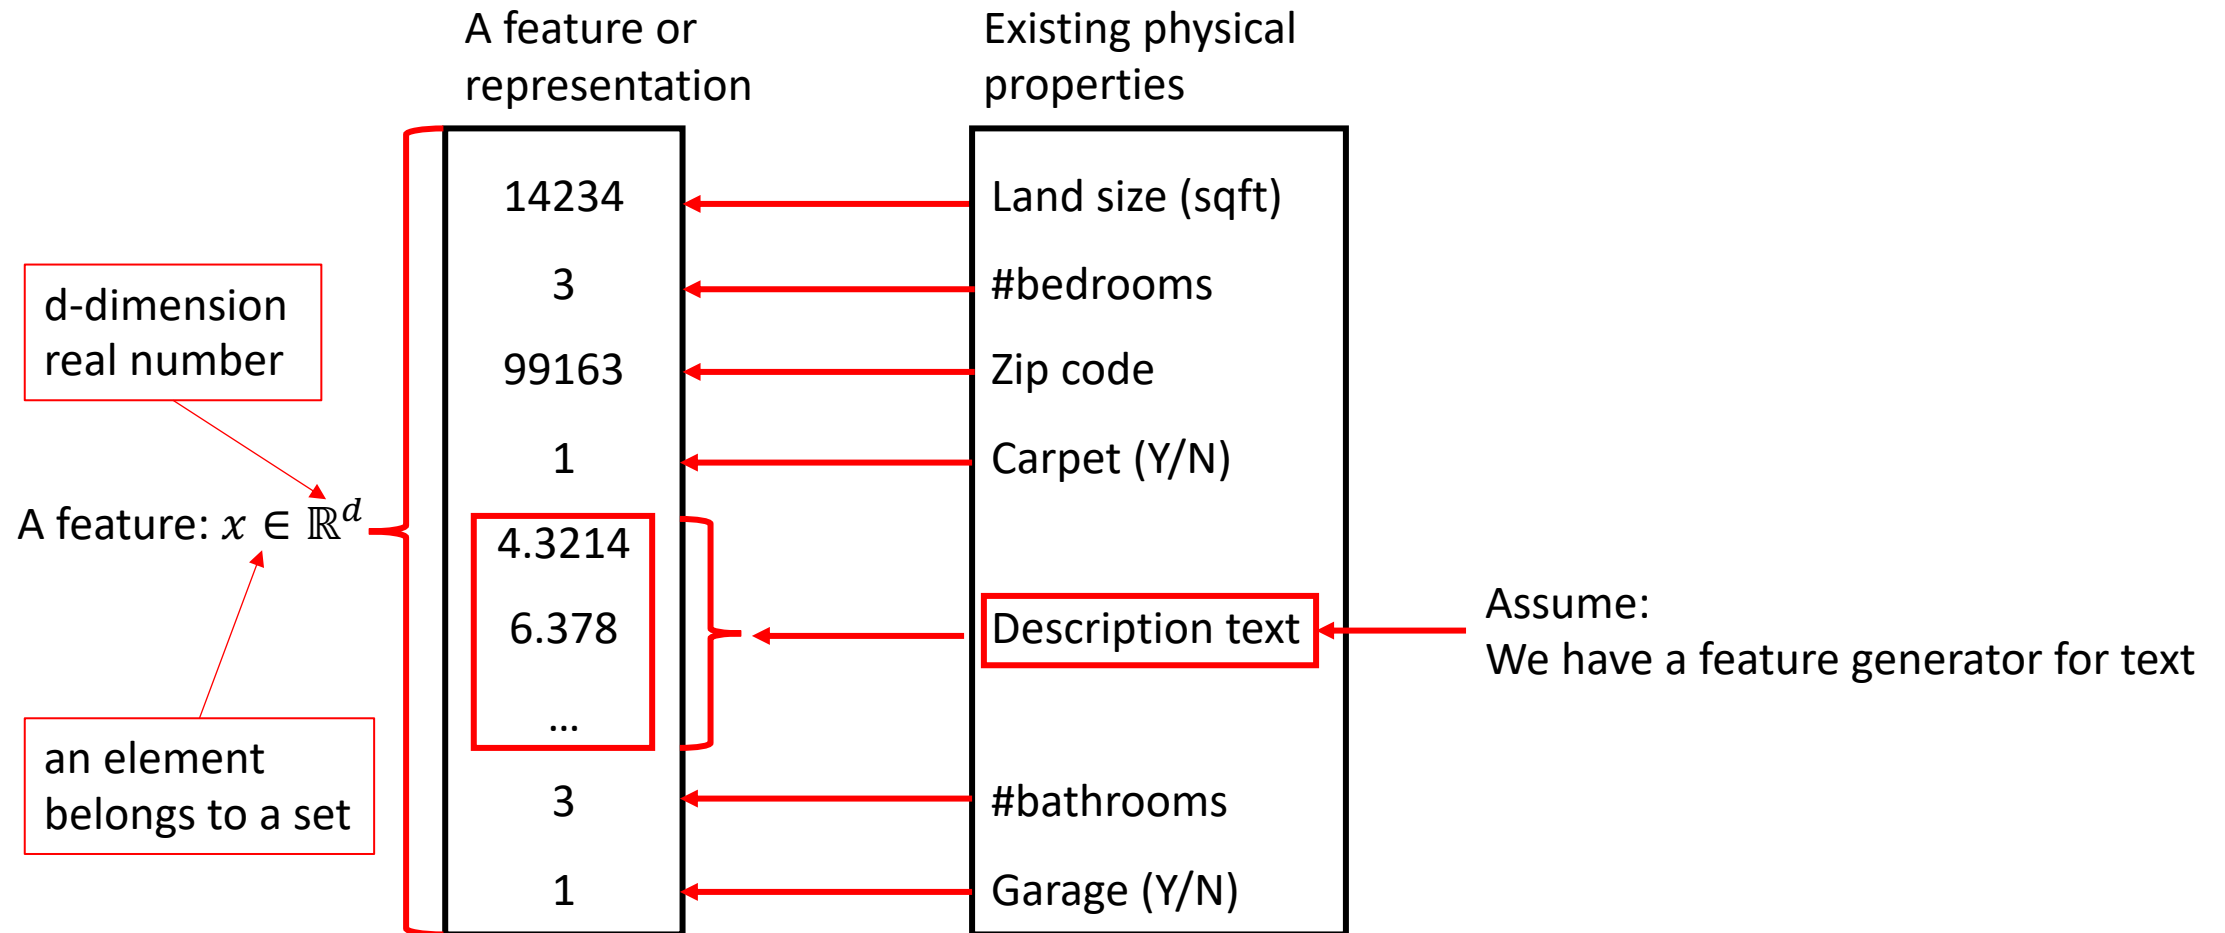
Neural Network Basics

CPT_S 434/534 Neural network design and application

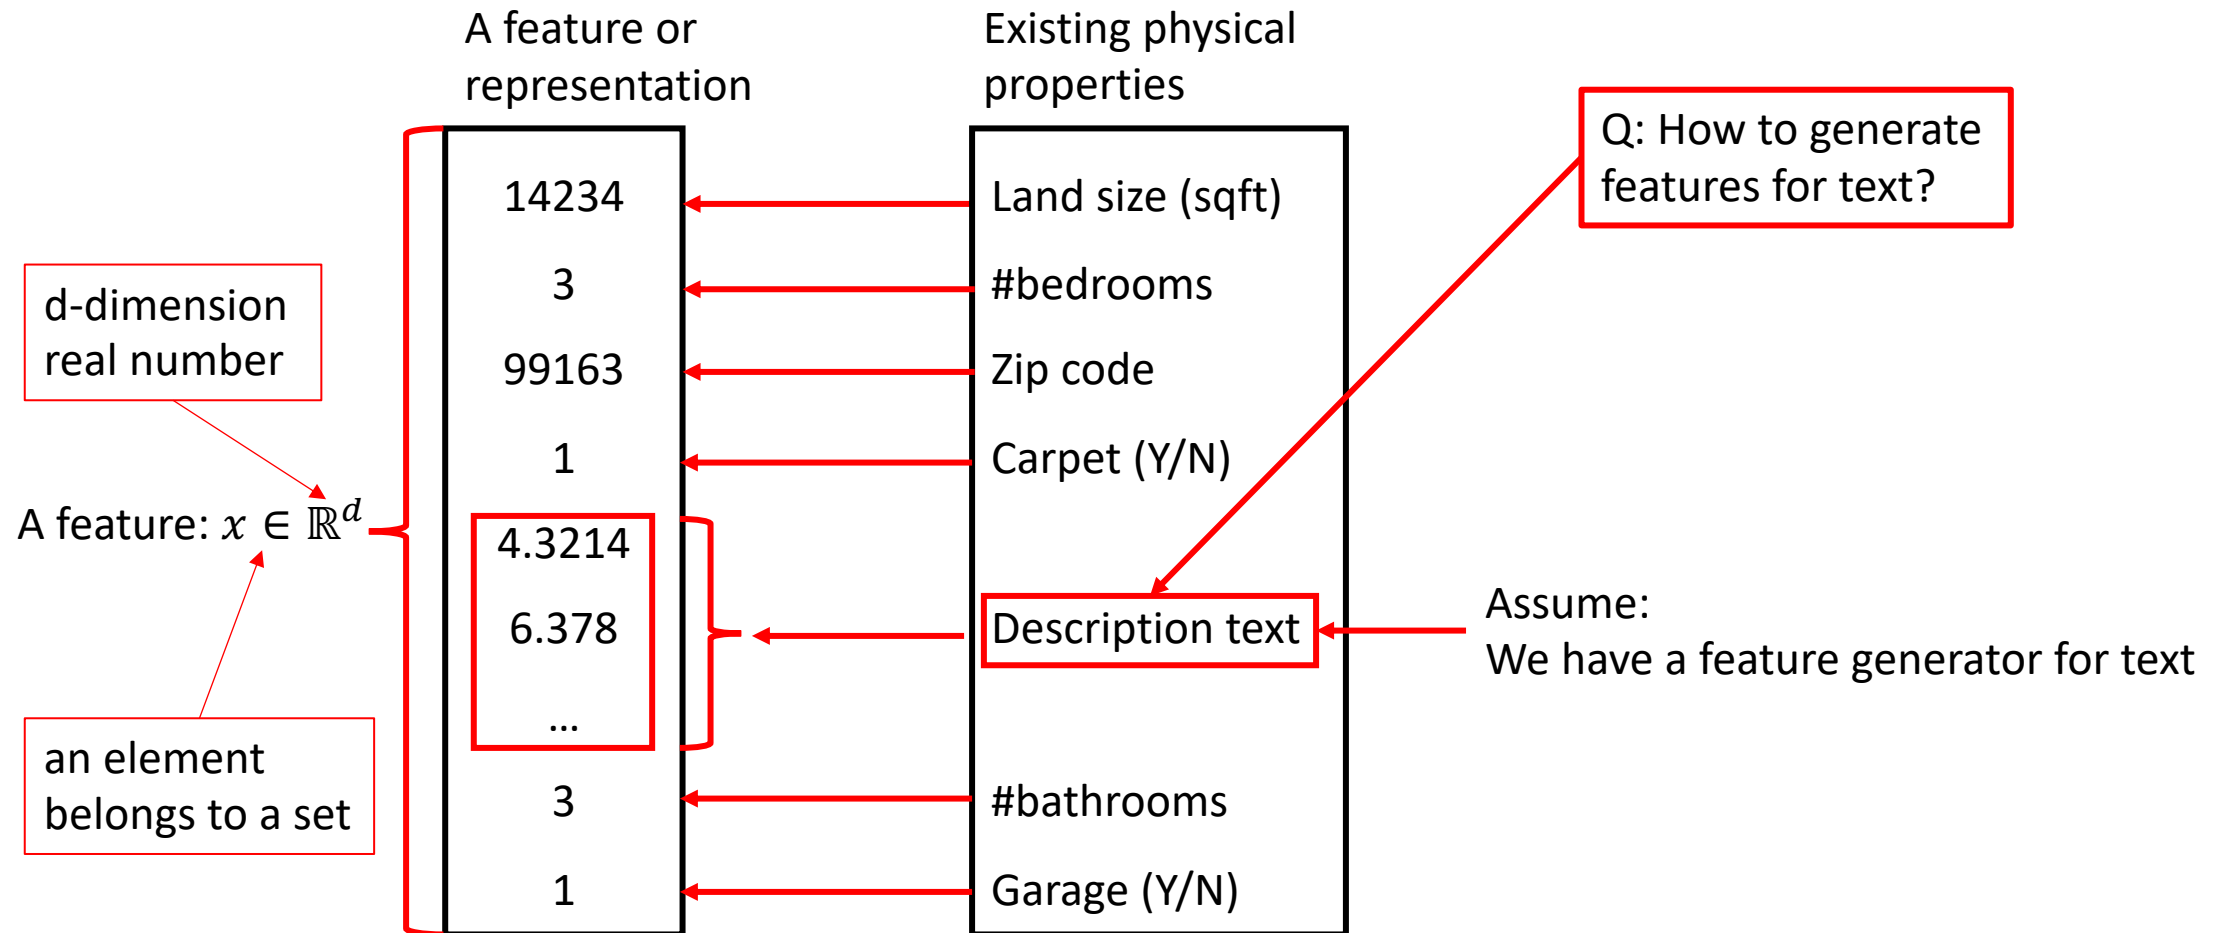
Today's class includes

- Bag-of-words features (hand-crafted)
 - TF-IDF for text data
 - HOG for image data
- History of convolutional neural networks
 - Difference from conventional machine learning methods such as linear models (from the viewpoint of feature generation)
- Feedforward networks: a simple kind of neural networks
 - Typical structure, properties and examples

House price prediction



House price prediction



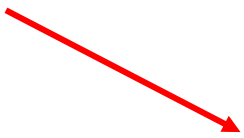
Bag-of-words features

- TF-IDF (term frequency–inverse document frequency)

t: a term

d: a document

D: a set of documents


$$\text{tfidf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$$

Bag-of-words features

- TF-IDF (term frequency–inverse document frequency)

t: a term

d: a document

D: a set of documents

$$\text{idf}(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$$

$$\text{tfidf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$$

Bag-of-words features

- TF-IDF (term frequency–inverse document frequency)

t: a term

d: a document

D: a set of documents

$$\text{tfidf}(t, d, D) = \text{tf}(t, d) \cdot \text{idf}(t, D)$$

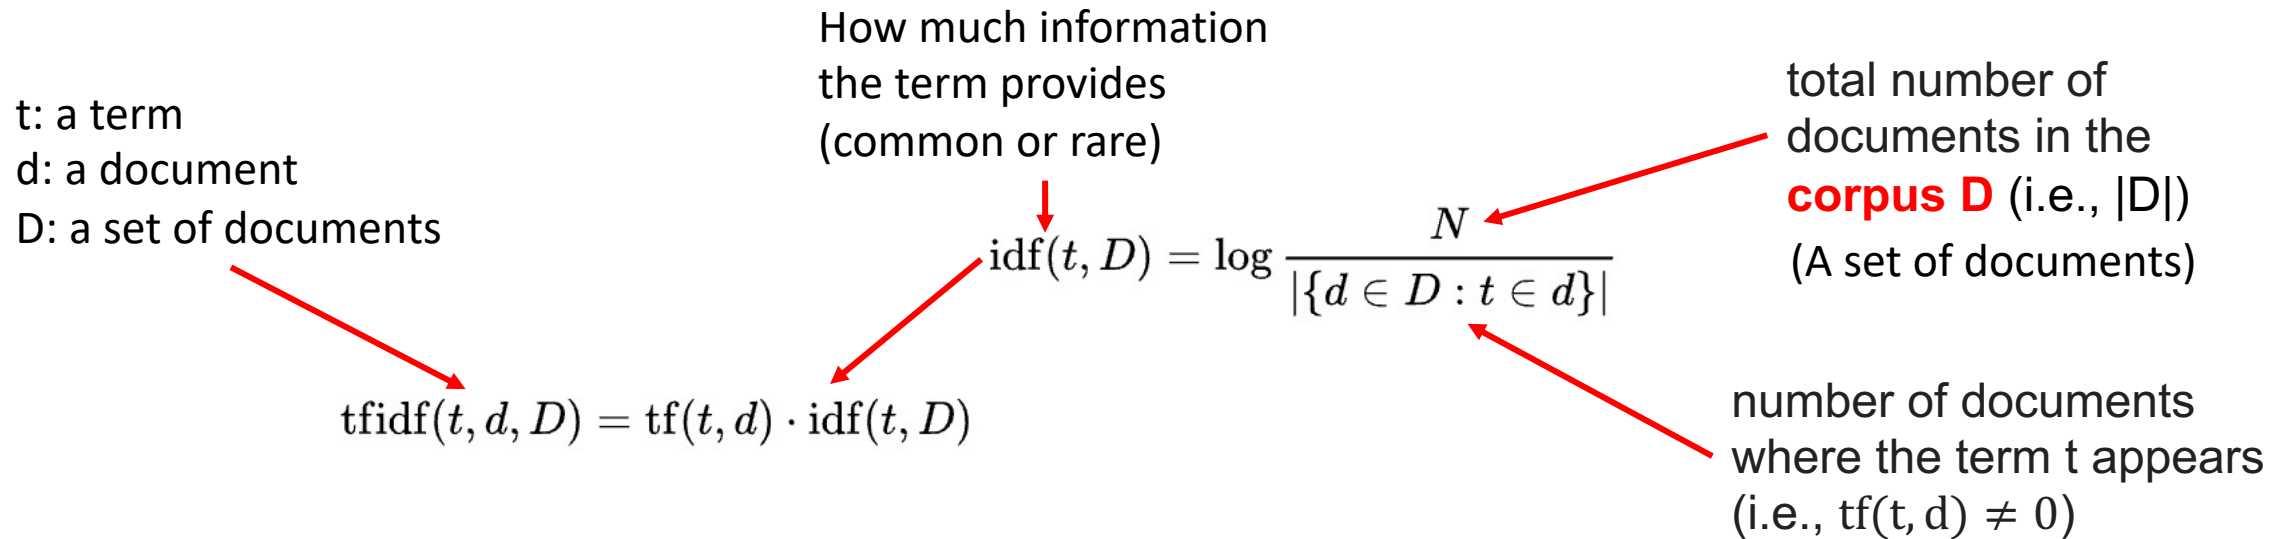
$$\text{idf}(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$$

total number of documents in the **corpus D** (i.e., $|D|$)
(A set of documents)

number of documents where the term t appears
(i.e., $\text{tf}(t, d) \neq 0$)

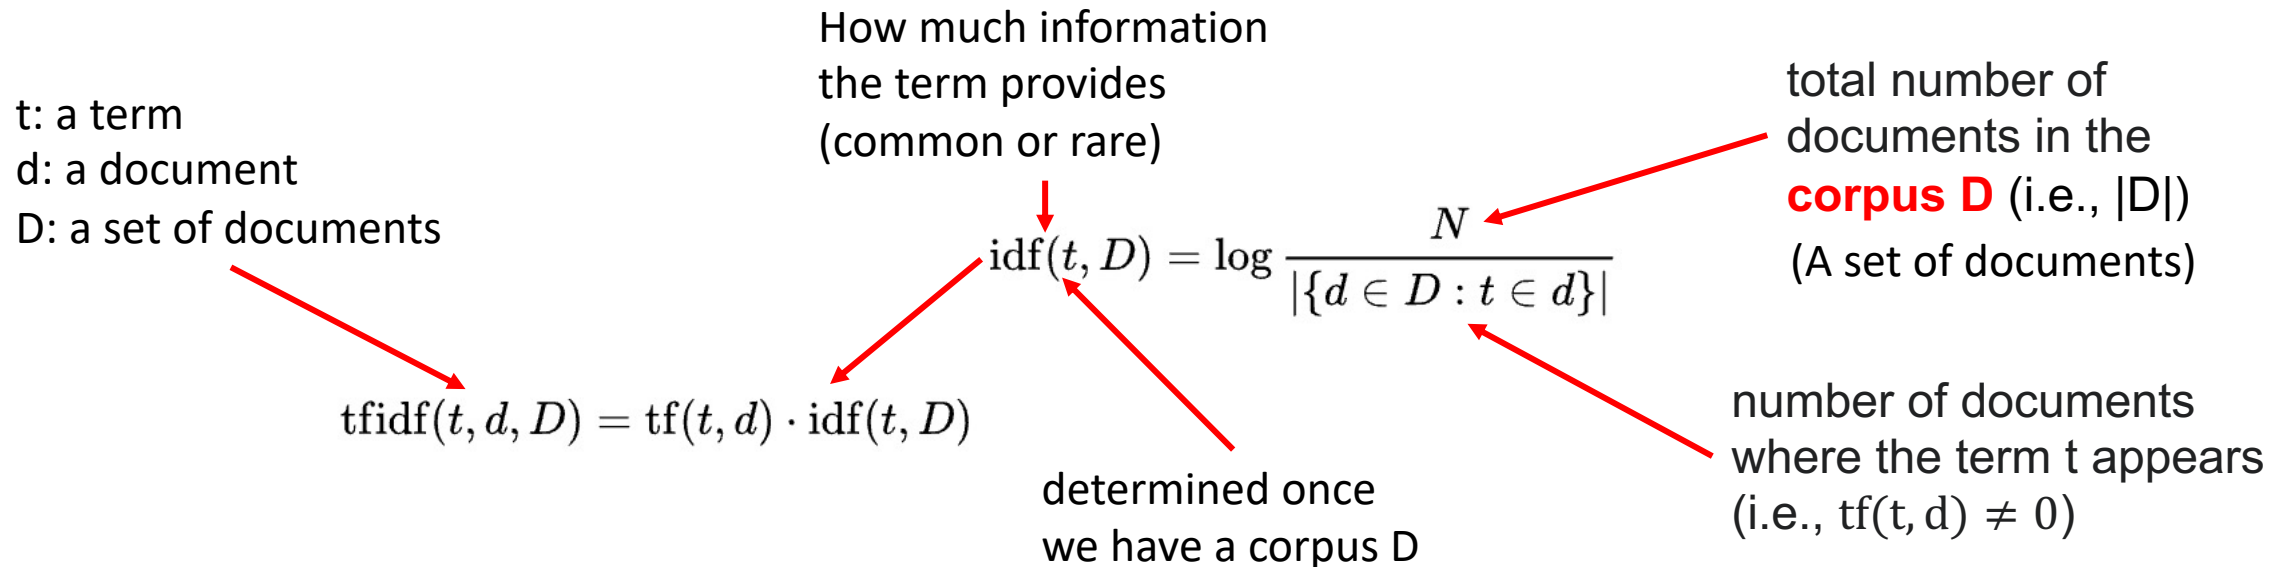
Bag-of-words features

- TF-IDF (term frequency–inverse document frequency)



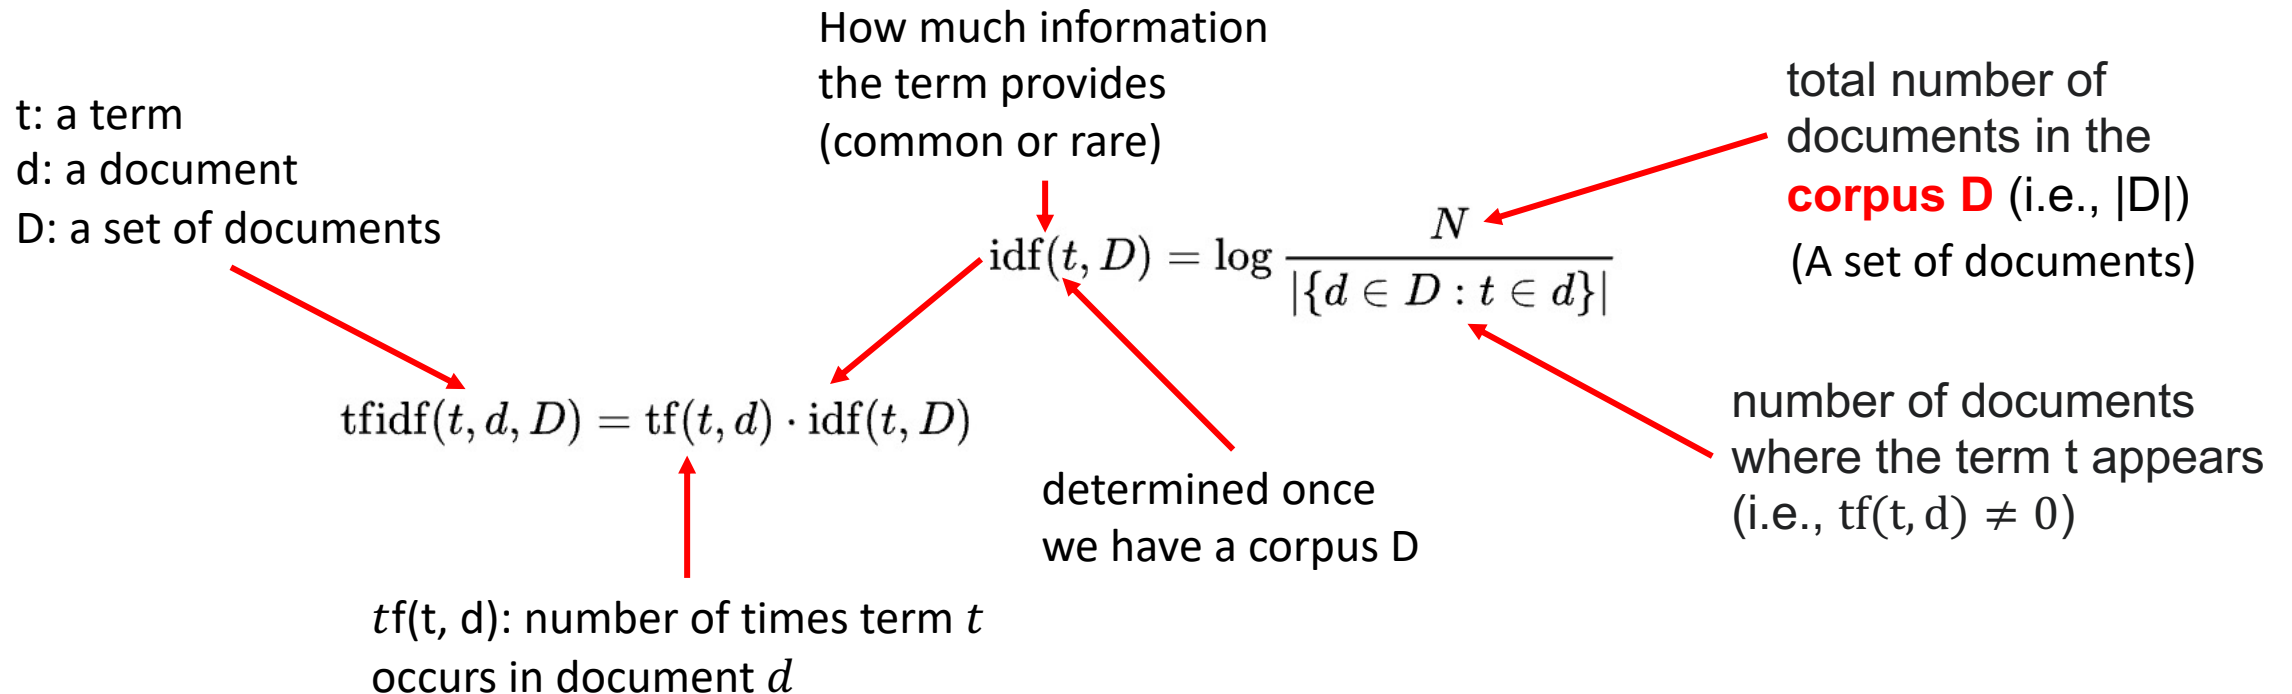
Bag-of-words features

- TF-IDF (term frequency–inverse document frequency)



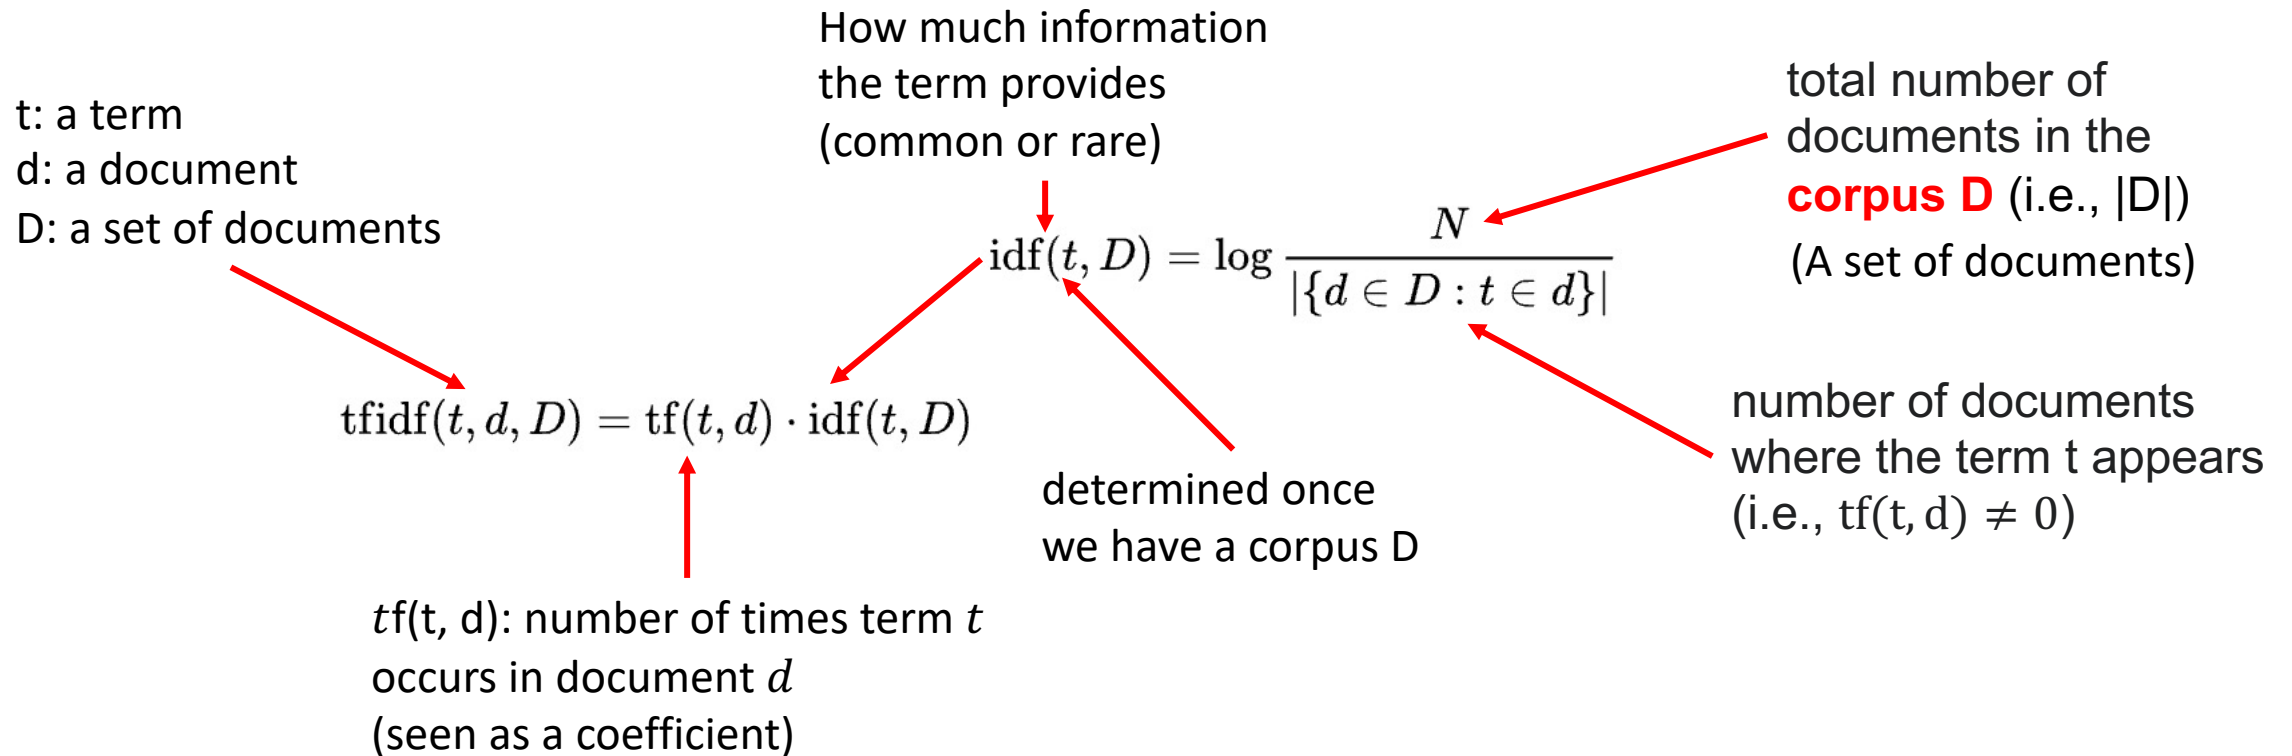
Bag-of-words features

- TF-IDF (term frequency–inverse document frequency)



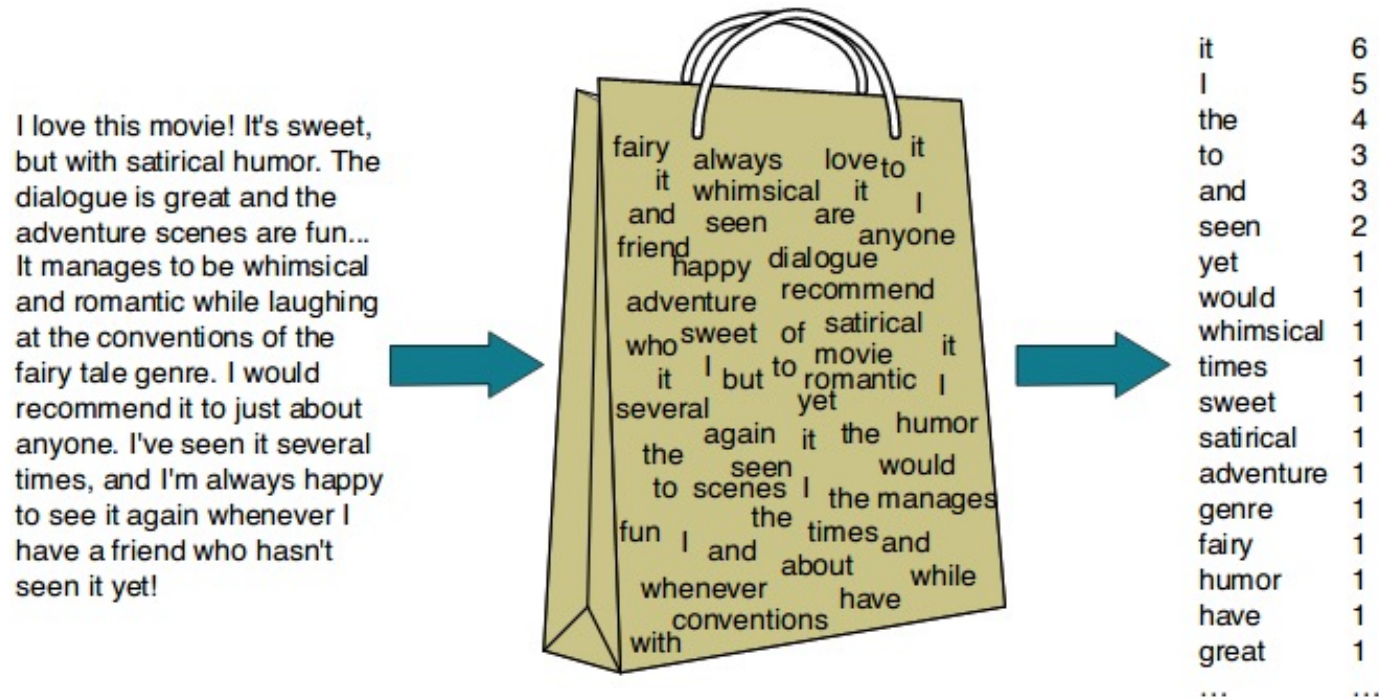
Bag-of-words features

- TF-IDF (term frequency–inverse document frequency)



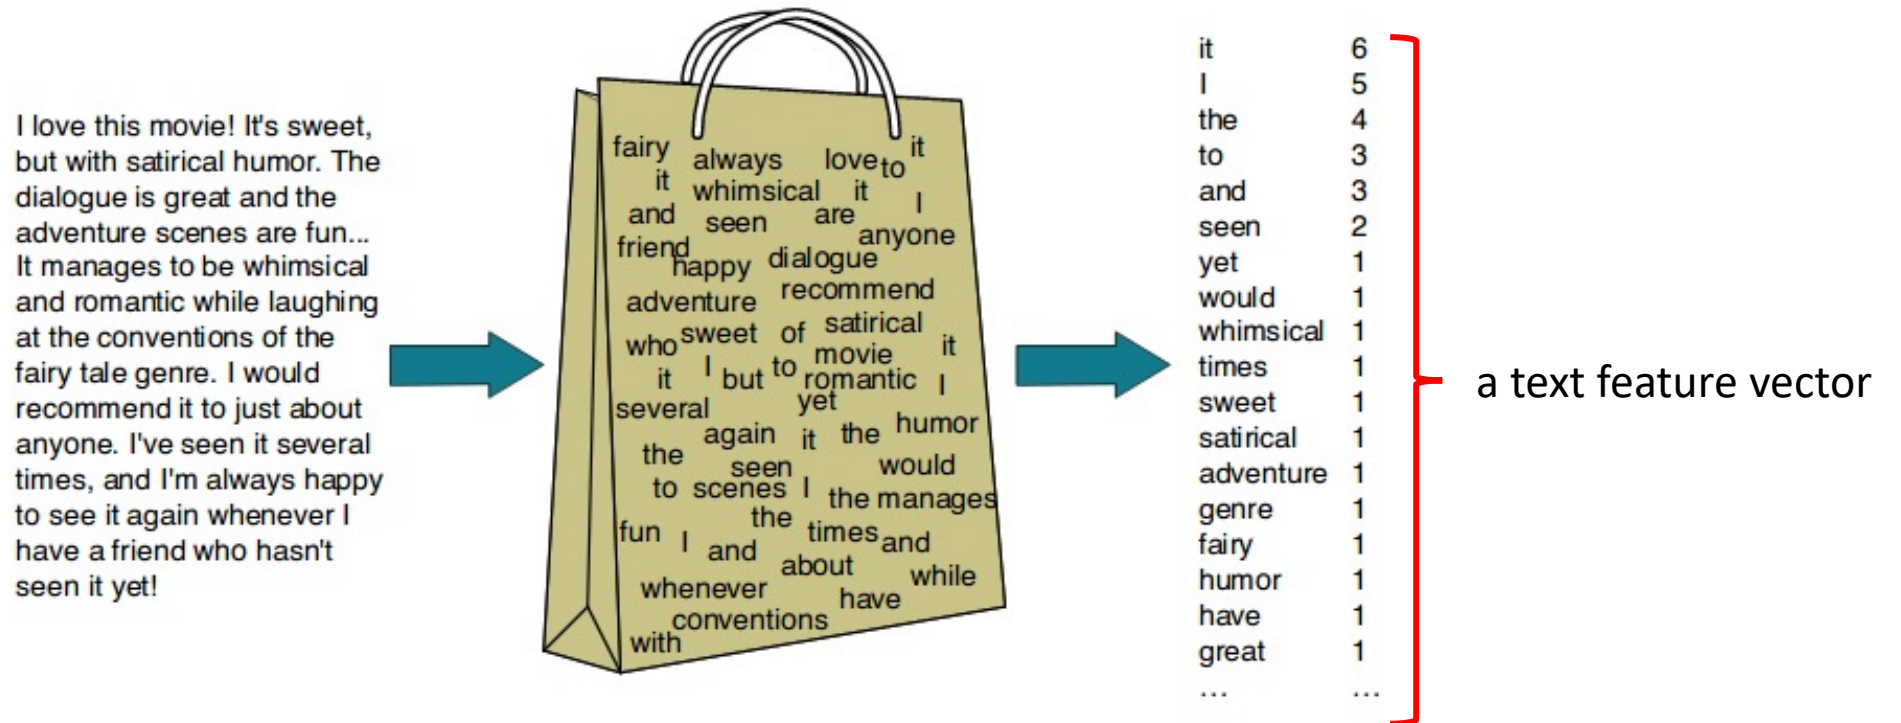
Bag-of-words features

- TF-IDF (term frequency–inverse document frequency)



Bag-of-words features

- TF-IDF (term frequency–inverse document frequency)




Bag-of-words features



Q: How to generate a bag-of-words feature for an image?

Histogram of oriented gradients

- Oriented gradients?
 - Gradients: changes in X and Y directions



Histogram of oriented gradients

- Oriented gradients?
 - Gradients: changes in X and Y directions




Pixel values in the image

121	10	78	96	125
48	152	68	125	111
145	78	85	89	65
154	214	56	200	66
214	87	45	102	45

Histogram of oriented gradients

- Oriented gradients?
 - Gradients: changes in X and Y directions



Pixel values in the image

121	10	78	96	125
48	152	68	125	111
145	78	85	89	65
154	214	56	200	66
214	87	45	102	45


X direction G_x

Subtract the value on the left from the pixel value on the right:

$$G_x = 89 - 78 = 11$$

Histogram of oriented gradients

- Oriented gradients?
 - Gradients: changes in X and Y directions



Pixel values in the image

121	10	78	96	125
48	152	68	125	111
145	78	85	89	65
154	214	56	200	66
214	87	45	102	45

X direction G_x

Subtract the value on the left from the pixel value on the right:

$$G_x = 89 - 78 = 11$$


Y direction G_y

Subtract the pixel value below from the pixel value above the selected pixel:

$$G_y = 68 - 56 = 12$$

Histogram of oriented gradients

- Oriented gradients?
 - Gradients: changes in X and Y directions
 - **Oriented?**



Pixel values in the image

121	10	78	96	125
48	152	68	125	111
145	78	85	89	65
154	214	56	200	66
214	87	45	102	45

X direction G_x

Subtract the value on the left from the pixel value on the right:

$$G_x = 89 - 78 = 11$$


Y direction G_y

Subtract the pixel value below from the pixel value above the selected pixel:

$$G_y = 68 - 56 = 12$$

Histogram of oriented gradients

- Oriented gradients?
 - Gradients: changes in X and Y directions
 - Oriented:



Pixel values in the image

121	10	78	96	125
48	152	68	125	111
145	78	85	89	65
154	214	56	200	66
214	87	45	102	45

X direction G_x
Subtract the value on the left from the pixel value on the right:


$$G_x = 89 - 78 = 11$$

Y direction G_y
Subtract the pixel value below from the pixel value above the selected pixel:

$$G_y = 68 - 56 = 12$$

Histogram of oriented gradients

- Oriented gradients?
 - Gradients: changes in X and Y directions
 - Oriented:



Pixel values in the image

121	10	78	96	125
48	152	68	125	111
145	78	85	89	65
154	214	56	200	66
214	87	45	102	45

X direction G_x

Subtract the value on the left from the pixel value on the right:

$$G_x = 89 - 78 = 11$$


Y direction G_y

Subtract the pixel value below from the pixel value above the selected pixel:

$$G_y = 68 - 56 = 12$$

Histogram of oriented gradients

- Oriented gradients?
 - Gradients: changes in X and Y directions
 - Oriented:




Pixel values in the image

121	10	78	96	125
48	152	68	125	111
145	78	85	89	65
154	214	56	200	66
214	87	45	102	45

X direction G_x
Subtract the value on the left from the pixel value on the right:
 $G_x = 89 - 78 = 11$

Y direction G_y
Subtract the pixel value below from the pixel value above the selected pixel:
 $G_y = 68 - 56 = 12$

Histogram of oriented gradients



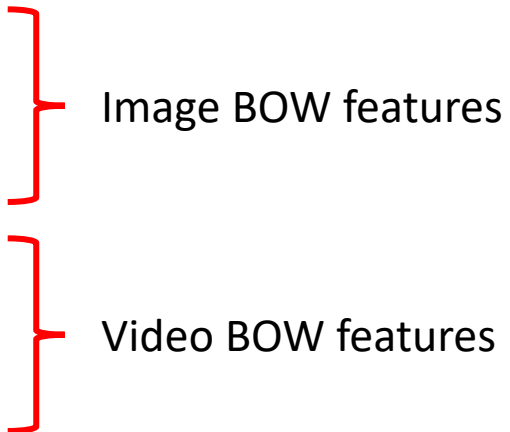
Histogram of oriented gradients

121	10	78	96	125
48	152	68	125	111
145	78	85	89	65
154	214	56	200	66
214	87	45	102	45

How many pixels with the corresponding value of angle \rightarrow a vector feature


Frequency						1										
Angle (Φ)	1	2	3	4 ...	35	36	37	38	39....		175	176	177	178	179	180

Many other hand-crafted features

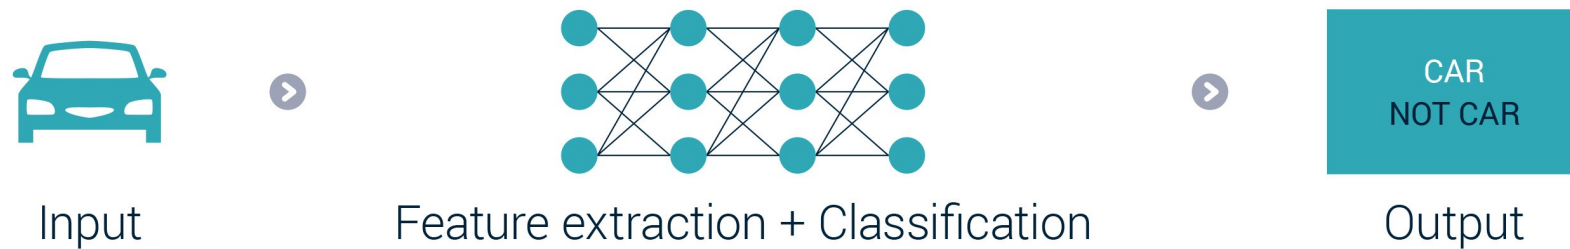
- Scale Invariant Feature Transform (SIFT)
 - Speeded-Up Robust Feature (SURF)
 - Histogram of Optical Flow (HOF)
 - Motion Boundary Histogram (MBH)
 - Fisher vector (FV, a similarity-based function)
 - Vector of Locally Aggregated Descriptors (VLAD)
 -
- 
- The diagram uses red brackets to categorize the features. A bracket on the right side groups the first two features, SIFT and SURF, under the label 'Image BOW features'. Another bracket groups the next two features, HOF and MBH, under the label 'Video BOW features'.
- Image BOW features
- Video BOW features

End-to-end training of neural networks

Machine Learning



Deep Learning



LeNet-5 in 1999




Fig. 1. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

ImageNet classification challenge 2012

Task 1

Team name	Filename	Error (5 guesses)	Description
SuperVision	test-preds-141-146.2009-131-137-145-146.2011-145f.	0.15315	Using extra training data from ImageNet Fall 2011 release
SuperVision	test-preds-131-137-145-135-145f.txt	0.16422	Using only supplied training data
ISI	pred_FVs_wLACs_weighted.txt	0.26172	Weighted sum of scores from each classifier with SIFT+FV, LBP+FV, GIST+FV, and CSIFT+FV, respectively.
ISI	pred_FVs_weighted.txt	0.26602	Weighted sum of scores from classifiers using each FV.
ISI	pred_FVs_summed.txt	0.26646	Naive sum of scores from classifiers using each FV.

AlexNet

ImageNet classification challenge 2012

Task 1

Team name	Filename	Error (5 guesses)	Description
SuperVision	test-preds-141-146.2009-131-137-145-146.2011-145f.	0.15315	Using extra training data from ImageNet Fall 2011 release
SuperVision	test-preds-131-137-145-135-145f.txt	0.16422	Using only supplied training data
ISI	pred_FVs_wLACs_weighted.txt	0.26172	Weighted sum of scores from each classifier with SIFT+FV, LBP+FV, GIST+FV, and CSIFT+FV, respectively.
ISI	pred_FVs_weighted.txt	0.26602	Weighted sum of scores from classifiers using each FV.
ISI	pred_FVs_summed.txt	0.26646	Naive sum of scores from classifiers using each FV.

AlexNet

AlexNet


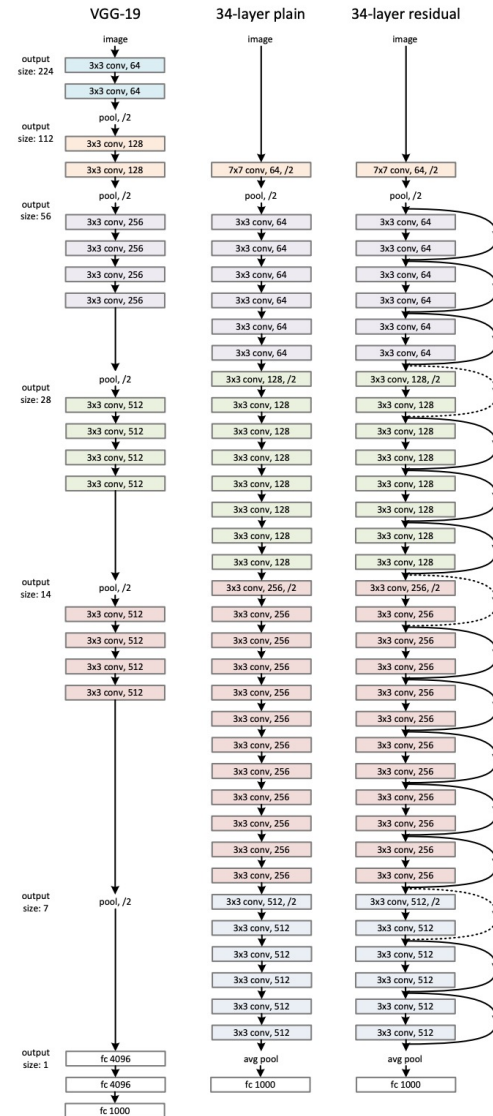


Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." *Advances in neural information processing systems* 25 (2012): 1097-1105.

VGG-19 and ResNet-34



[ResNet]


Feedforward networks

- Or multilayer perceptrons (MLPs)

$$f(w; x_i) = x_i'w + b \rightarrow y_i \quad \text{Linear model}$$

Feedforward networks

- Or multilayer perceptrons (MLPs)

$$f(w; x_i) = x_i'w + b \rightarrow y_i \quad \text{Linear model}$$


Weight Bias

Feedforward networks

- Or multilayer perceptrons (MLPs)


$$f(w; x_i) = x_i'w + b \rightarrow y_i \quad \text{Linear model}$$

Weight Bias

$$f(w; x_i)$$


Feedforward networks

- Or multilayer perceptrons (MLPs)




Feedforward networks

- Or multilayer perceptrons (MLPs)




Feedforward networks

- Or multilayer perceptrons (MLPs)




Feedforward networks

- Or multilayer perceptrons (MLPs)




Feedforward networks

- Or multilayer perceptrons (MLPs)



Feedforward networks

- Or multilayer perceptrons (MLPs)



Feedforward networks

- Deep:
 - Many compositional layers

Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear

Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear

e.g., activation layers

Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions

Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions

$$f_m \left(\dots \left(f_2 \left(f_1(w; x_i) \right) \right) \right) \rightarrow y_i$$

Composition may not preserve convexity

Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions
 - Some functions f_i can be nonconvex

$$f_m \left(\dots \left(f_2 \left(f_1(w; x_i) \right) \right) \right) \rightarrow y_i$$

Composition may not preserve convexity

Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions
 - Some functions f_i can be nonconvex
- Feedforward **Q**: why this name?


Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions
 - Some functions f_i can be nonconvex
- Feedforward
 - Information feedforward from input to output layer

Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions
 - Some functions f_i can be nonconvex
- Feedforward
 - Information feedforward from input to output layer

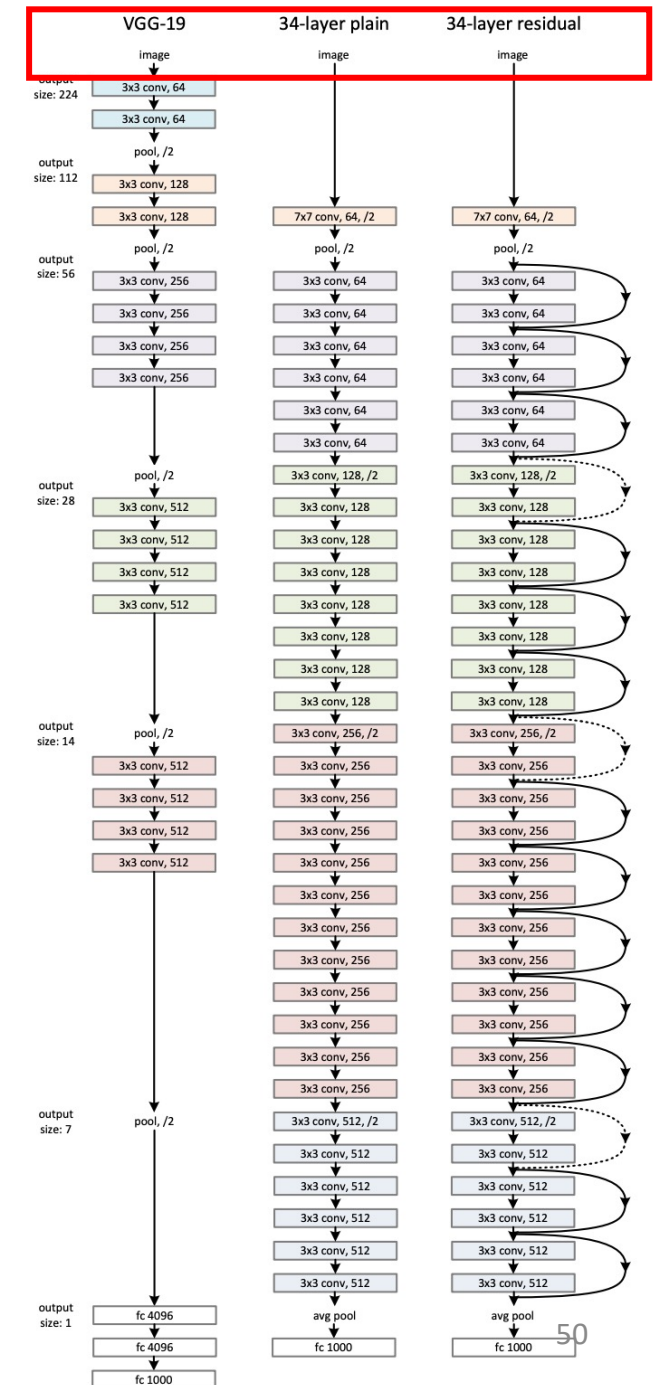
$$f_m \left(\dots \left(f_2(f_1(w; x_i)) \right) \right) = \hat{y}_i$$



Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions
 - Some functions f_i can be nonconvex
- Feedforward
 - Information feedforward from input to output layer

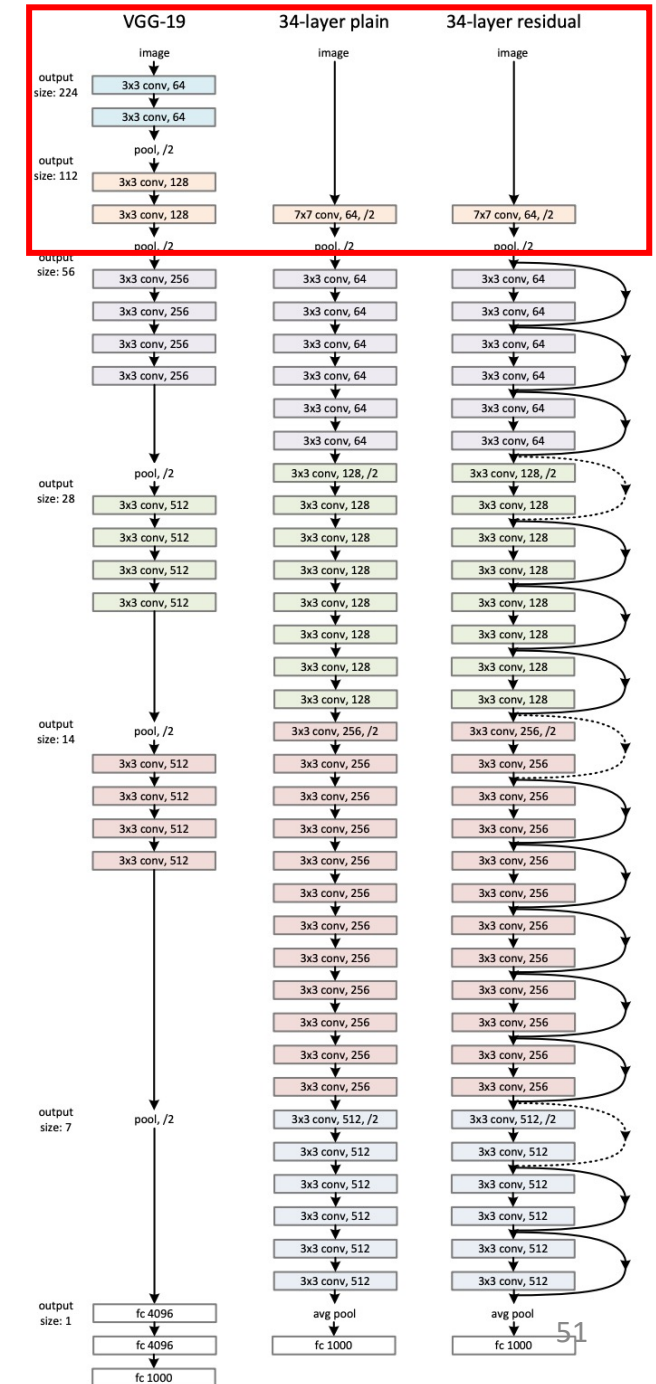
$$f_m \left(\dots \left(f_2 \left(f_1(w; \mathbf{x}_i) \right) \right) \right) = \hat{y}_i$$



Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions
 - Some functions f_i can be nonconvex
- Feedforward
 - Information feedforward from input to output layer

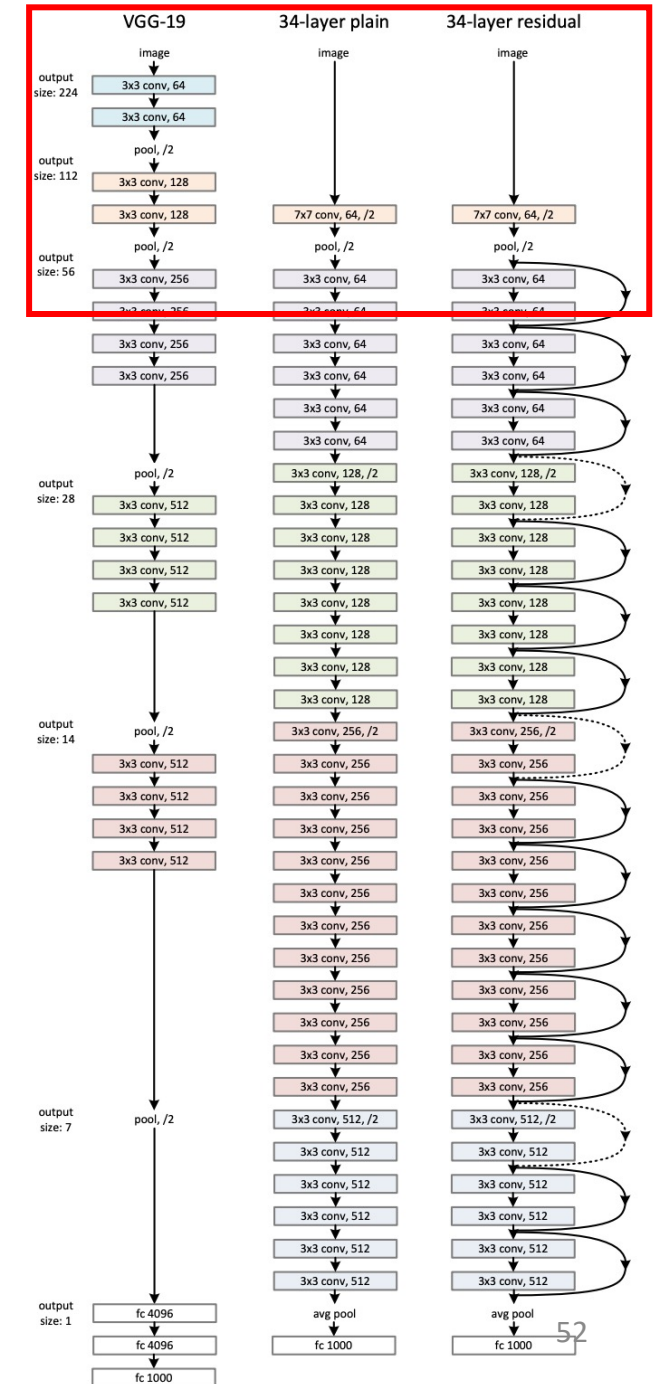
$$f_m \left(\dots \left(f_2 \left(f_1(w; x_i) \right) \right) \right) = \hat{y}_i$$



Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions
 - Some functions f_i can be nonconvex
- Feedforward
 - Information feedforward from input to output layer

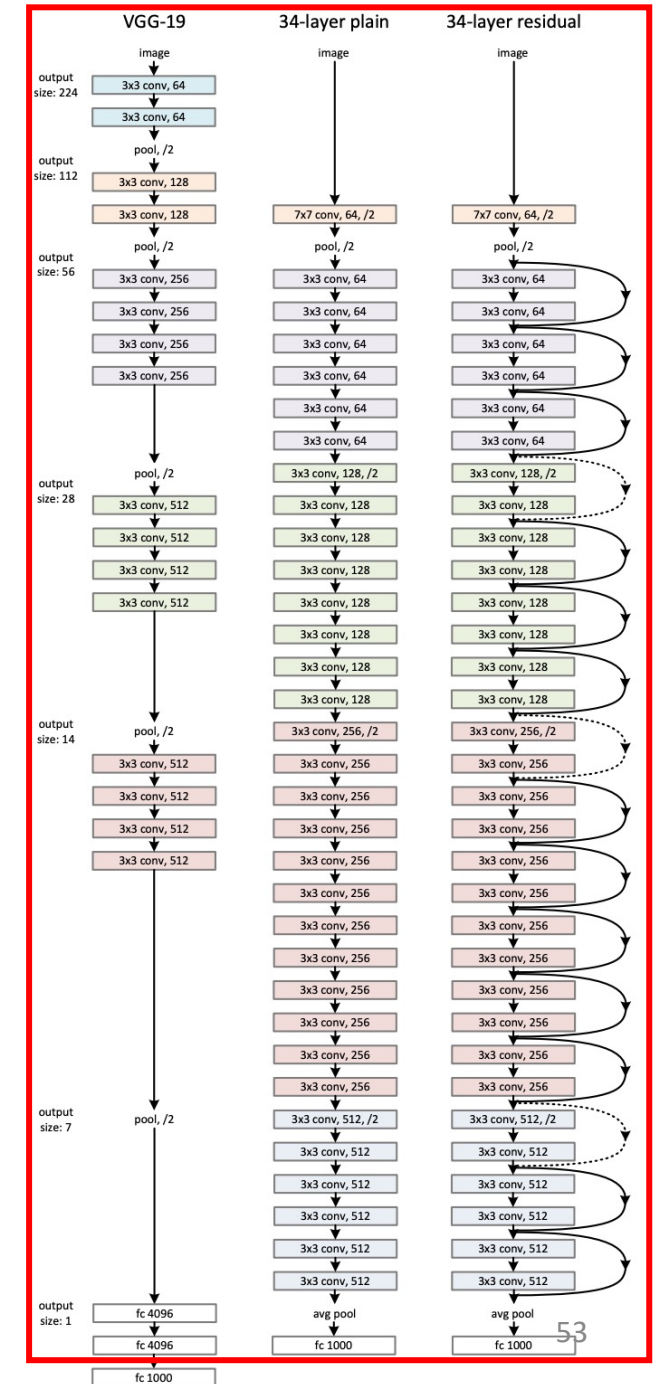
$$f_m \left(\dots \left(f_2(f_1(w; x_i)) \right) \right) = \hat{y}_i$$



Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions
 - Some functions f_i can be nonconvex
- Feedforward
 - Information feedforward from input to output layer

$$f_m \left(\dots \left(f_2 \left(f_1(w; x_i) \right) \right) \right) = \hat{y}_i$$



Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions
 - Some functions f_i can be nonconvex
- Feedforward
 - Information feedforward from input to output layer

Q: Any non-feedforward networks?

Feedforward networks

- Deep:
 - Many compositional layers
- Nonlinearity
 - Some functions f_i can be nonlinear
- Nonconvexity
 - Composition of functions
 - Some functions f_i can be nonconvex
- Feedforward
 - Information feedforward from input to output layer

Q: Any non-feedforward networks?
Contains circles → recurrent networks

References

- [ResNet] He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. "Deep residual learning for image recognition." In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770-778. 2016.