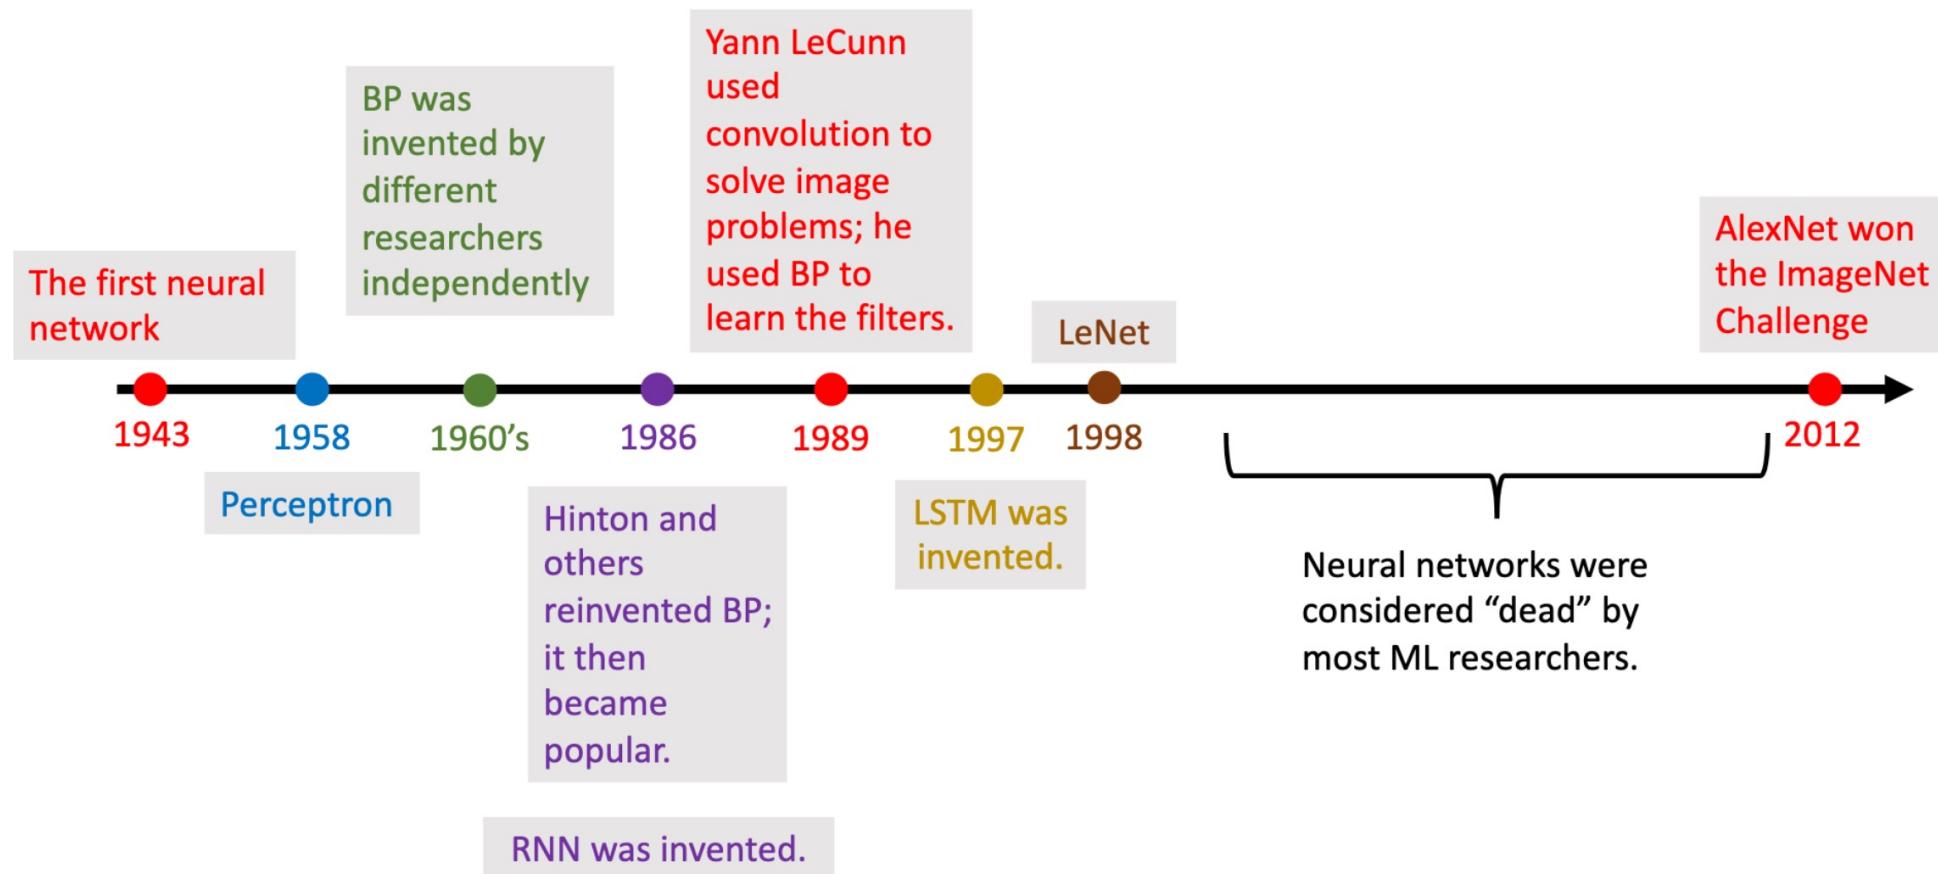


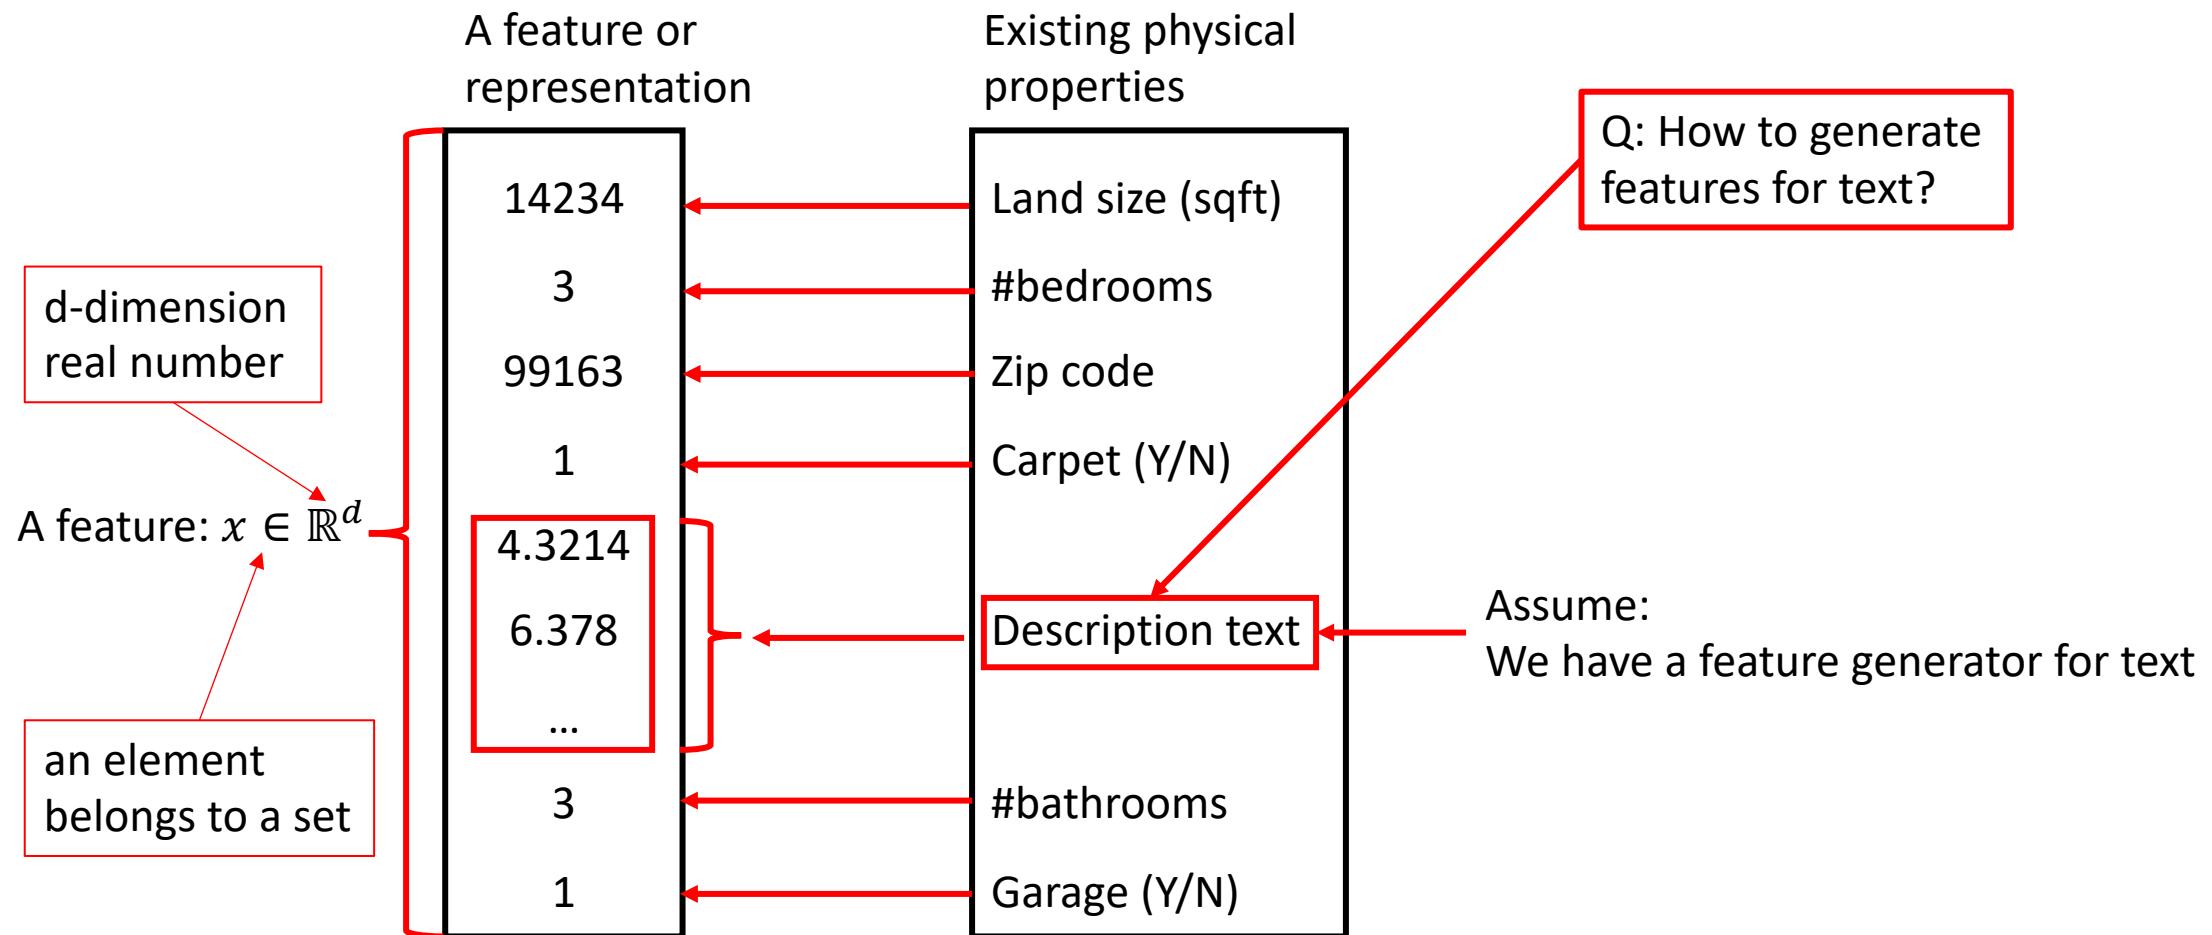
# Convolutional Neural Networks

Neural Networks Design And Application

# History

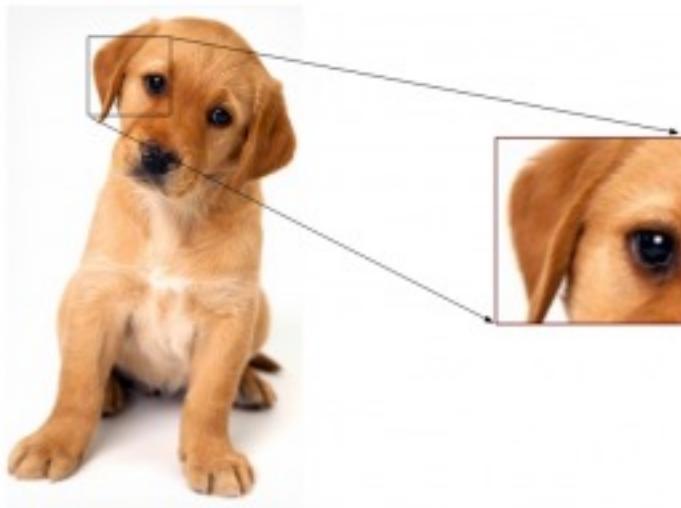


# Review: house price prediction



# Review: histogram of oriented gradients

- Oriented gradients?
  - Gradients: changes in X and Y directions
  - Oriented:



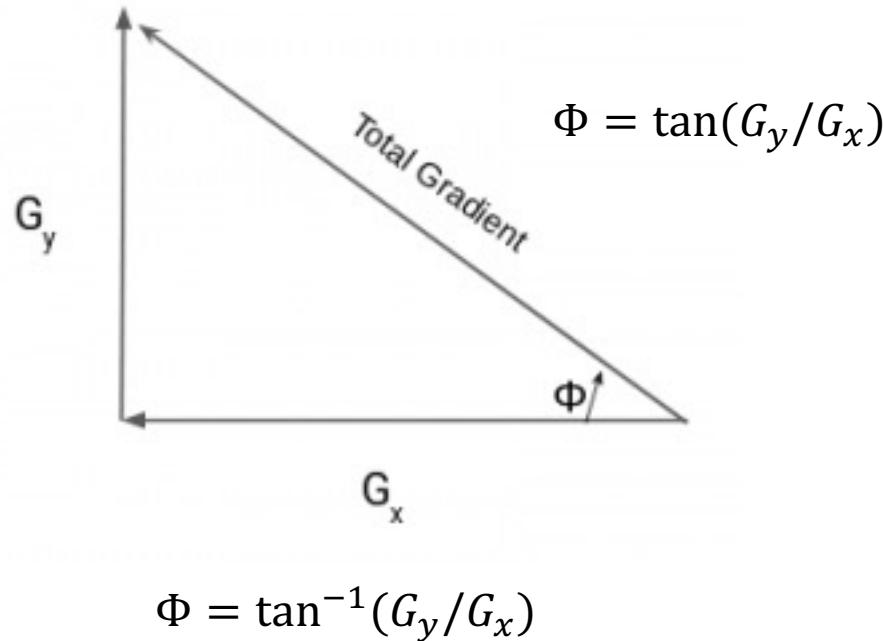
X direction  $G_x$   
Subtract the value on the  
left from the pixel value  
on the right:  
 $89 - 78 = 11$

121	10	78	96	125
48	152	68	125	111
145	78	85	89	65
154	214	56	200	66
214	87	45	102	45

Y direction  $G_y$   
Subtract the pixel value  
below from the pixel value  
above the selected pixel:  
 $68 - 56 = 8$

# Review: histogram of oriented gradients

- Oriented gradients?
  - Gradients: changes in X and Y directions
  - Oriented:



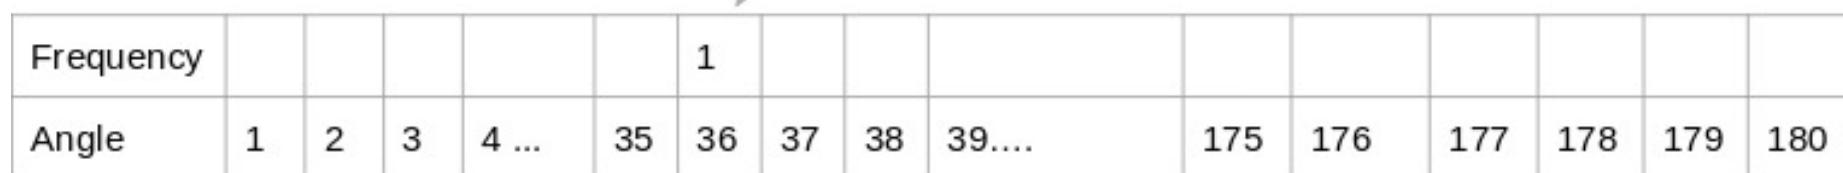
X direction  $G_x$   
Subtract the value on the  
left from the pixel value  
on the right:  
 $89 - 78 = 11$

121	10	78	96	125
48	152	68	125	111
145	78	85	89	65
154	214	56	200	66
214	87	45	102	45

Y direction  $G_y$   
Subtract the pixel value  
below from the pixel value  
above the selected pixel:  
 $68 - 56 = 8$

# Review: histogram of oriented gradients

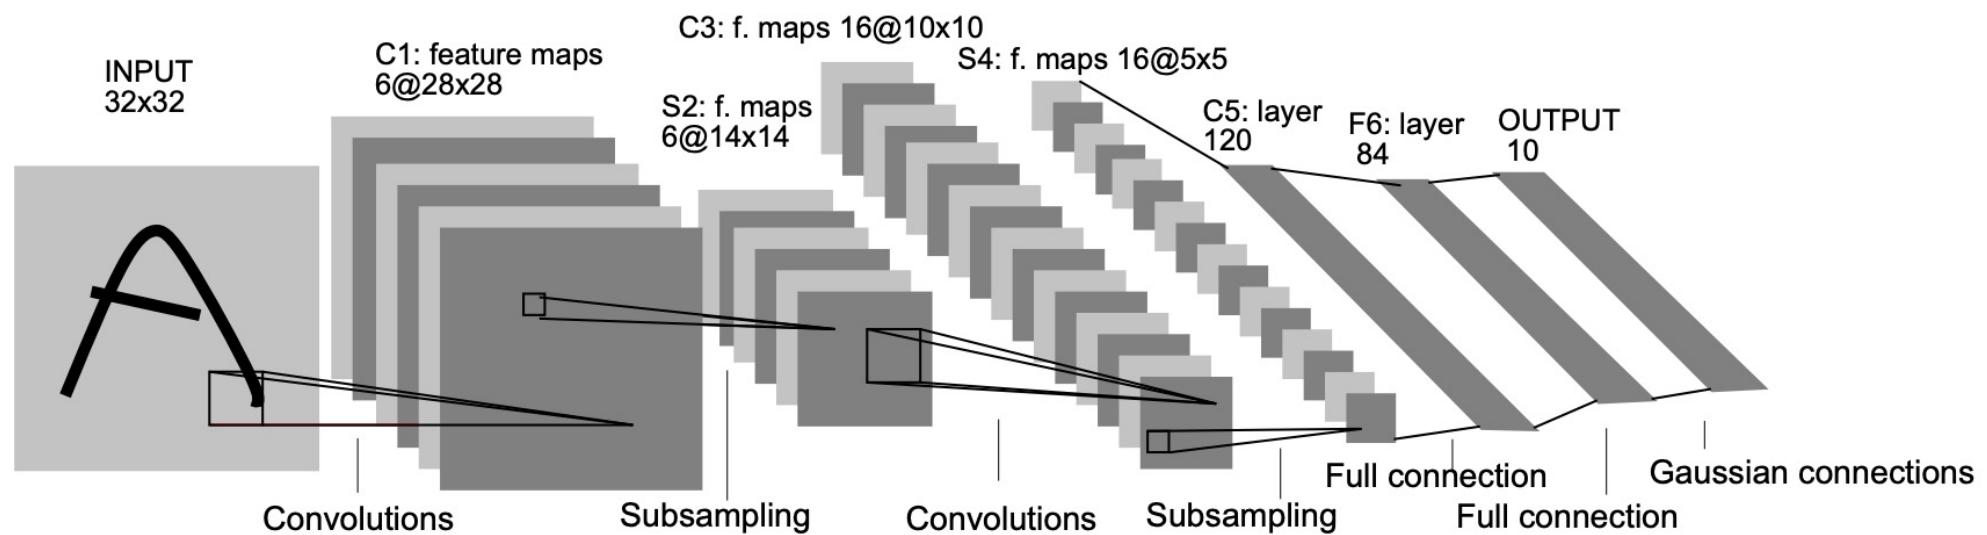
121	10	78	96	125
48	152	68	125	111
145	78	85	89	65
154	214	56	200	66
214	87	45	102	45



# Review: ImageNet challenge 2012

Task 1				
Team name	Filename	Error (5 guesses)		Description
SuperVision	test-preds-141-146.2009-131-137-145-146.2011-145f.	0.15315		Using extra training data from ImageNet Fall 2011 release
SuperVision	test-preds-131-137-145-135-145f.txt	0.16422		Using only supplied training data
ISI	pred_FVs_wLACs_weighted.txt	0.26172		Weighted sum of scores from each classifier with SIFT+FV, LBP+FV, GIST+FV, and CSIFT+FV, respectively.
ISI	pred_FVs_weighted.txt	0.26602		Weighted sum of scores from classifiers using each FV.
ISI	pred_FVs_summed.txt	0.26646		Naive sum of scores from classifiers using each FV.

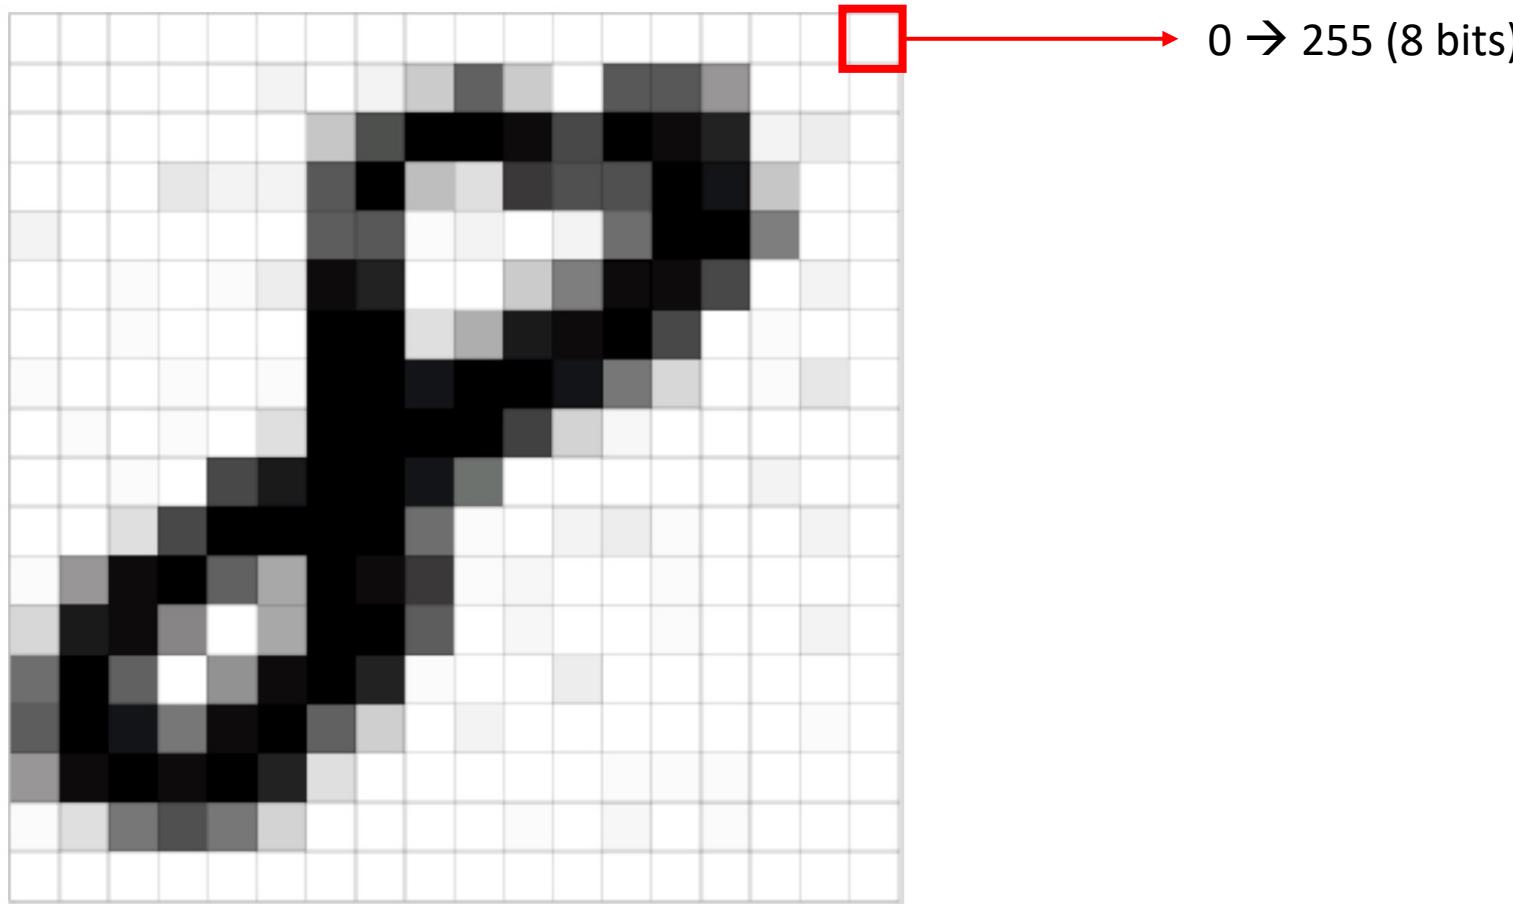
# Review: LeNet-5 in 1999



**Fig. 1.** Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

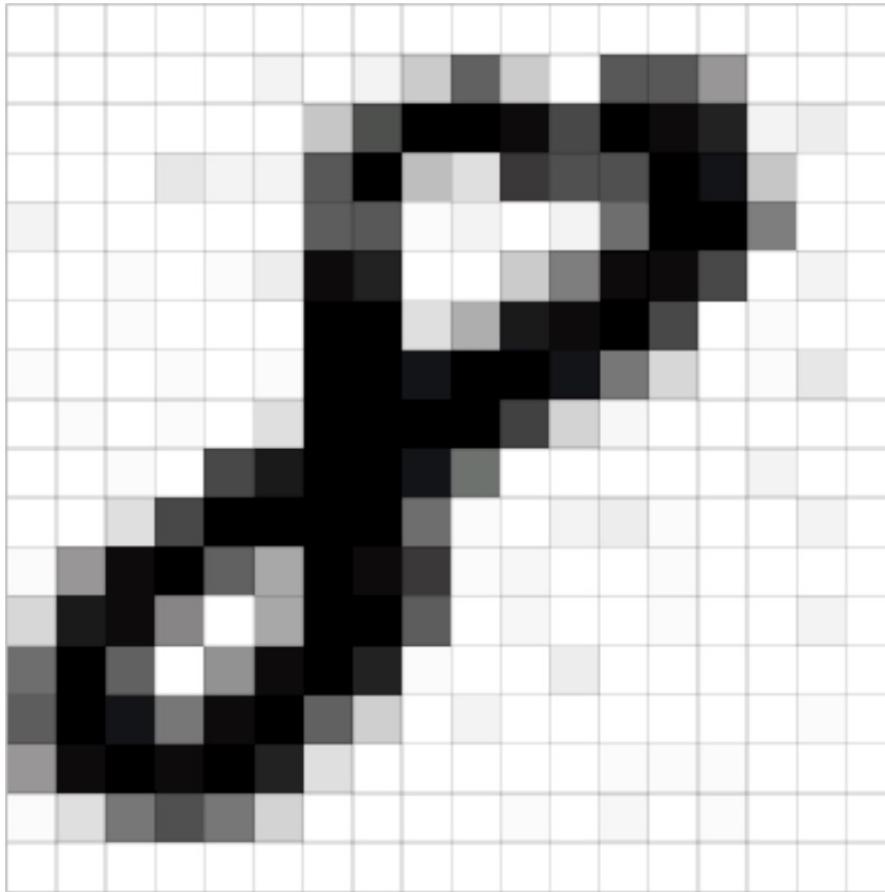
LeCun, Yann, Patrick Haffner, Léon Bottou, and Yoshua Bengio. "Object recognition with gradient-based learning." In *Shape, contour and grouping in computer vision*, pp. 319-345. Springer, Berlin, Heidelberg, 1999.

# What is convolutional neural network?



A grayscale image

# What is convolutional neural network?



A grayscale image

An image → a matrix

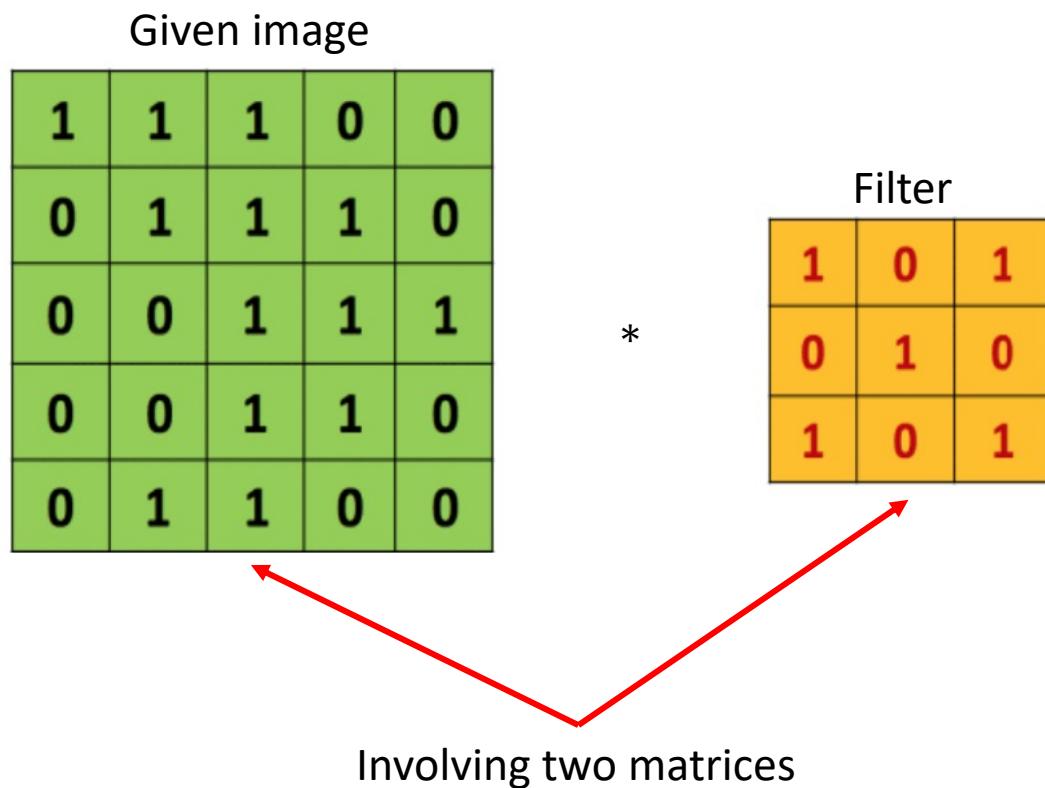
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

# Convolution for images (matrices)

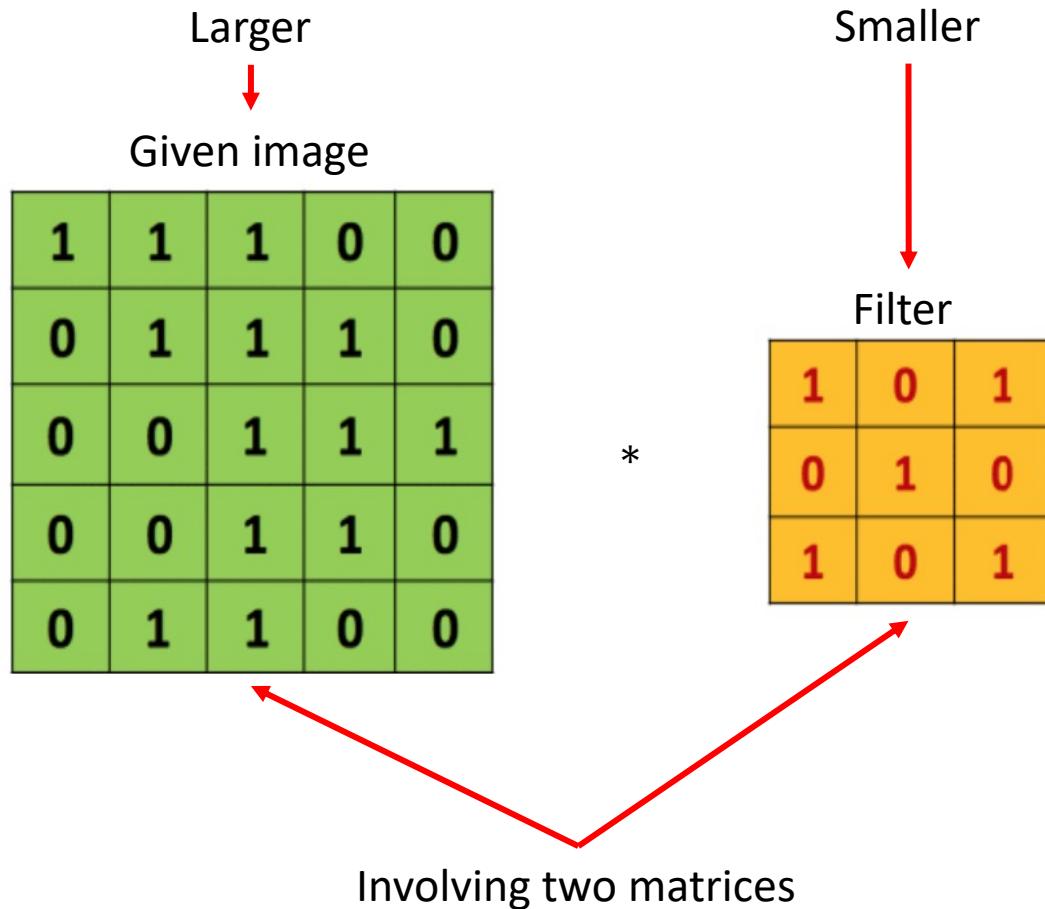
$$\begin{matrix} \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix}$$

Involving two matrices

# Convolution for images (matrices)



# Convolution for images (matrices)



# Convolution for images (matrices)

$$\begin{matrix} \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix}$$

Finding pairs

# Convolution for images (matrices)

$$\begin{matrix} & \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & \\ \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & \times & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix}$$

Finding pairs

# Convolution for images (matrices)

$$\begin{matrix} \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix}$$

Finding pairs

# Convolution for images (matrices)

$$\begin{matrix} & \begin{matrix} 1 & 1 & 1 & 0 & 0 \end{matrix} \\ \begin{matrix} 1 & 1 & 1 & 1 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & \begin{matrix} * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix} \end{matrix}$$

Finding pairs

# Convolution for images (matrices)

$$\begin{matrix} \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix}$$

Finding pairs

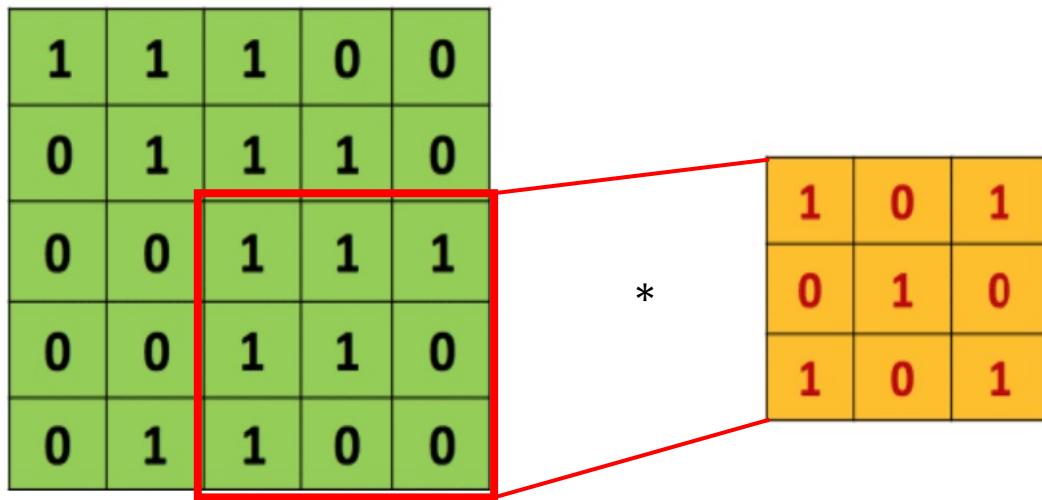
# Convolution for images (matrices)

$$\begin{matrix} \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix}$$

Finding pairs

Q: how many pairs we have?

# Convolution for images (matrices)



Finding pairs

Q: how many pairs we have?

$$(5-3+1) * (5-3+1)=9$$

# Convolution for images (matrices)

$$\begin{matrix} \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix}$$

Inner product of each pair

# Convolution for images (matrices)

$$\begin{matrix} \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix}$$

Inner product of each pair

Elementwise multiplication + summation

# Convolution for images (matrices)

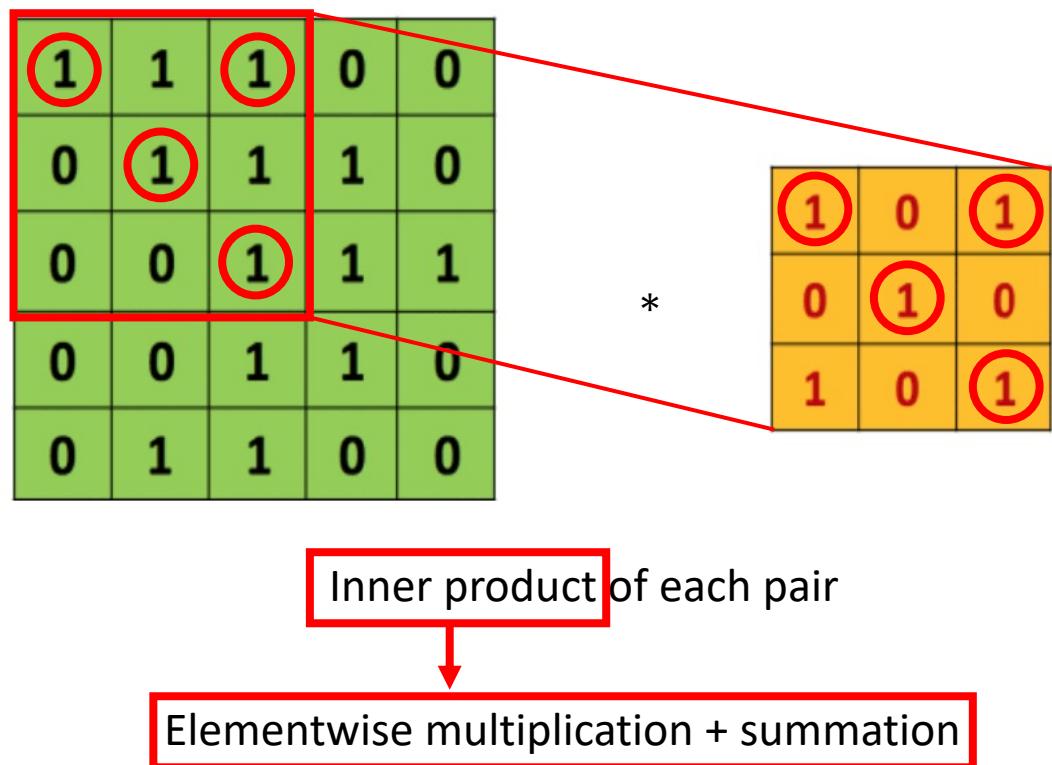
$$\begin{matrix} \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix}$$

Inner product of each pair

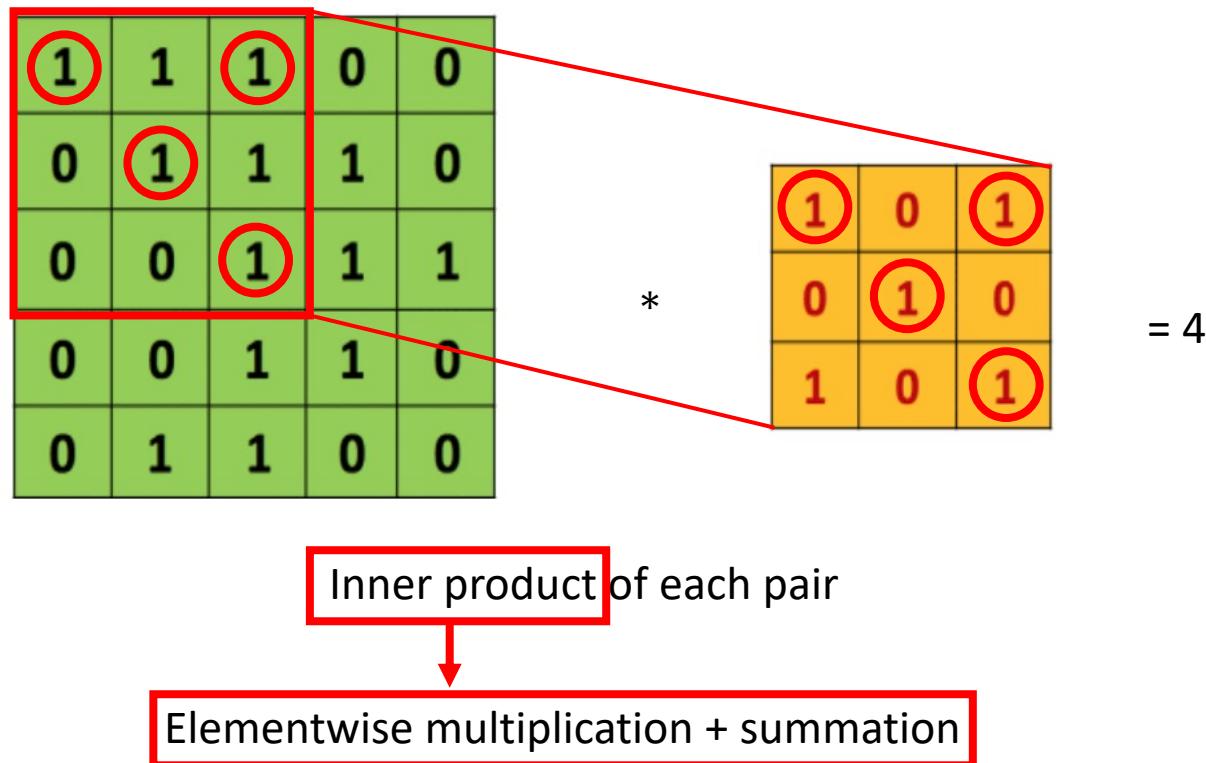
Elementwise multiplication + summation

Q: what is your result for the first pair?

# Convolution for images (matrices)



# Convolution for images (matrices)



# Convolution for images (matrices)

$$\begin{matrix} & \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & \\ \begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} & \times & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix}$$

Q: the second pair?

# Convolution for images (matrices)

$$\begin{matrix} & \begin{matrix} 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 0 & 1 & 1 & 0 \end{matrix} & * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \end{matrix}$$

The diagram illustrates a convolution operation between two matrices. The input matrix (left) is a 5x5 grid with values: [1, 1, 0, 0], [0, 1, 1, 0], [0, 0, 1, 1], [0, 0, 1, 1], [0, 1, 1, 0]. A 3x3 kernel (right) has values: [1, 0, 1], [0, 1, 0], [1, 0, 1]. Red circles highlight the elements 1, 1, and 1 in the top-left 3x3 subgrid of the input, which are being multiplied by the corresponding 1 in the top-left element of the kernel. Red lines connect these highlighted elements.

# Convolution for images (matrices)

$$\begin{matrix} & \begin{matrix} 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 0 & 1 & 1 & 0 \end{matrix} & * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} & = 3 \end{matrix}$$

The diagram illustrates a convolution operation between two matrices. The input matrix (left) is a 5x5 grid with values: [1, 1, 0, 0], [0, 1, 1, 0], [0, 0, 1, 1], [0, 0, 1, 1], [0, 1, 1, 0]. A red box highlights a 3x3 submatrix in the top-left corner: [1, 1, 0], [0, 1, 1], [0, 0, 1]. The kernel matrix (right) is a 3x3 grid with values: [1, 0, 1], [0, 1, 0], [1, 0, 1]. Red circles highlight the elements 1, 1, and 1 in the kernel, which correspond to the highlighted elements in the input submatrix. The result of the convolution step is 3.

# Convolution for images (matrices)

$$\begin{matrix} & \begin{matrix} 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 0 & 1 & 1 & 0 \end{matrix} & * & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} & = 3 \end{matrix}$$

The diagram illustrates a convolution operation between two matrices. The input matrix (left) is a 5x5 grid with values: [1, 1, 0, 0], [0, 1, 1, 0], [0, 0, 1, 1], [0, 0, 1, 1], [0, 1, 1, 0]. A 3x3 kernel (right) has values: [1, 0, 1], [0, 1, 0], [1, 0, 1]. The result of the convolution is 3. Red circles highlight the elements being multiplied in the first step of the kernel's slide across the input.

We can repeat for each pair

# Convolution for images (matrices)

$$\begin{matrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{matrix} * \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix} \rightarrow \begin{matrix} 4 & 3 & 4 \\ 2 & 4 & 3 \\ 2 & 3 & 4 \end{matrix}$$

# Convolution for images (matrices)

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

\*

1	0	1
0	1	0
1	0	1



4	3	4
2	4	3
2	3	4

Place each element  
according to their positions

# Convolution for images (matrices)

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

\*

1	0	1
0	1	0
1	0	1



Row: 1  
Column: 1

4	3	4
2	4	3
2	3	4

Place each element  
according to their positions

# Convolution for images (matrices)

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

\*

1	0	1
0	1	0
1	0	1



4	3	4
2	4	3
2	3	4

Row: 1  
Column: 2

Place each element  
according to their positions

# Convolution for images (matrices)

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

\*

1	0	1
0	1	0
1	0	1



Row: 3  
Column: 3

4	3	4
2	4	3
2	3	4

Place each element  
according to their positions

# Convolution for images (matrices)

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

n=5

\*

1	0	1
0	1	0
1	0	1

m=3

→

4	3	4
2	4	3
2	3	4

Q: dimension?

# Convolution for images (matrices)

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

n=5

\*

1	0	1
0	1	0
1	0	1

m=3

→

4	3	4
2	4	3
2	3	4

n-m+1=3

# Convolution for images (matrices)

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

n=5

One matrix

\*

1	0	1
0	1	0
1	0	1

m=3

One matrix

→

4	3	4
2	4	3
2	3	4

n-m+1=3

One matrix

# Convolution for images (matrices)

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

n=5

One matrix

1	0	1
0	1	0
1	0	1

\*

m=3

One matrix



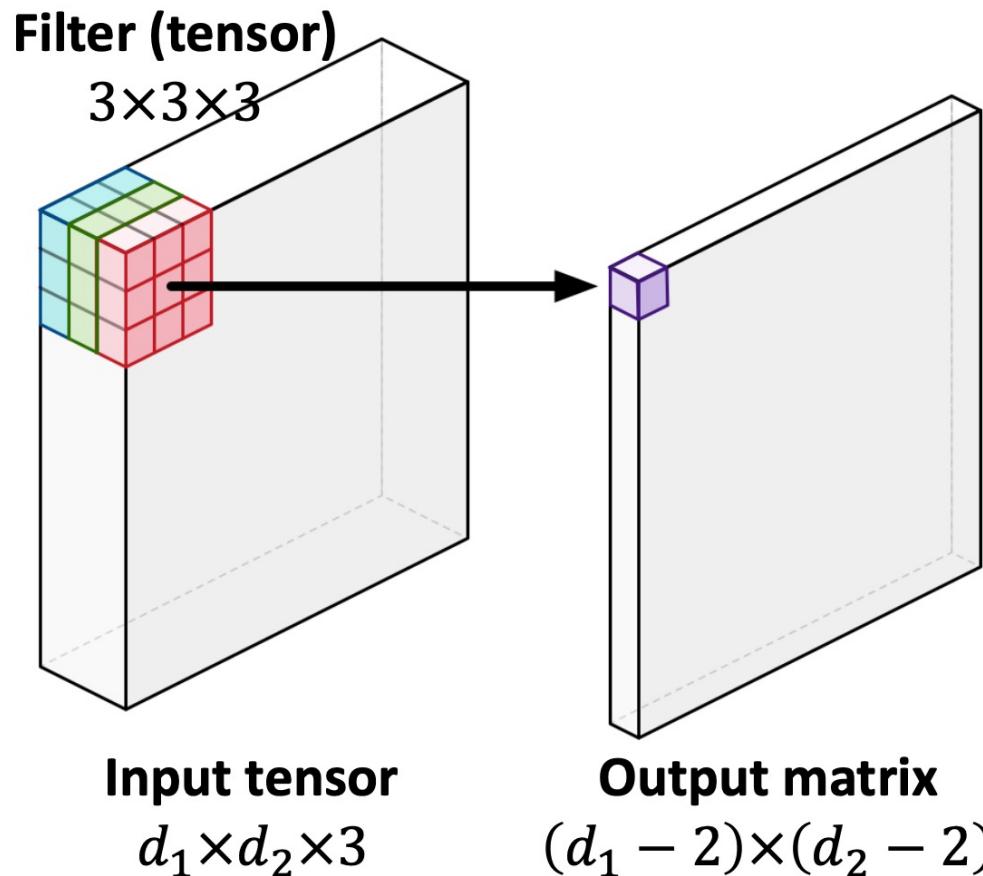
4	3	4
2	4	3
2	3	4

n-m+1=3

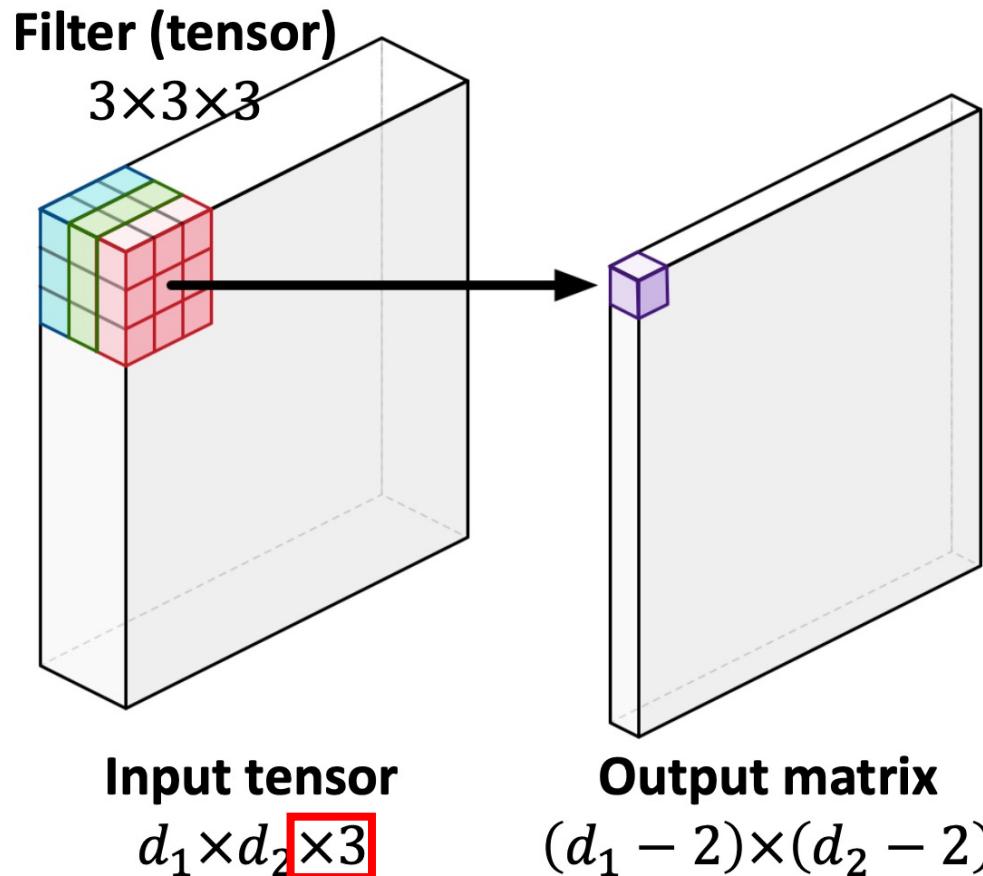
One matrix

One input matrix \* one filter → one feature matrix

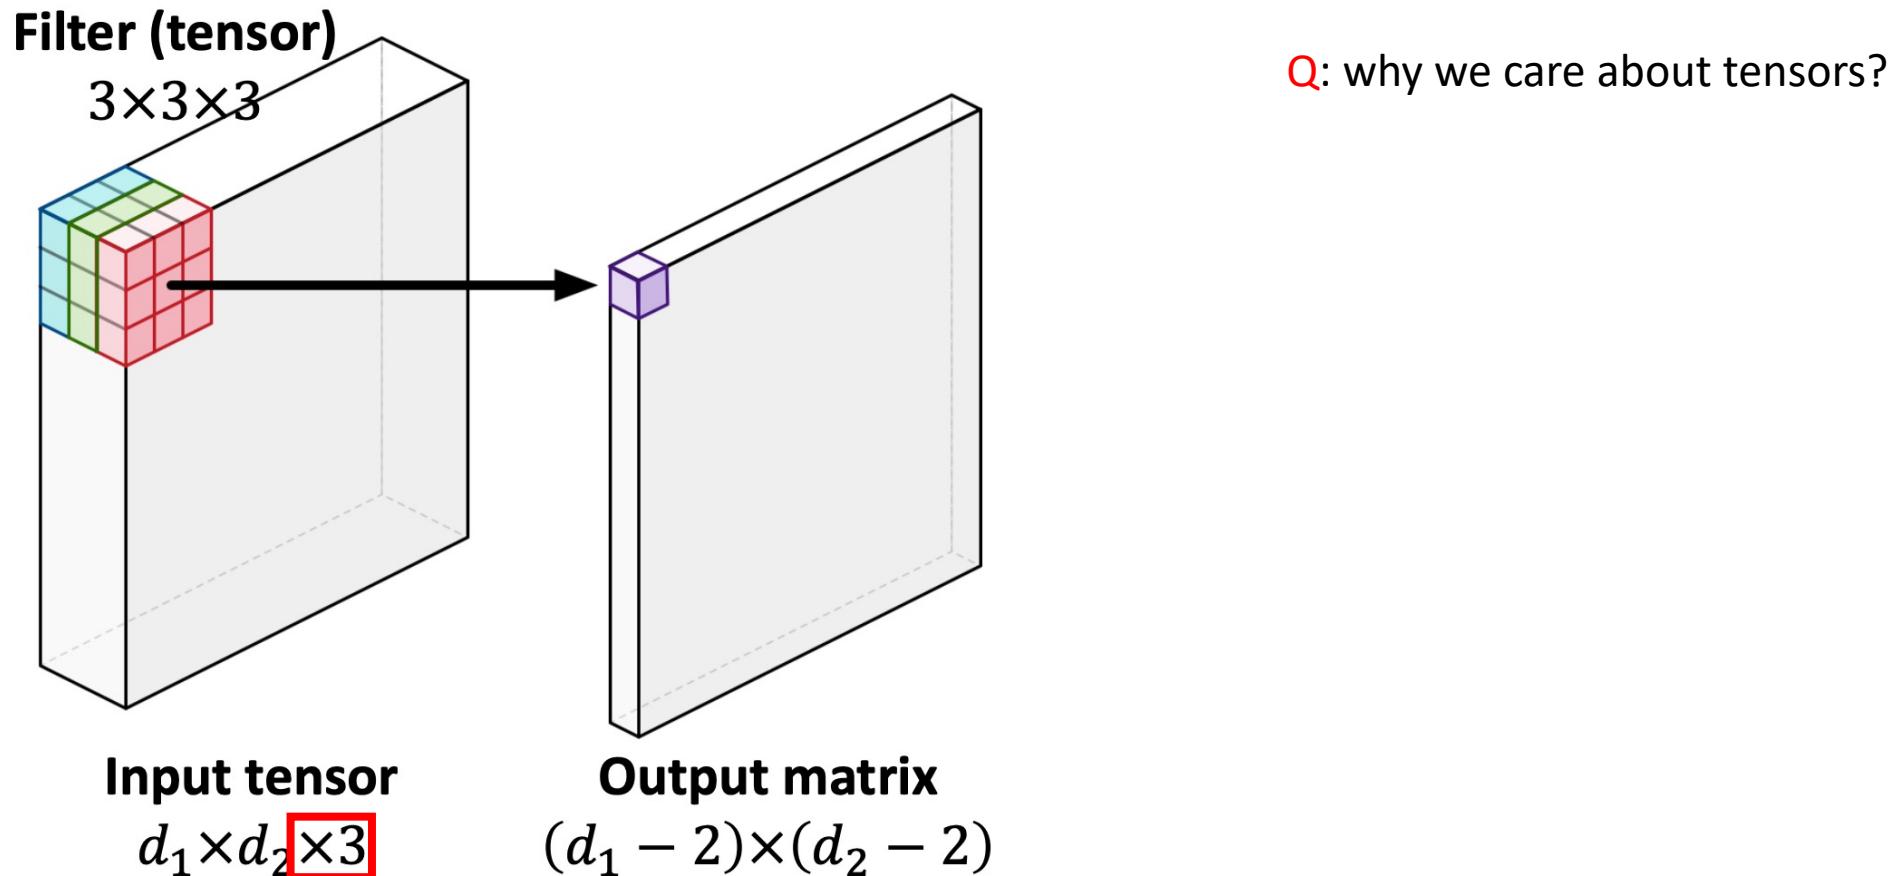
# Convolution for images (**tensors**)



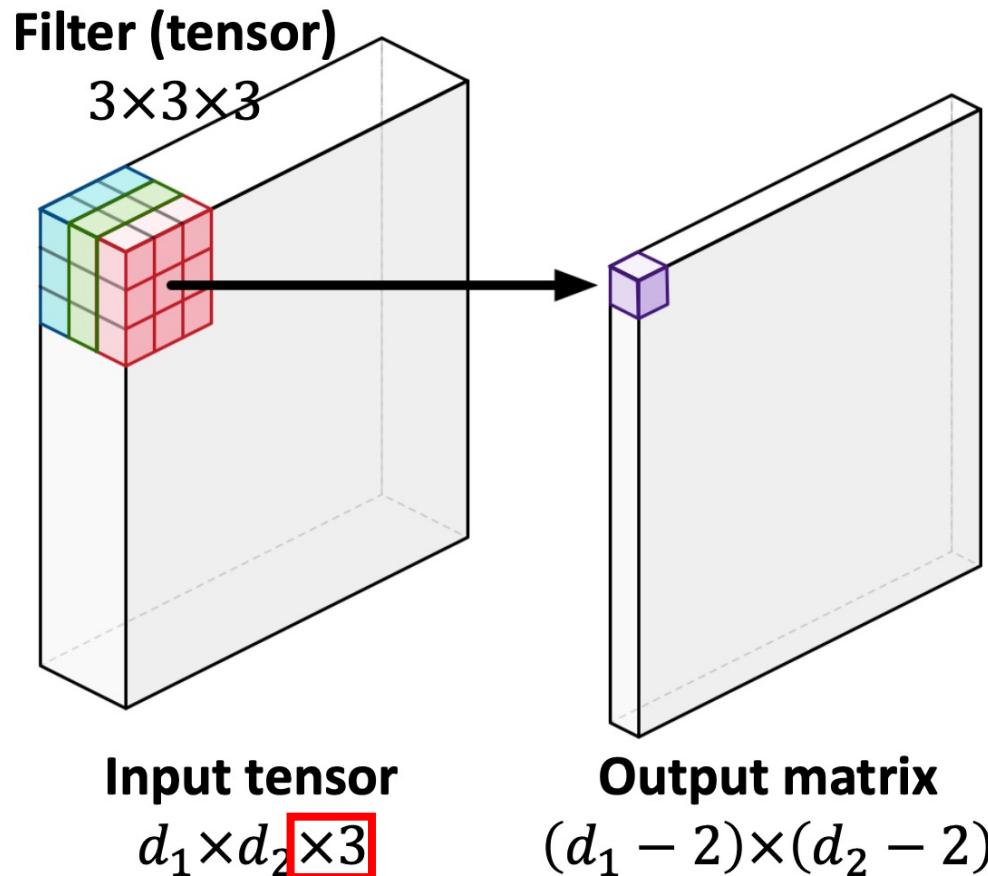
# Convolution for images (**tensors**)



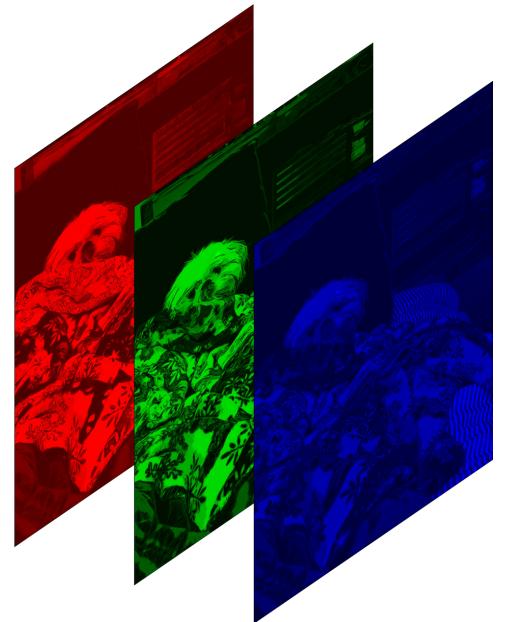
# Convolution for images (tensors)



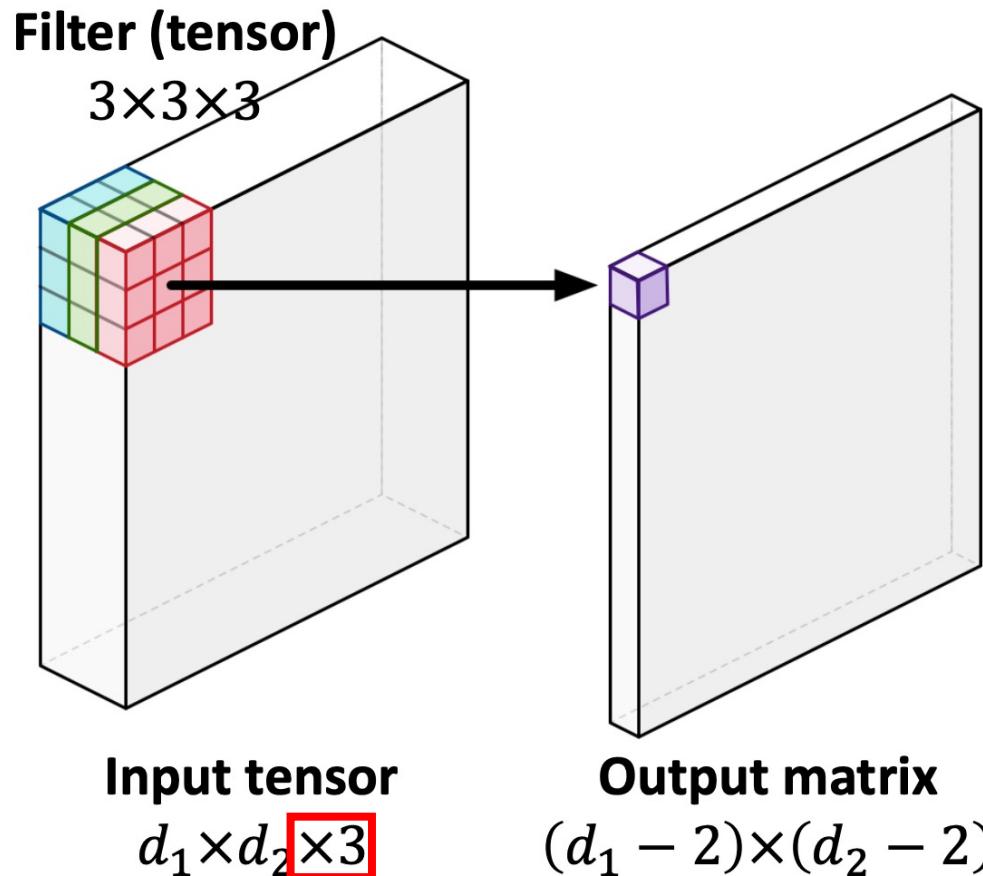
# Convolution for images (tensors)



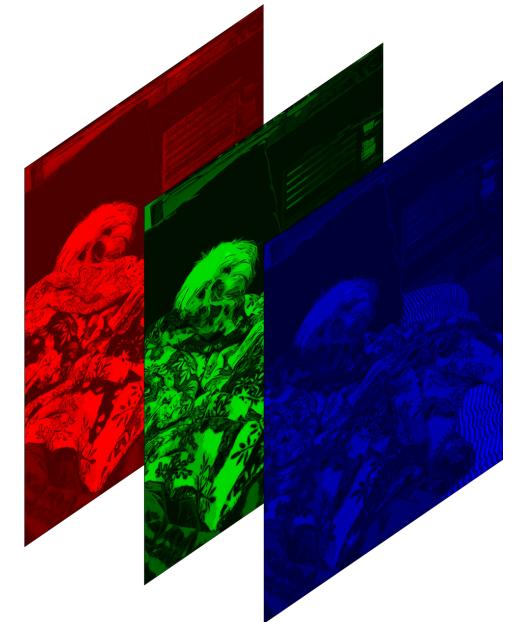
Q: why we care about tensors?



# Convolution for images (tensors)

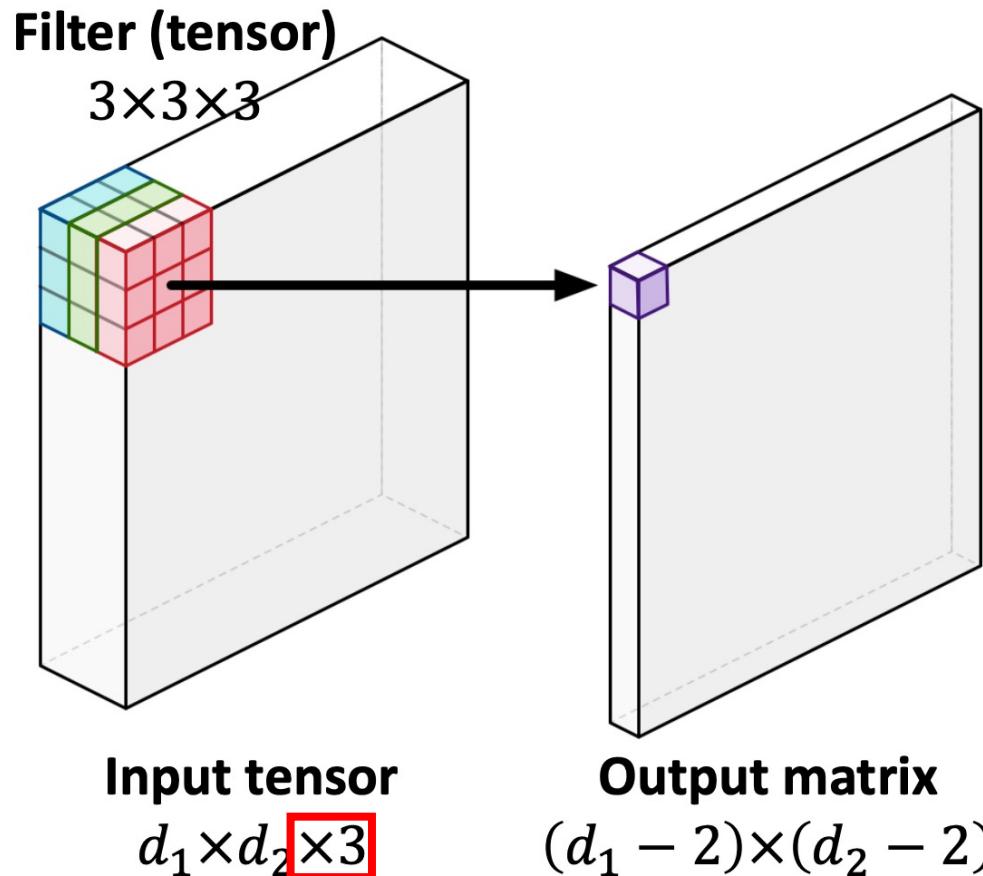


Q: why we care about tensors?

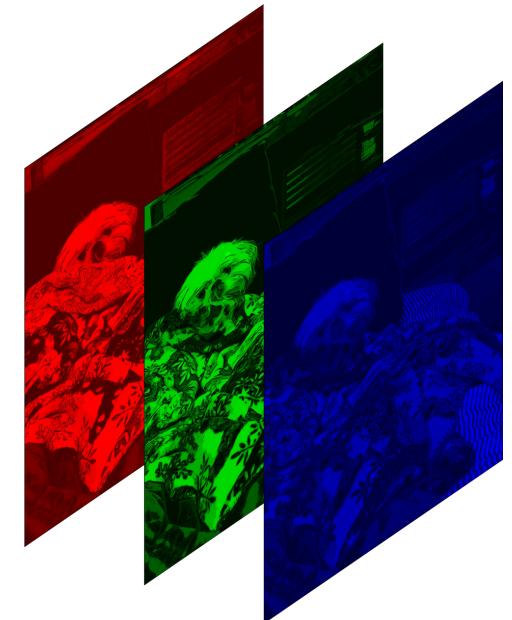


Reason 1:  
RGB channels are more common

# Convolution for images (tensors)

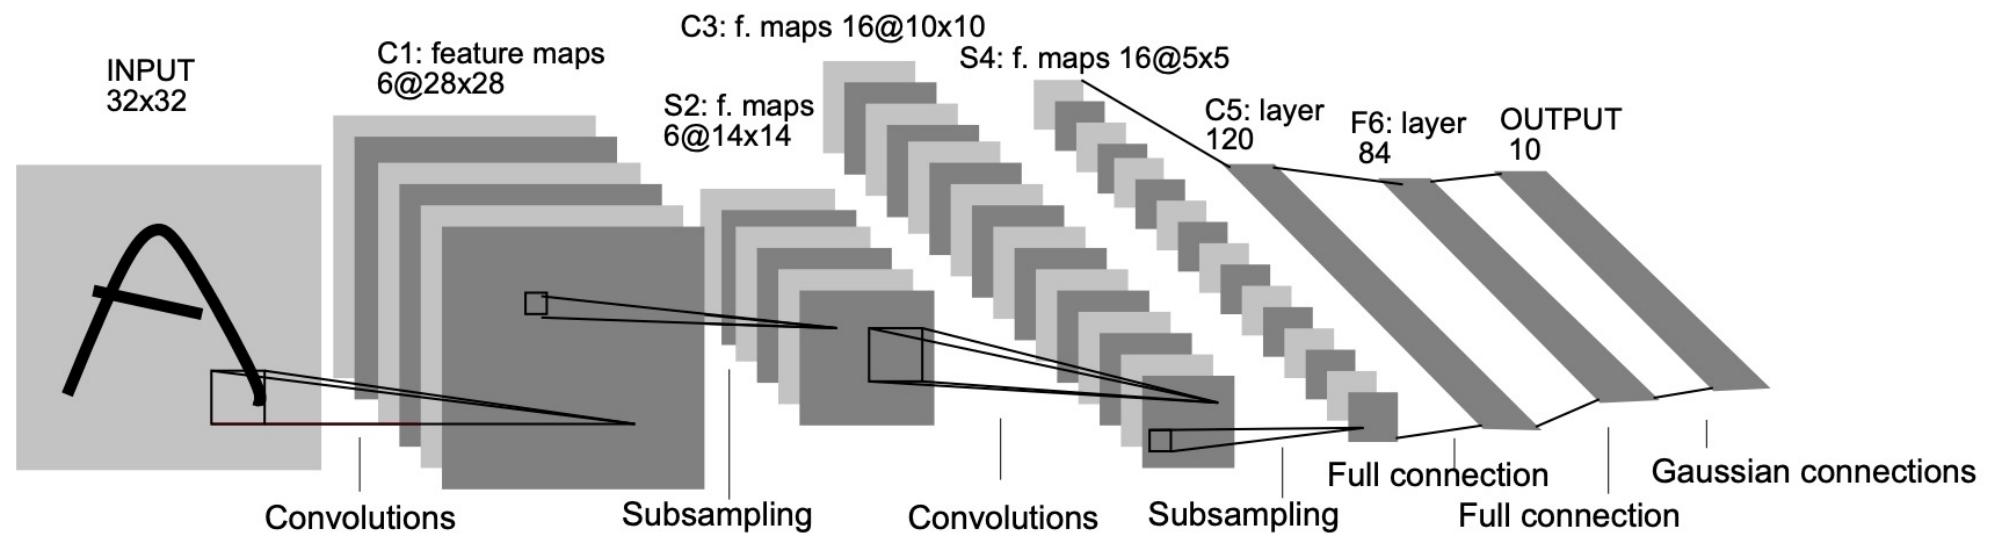


Q: why we care about tensors?



Reason 1:  
RGB channels are more common  
Each channel → a matrix

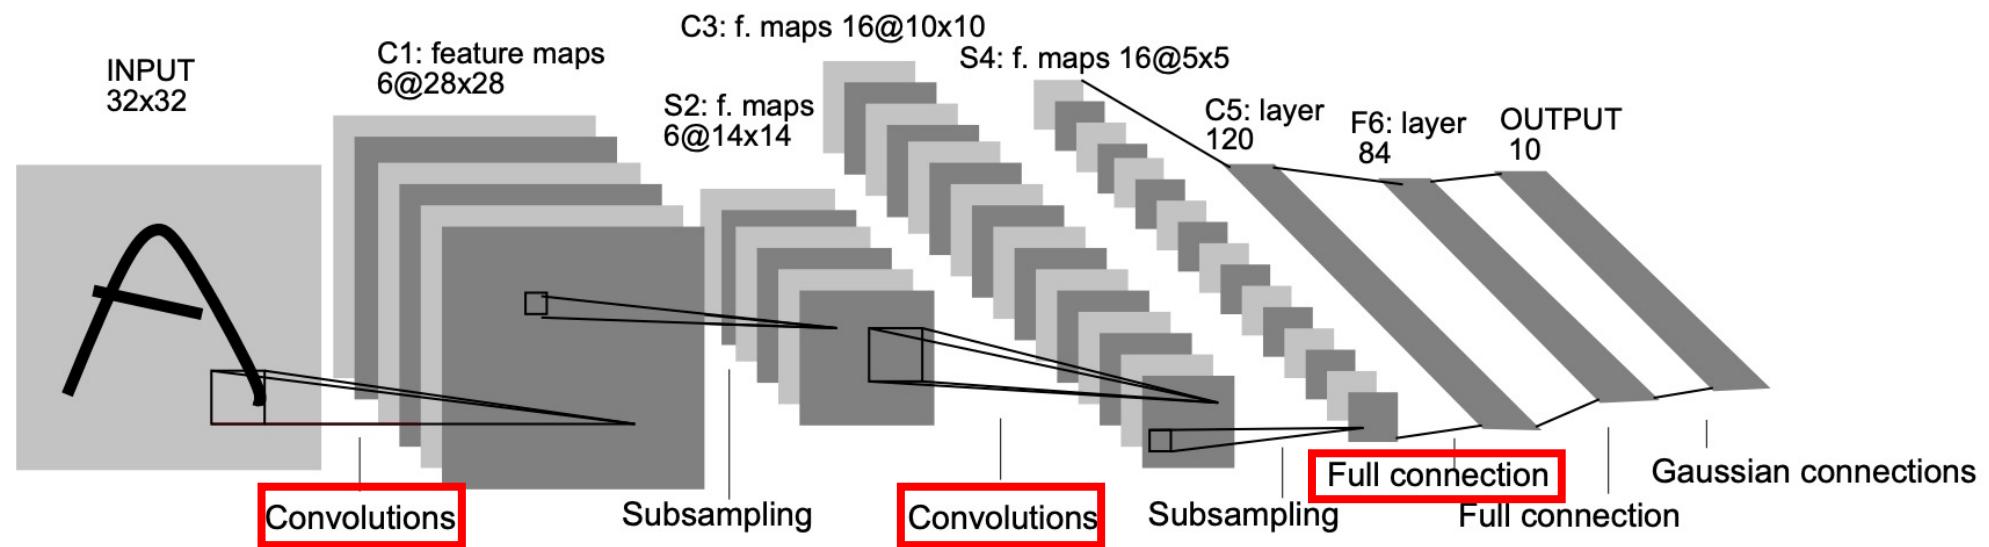
# LeNet-5 in 1999



**Fig. 1.** Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeCun, Yann, Patrick Haffner, Léon Bottou, and Yoshua Bengio. "Object recognition with gradient-based learning." In *Shape, contour and grouping in computer vision*, pp. 319-345. Springer, Berlin, Heidelberg, 1999. <sup>45</sup>

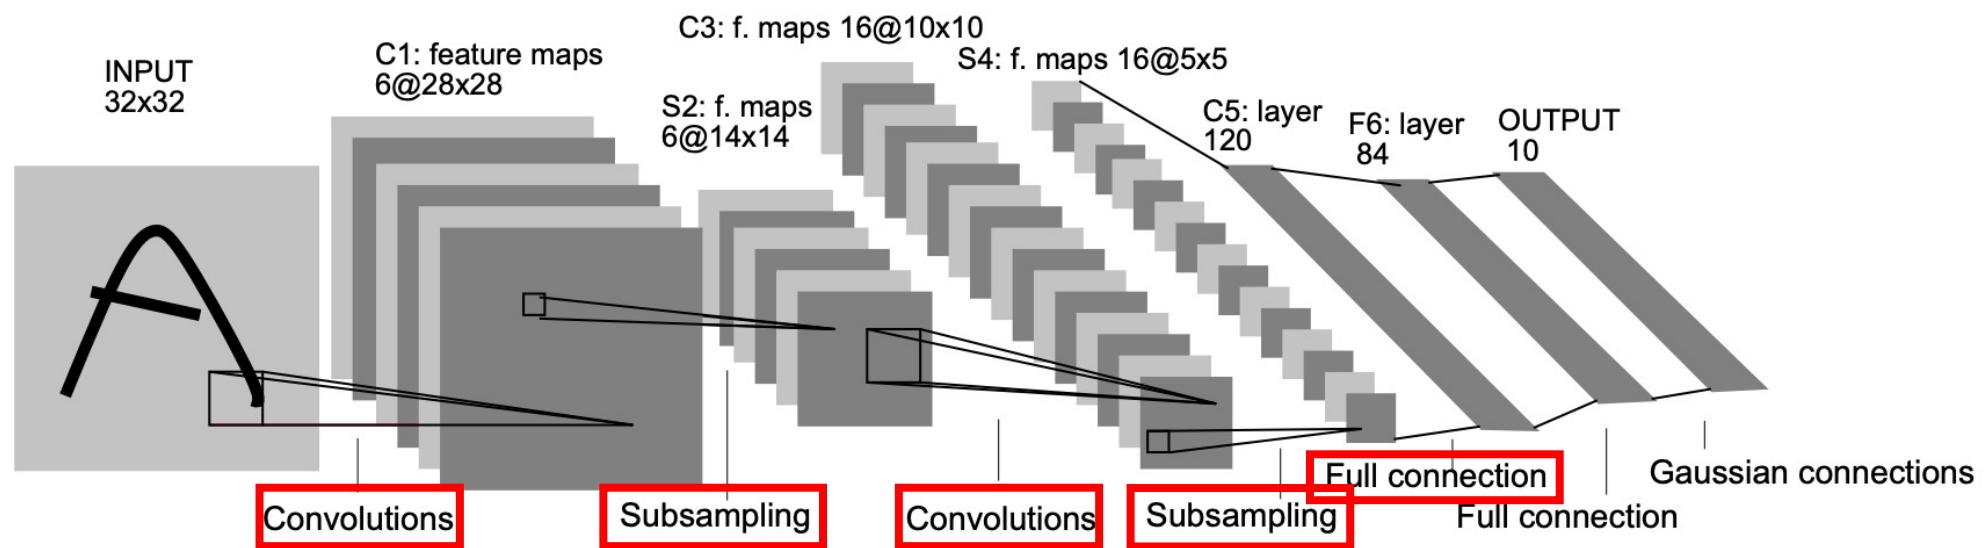
# LeNet-5 in 1999



**Fig. 1.** Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeCun, Yann, Patrick Haffner, Léon Bottou, and Yoshua Bengio. "Object recognition with gradient-based learning." In *Shape, contour and grouping in computer vision*, pp. 319-345. Springer, Berlin, Heidelberg, 1999. <sup>46</sup>

# LeNet-5 in 1999

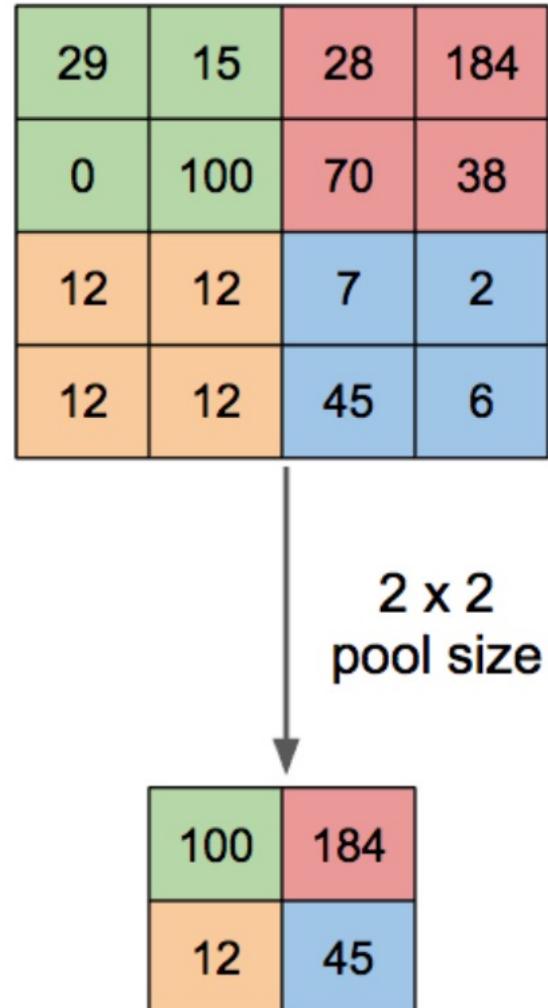


**Fig. 1.** Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeCun, Yann, Patrick Haffner, Léon Bottou, and Yoshua Bengio. "Object recognition with gradient-based learning." In *Shape, contour and grouping in computer vision*, pp. 319-345. Springer, Berlin, Heidelberg, 1999. <sup>47</sup>

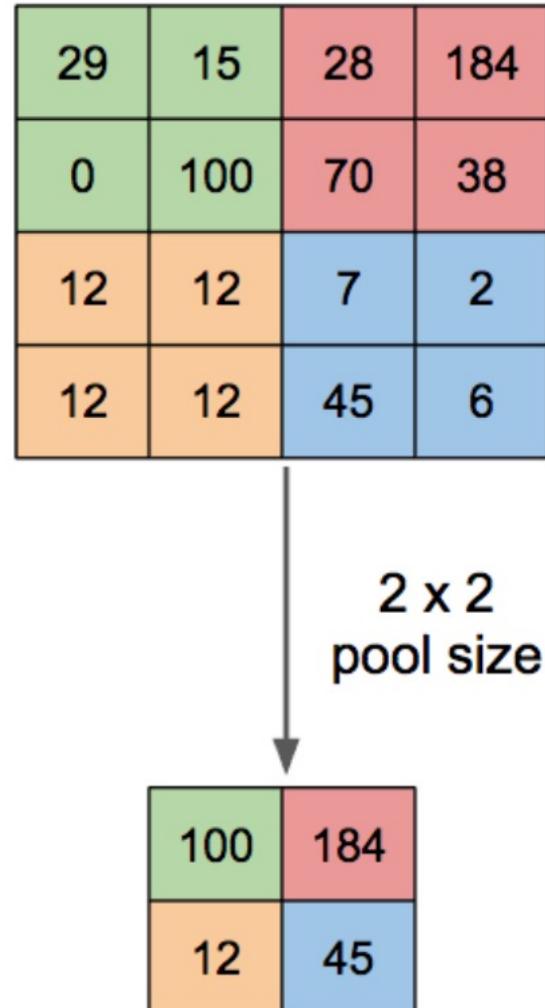
# Subsampling operations

- Max pooling



# Subsampling operations

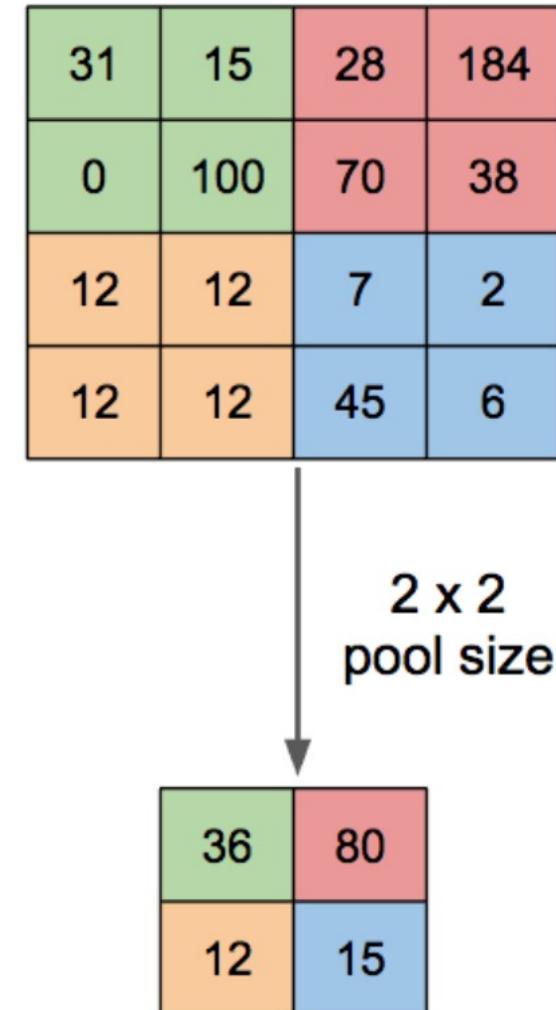
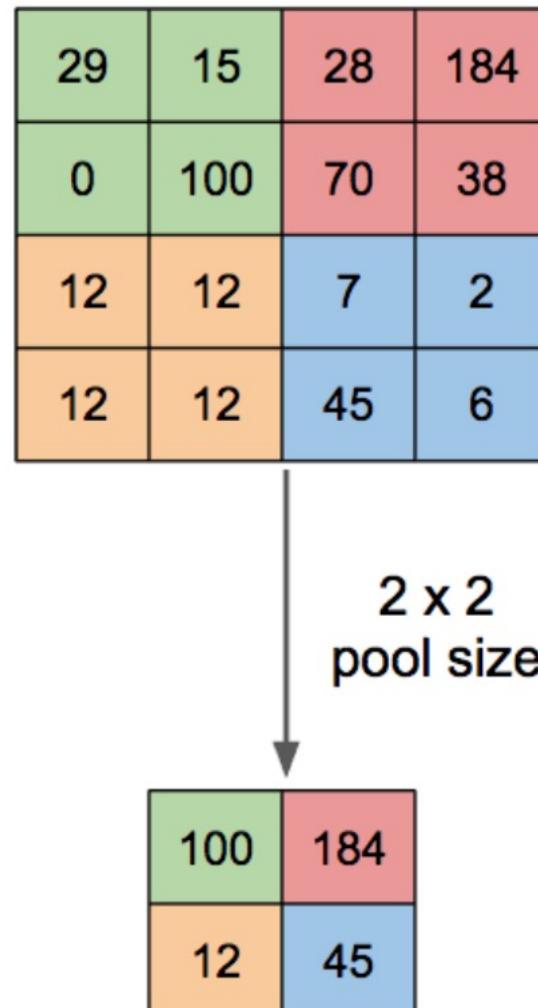
- Max pooling



Q: what does max Pooling really do?

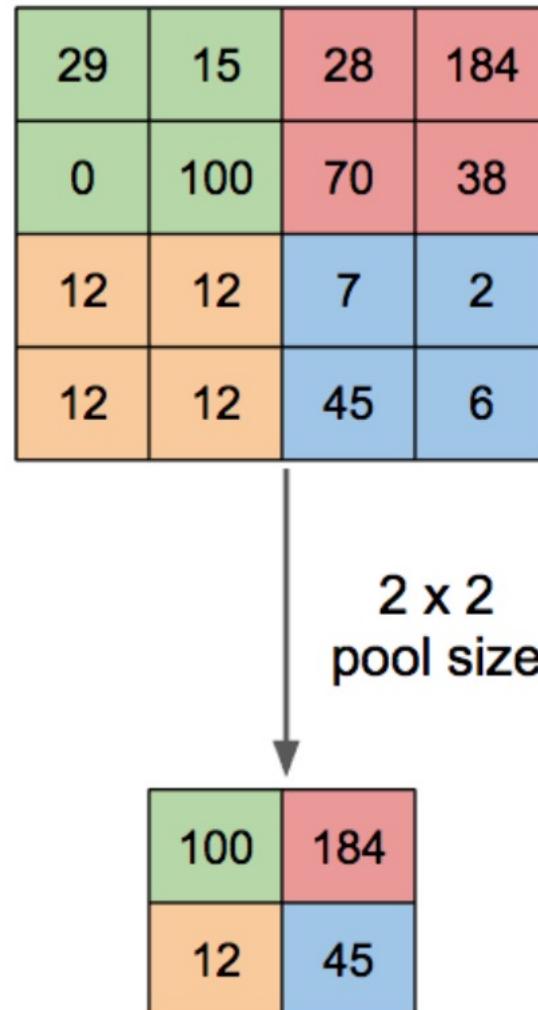
# Subsampling operations

- Max pooling
- Average pooling

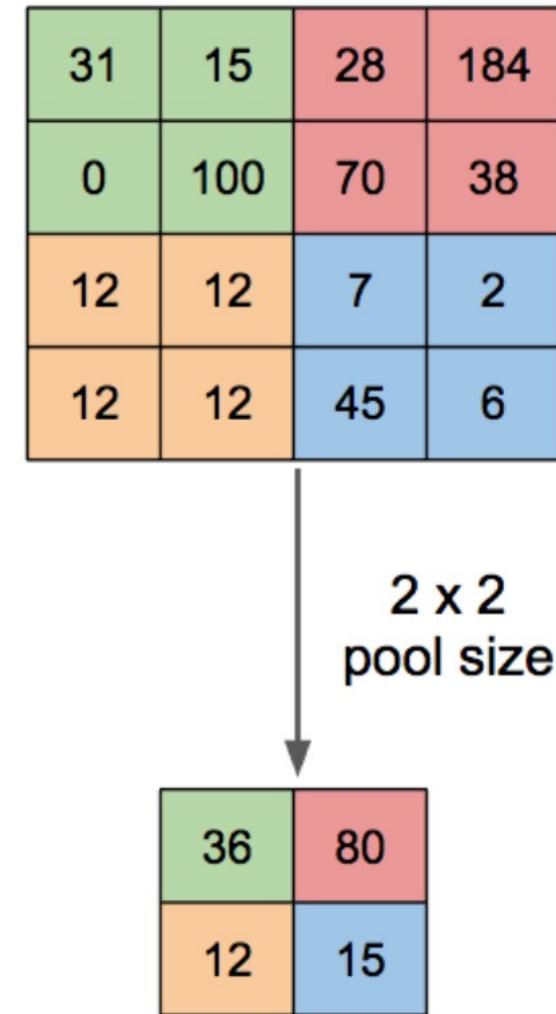


# Subsampling operations

- Max pooling
- Average pooling

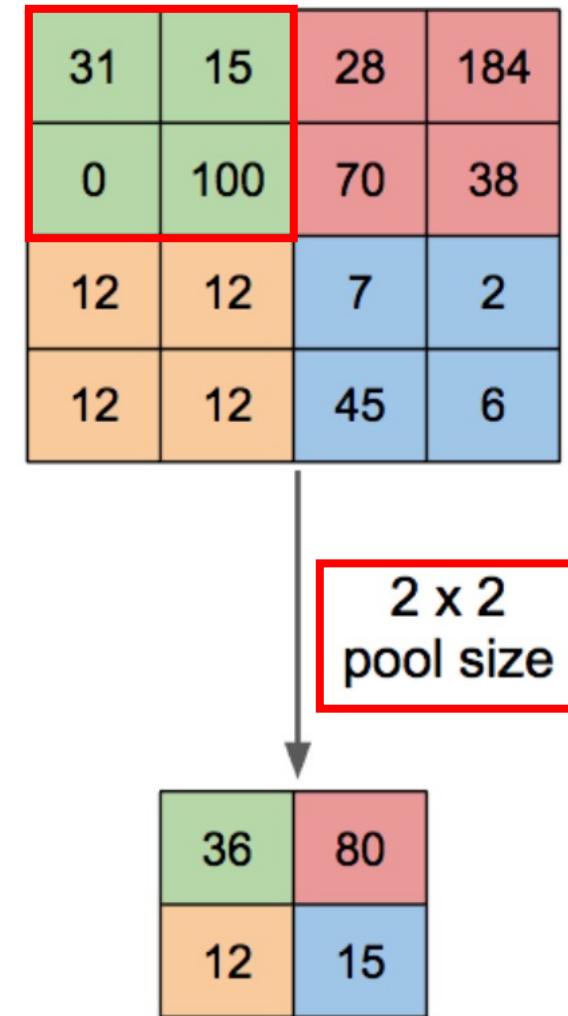
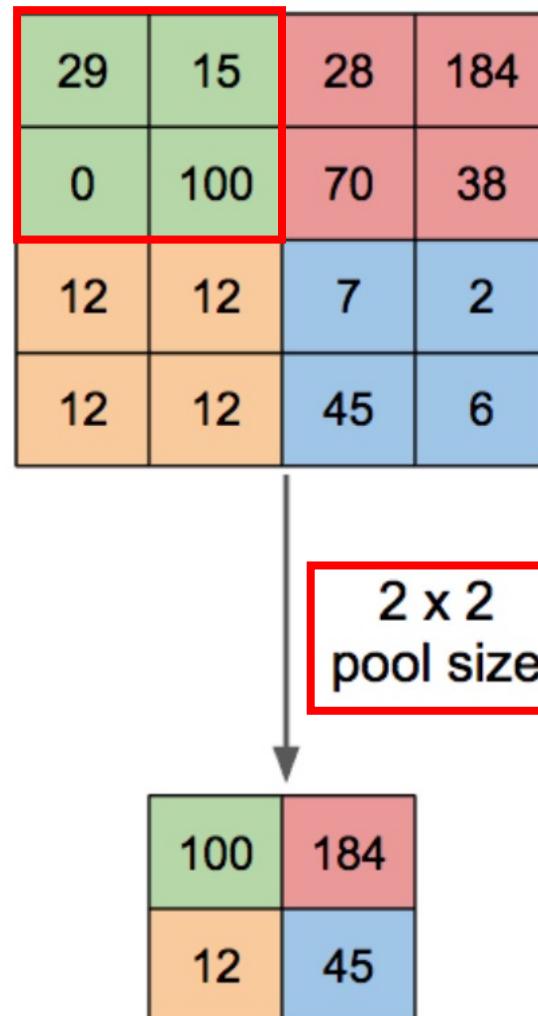


Q: what does average Pooling really do?



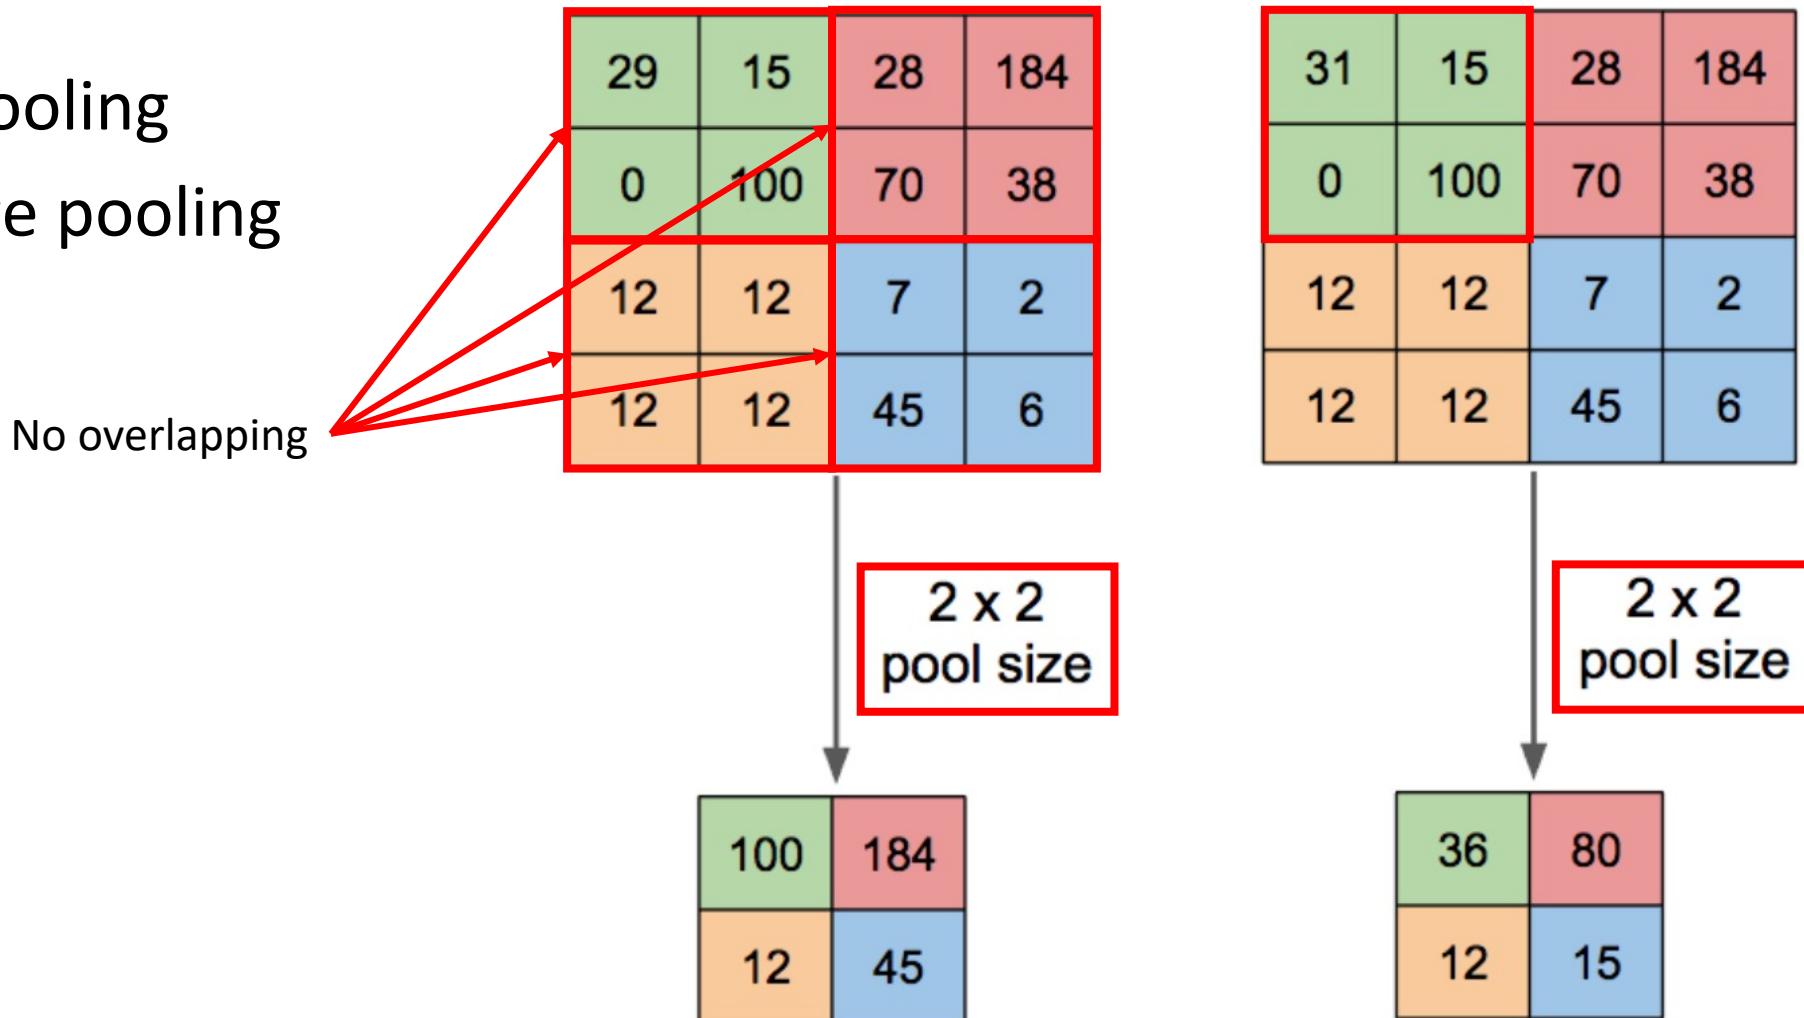
# Subsampling operations

- Max pooling
- Average pooling



# Subsampling operations

- Max pooling
- Average pooling



# Subsampling operations

- Max pooling
- Average pooling

No overlapping  
(stride=2\*2)

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2  
pool size

100	184
12	45

31	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2  
pool size

36	80
12	15

# Subsampling operations

- Max pooling
- Average pooling

No overlapping  
(stride=2\*2)

Row stride = 2  
Column stride = 2

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2  
pool size

100	184
12	45

31	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2  
pool size

36	80
12	15

# Subsampling operations

- Max pooling
- Average pooling

No overlapping  
(stride=2\*2)

Row stride = 2  
Column stride = 2

Q: Why pooling?  
Connection to subsampling?

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2  
pool size

100	184
12	45

31	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2  
pool size

36	80
12	15

# Subsampling operations

- Max pooling
- Average pooling

No overlapping  
(stride=2\*2)

Row stride = 2  
Column stride = 2

Q: Why pooling?  
Connection to subsampling?

$4 \times 4 \rightarrow 2 \times 2$

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2  
pool size

100	184
12	45

31	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2  
pool size

36	80
12	15

# Subsampling operations

- Max pooling
- Average pooling

No overlapping  
(stride=2\*2)

Row stride = 2  
Column stride = 2

Q: Why pooling?  
Connection to subsampling?

$4 \times 4 \rightarrow 2 \times 2$

Dimension reduced

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2  
pool size

100	184
12	45

31	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2  
pool size

36	80
12	15

# Subsampling operations

- Max pooling
- Average pooling

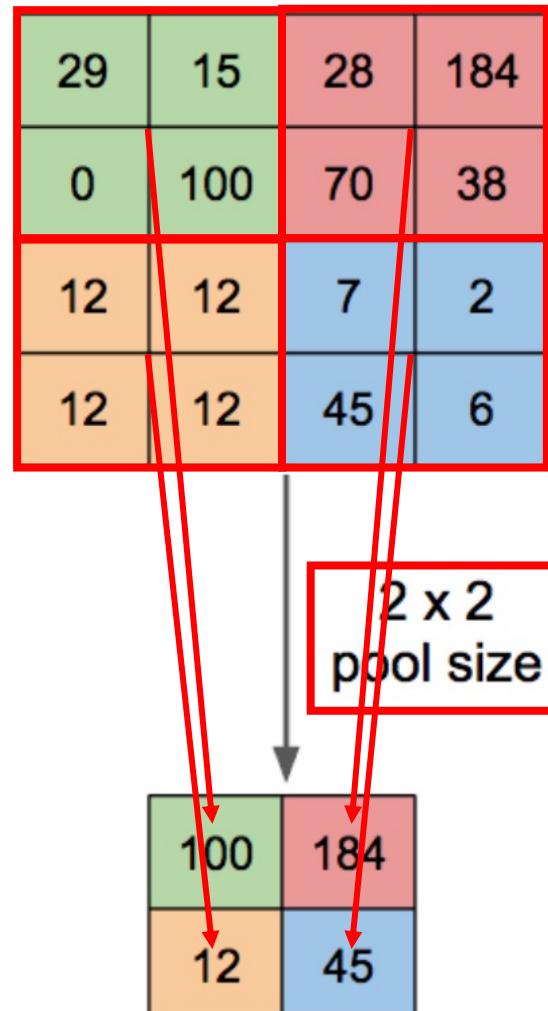
No overlapping  
(stride=2\*2)

Row stride = 2  
Column stride = 2

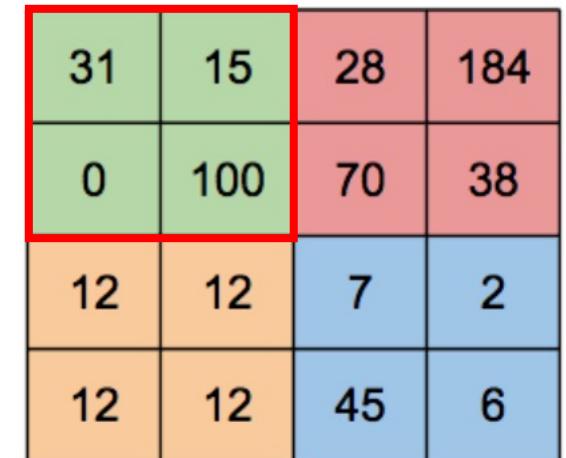
Q: Why pooling?  
Connection to subsampling?

$4 \times 4 \rightarrow 2 \times 2$

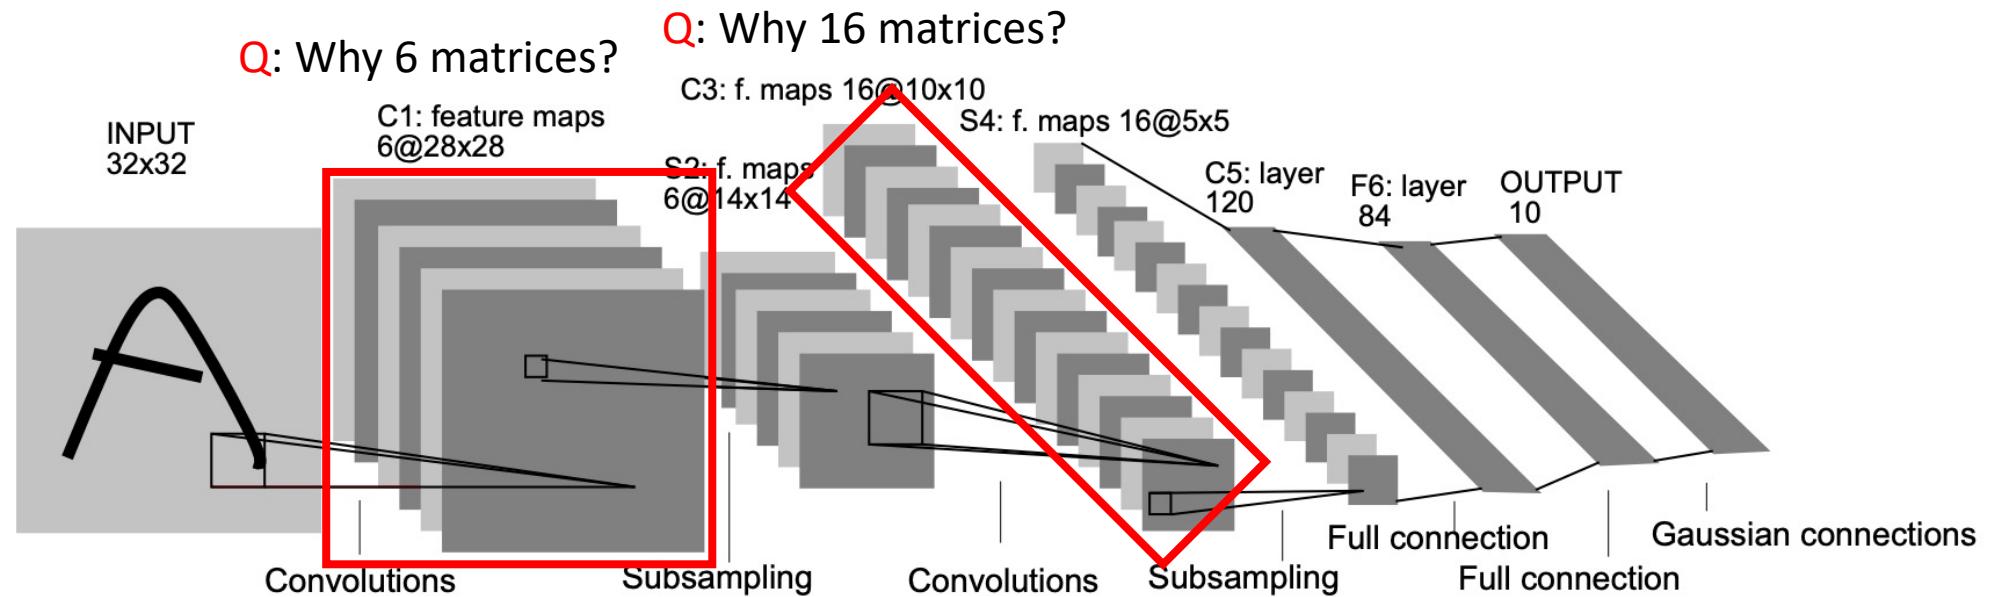
Dimension reduced



Use one to represent all



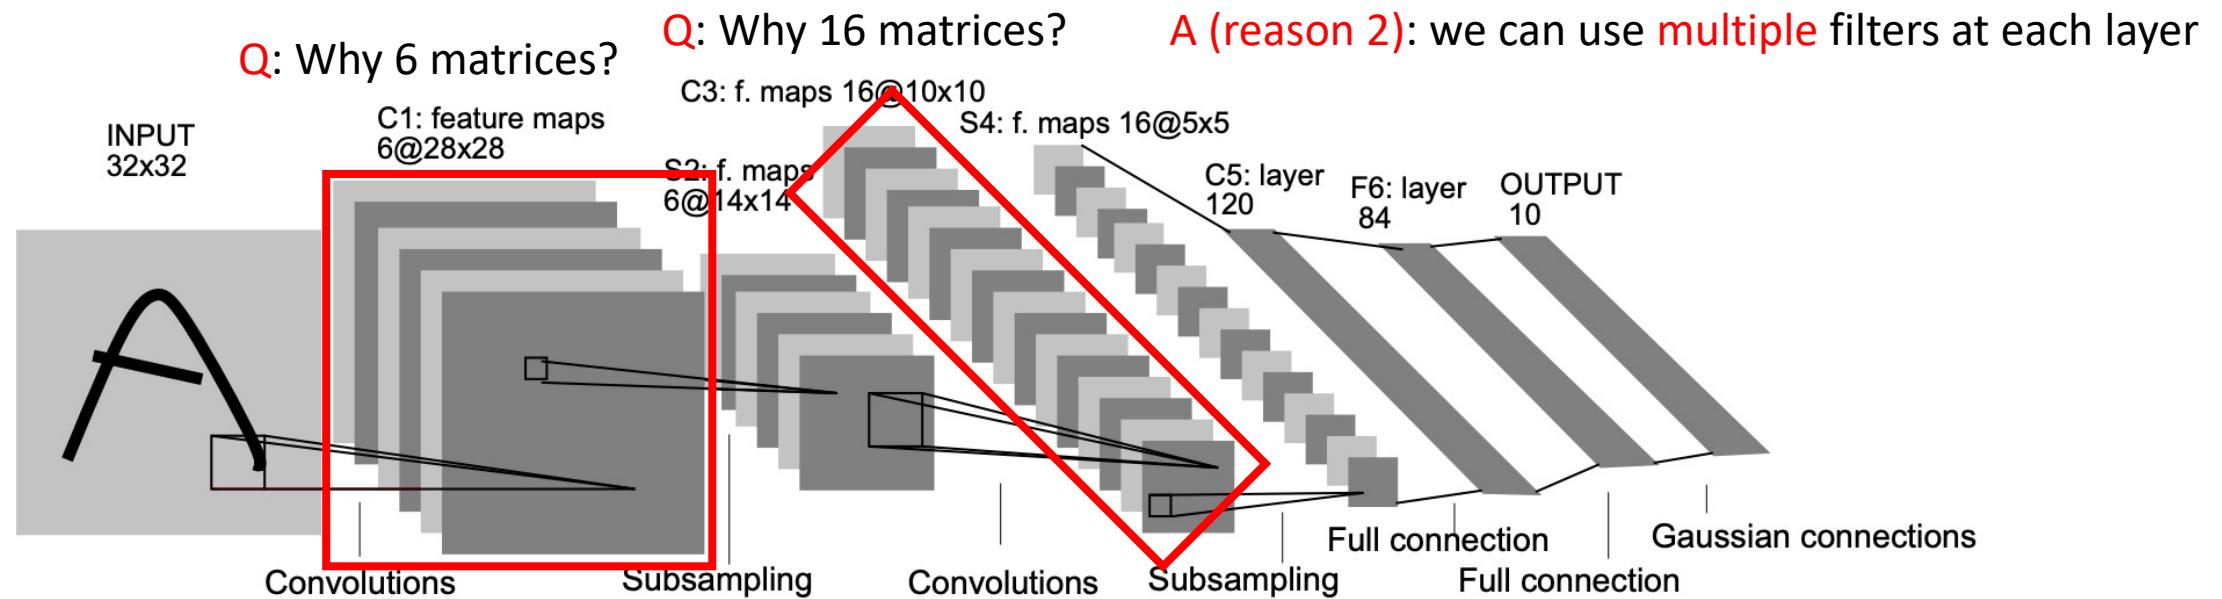
# LeNet-5 in 1999



**Fig. 1.** Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeCun, Yann, Patrick Haffner, Léon Bottou, and Yoshua Bengio. "Object recognition with gradient-based learning." In *Shape, contour and grouping in computer vision*, pp. 319-345. Springer, Berlin, Heidelberg, 1999. <sup>60</sup>

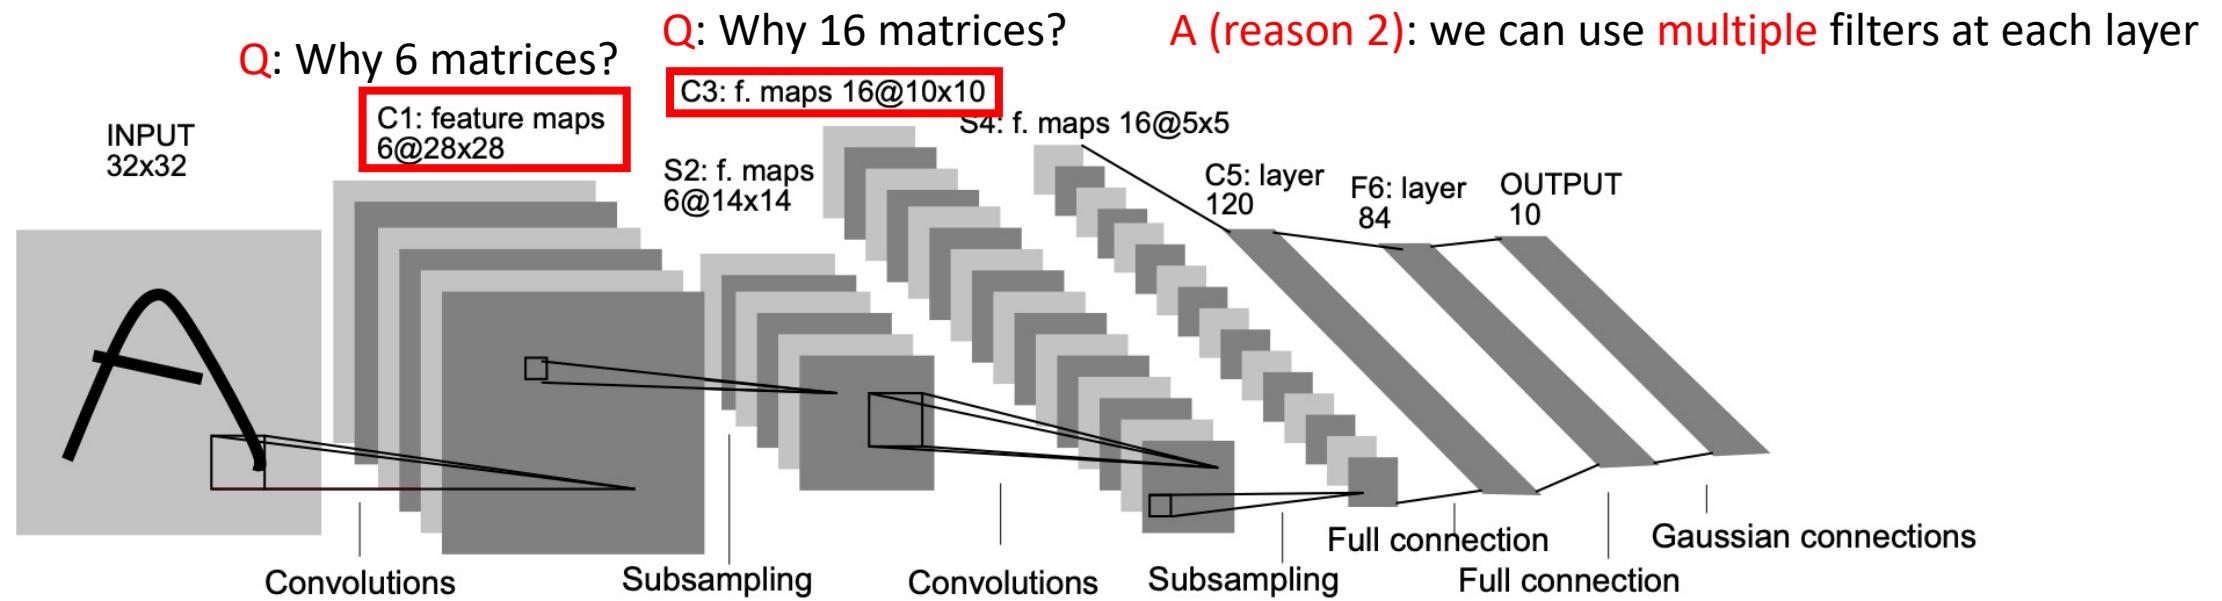
# LeNet-5 in 1999



**Fig. 1.** Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeCun, Yann, Patrick Haffner, Léon Bottou, and Yoshua Bengio. "Object recognition with gradient-based learning." In *Shape, contour and grouping in computer vision*, pp. 319-345. Springer, Berlin, Heidelberg, 1999. <sup>61</sup>

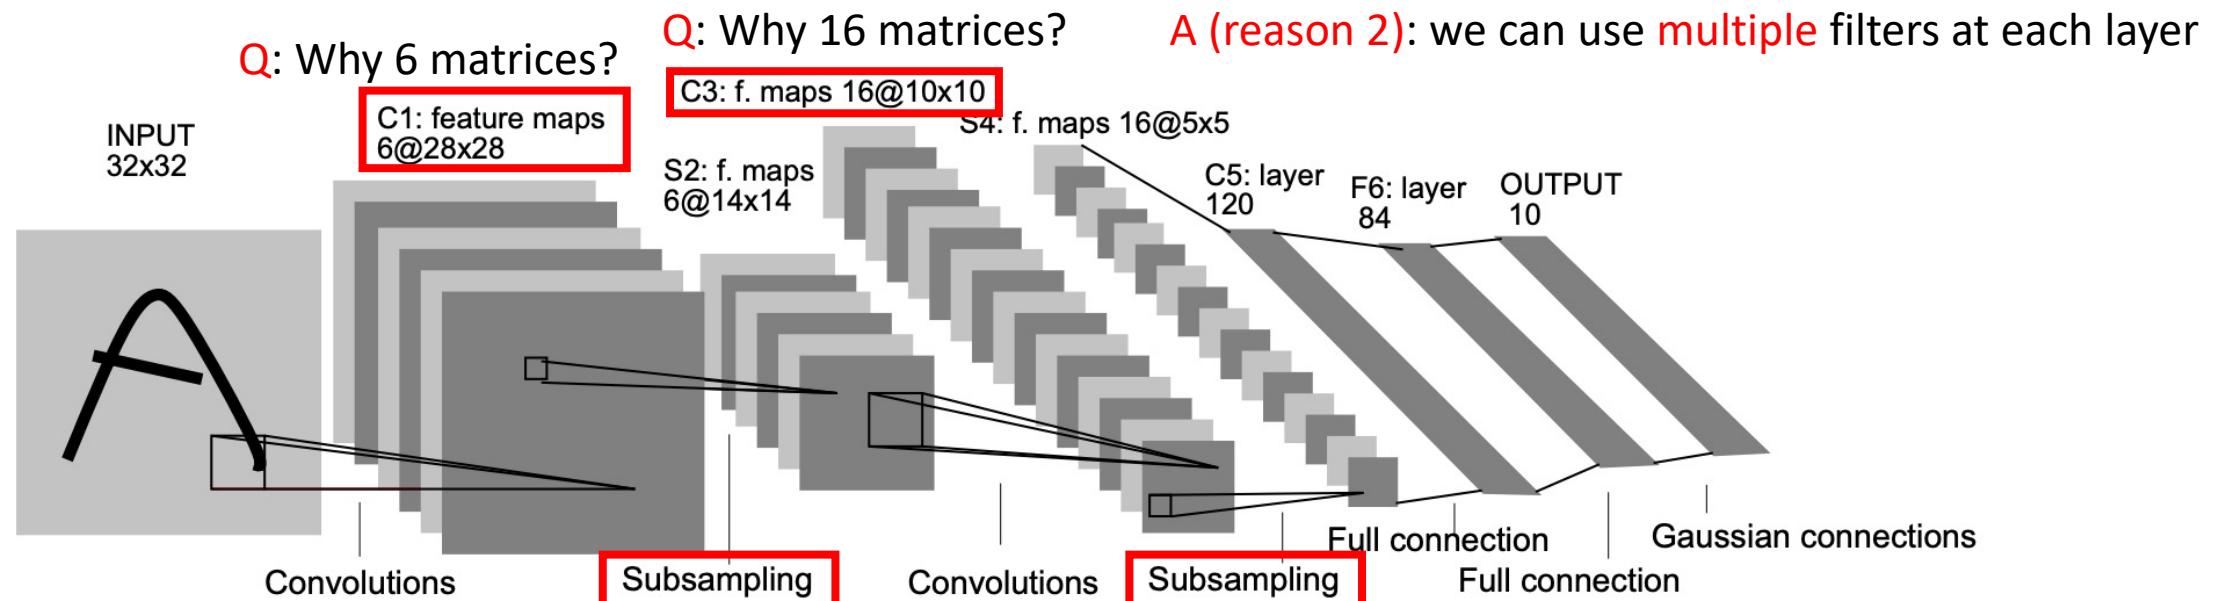
# LeNet-5 in 1999



**Fig. 1.** Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeCun, Yann, Patrick Haffner, Léon Bottou, and Yoshua Bengio. "Object recognition with gradient-based learning." In *Shape, contour and grouping in computer vision*, pp. 319-345. Springer, Berlin, Heidelberg, 1999. <sup>62</sup>

# LeNet-5 in 1999

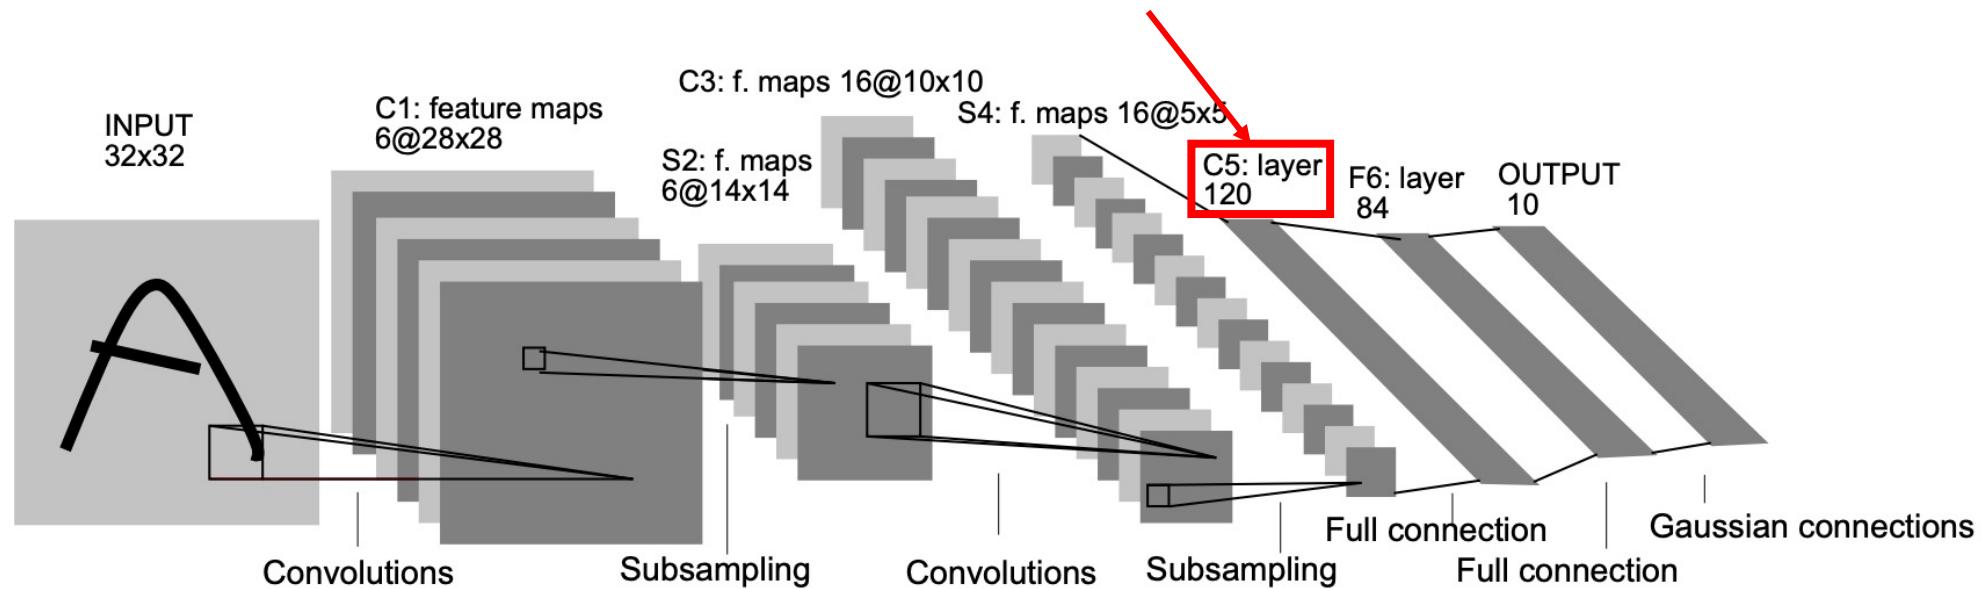


**Fig. 1.** Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Subsampling layer: max/average pooling

# LeNet-5 in 1999

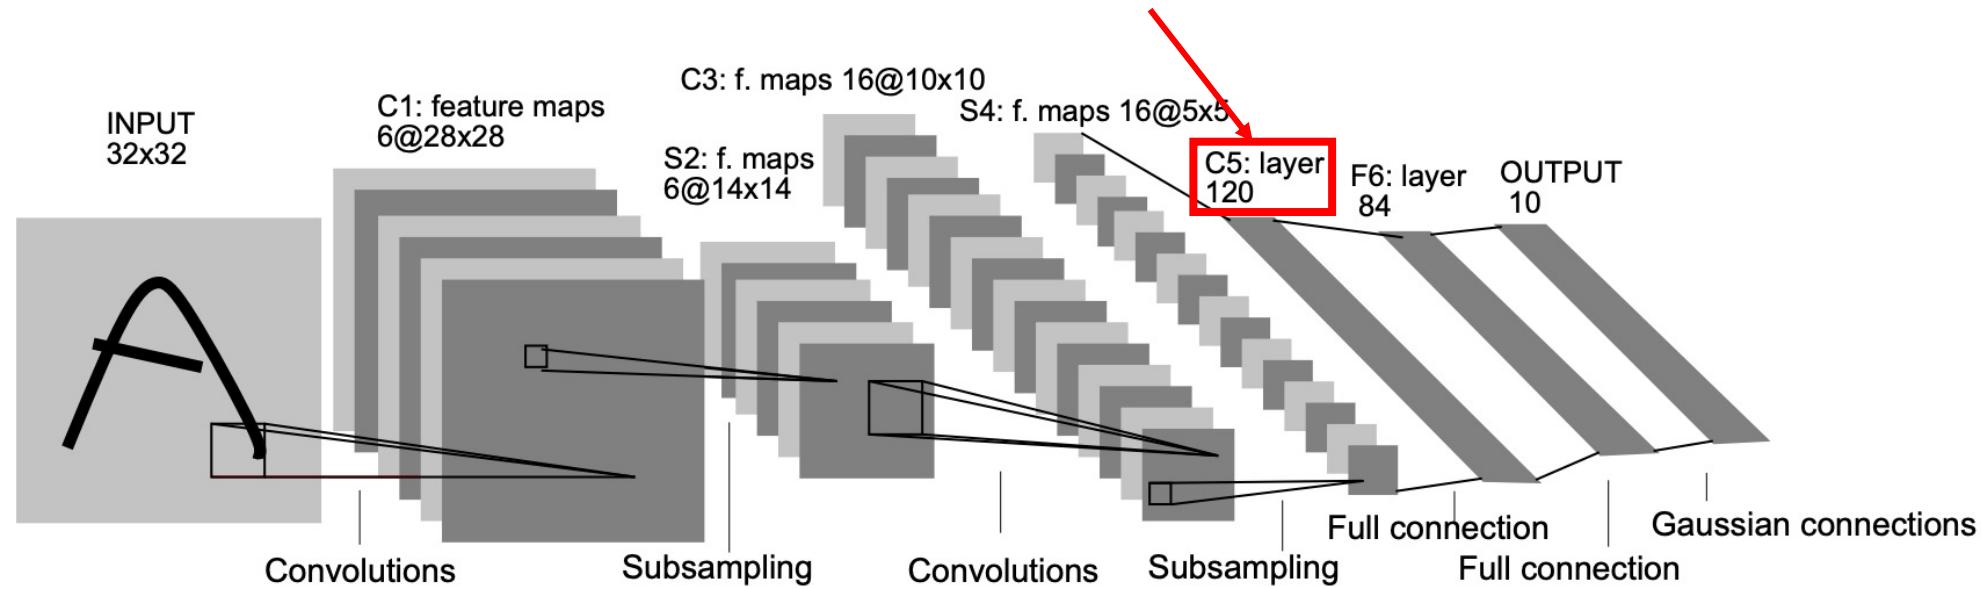
One more **question**:  
How C5 comes from?



**Fig. 1.** Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

# LeNet-5 in 1999

One more **question**:  
How C5 comes from? Matrices → a vector?



**Fig. 1.** Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

# External reading materials

- LeCun, Yann, Patrick Haffner, Léon Bottou, and Yoshua Bengio. "Object recognition with gradient-based learning." In *Shape, contour and grouping in computer vision*, pp. 319-345. Springer, Berlin, Heidelberg, 1999.
  - Online at <http://yann.lecun.com/exdb/publis/pdf/lecun-99.pdf>
  - Section 2.2
  - Understand architecture of LeNet-5
- LeCun, Yann, Léon Bottou, Yoshua Bengio, and Patrick Haffner. "Gradient-based learning applied to document recognition." *Proceedings of the IEEE* 86, no. 11 (1998): 2278-2324.
  - Online at [http://vision.stanford.edu/cs598\\_spring07/papers/Lecun98.pdf](http://vision.stanford.edu/cs598_spring07/papers/Lecun98.pdf)
  - Section II.B