

# Practical Tricks for CNNs

Neural Networks Design And Application

# Practical tricks

- Batch normalization and local response normalization
- Data augmentation
- Dropout
- Regularization/weight decay
- Pre-train
- Stagewise training

# Rescaling images

```
# calculate mu and sig using the training set
d = x_train.shape[1]
mu = numpy.mean(x_train, axis=0).reshape(1, d)
sig = numpy.std(x_train, axis=0).reshape(1, d)

# transform the training features
x_train = (x_train - mu) / (sig + 1E-6)

# transform the test features
x_test = (x_test - mu) / (sig + 1E-6)
```

# Rescaling images

```
# calculate mu and sig using the training set
d = x_train.shape[1]
mu = numpy.mean(x_train, axis=0).reshape(1, d)
sig = numpy.std(x_train, axis=0).reshape(1, d)

# transform the training features
x_train = (x_train - mu) / (sig + 1E-6)

# transform the test features
x_test = (x_test - mu) / (sig + 1E-6)
```

# Rescaling images

```
# calculate mu and sig using the training set
d = x_train.shape[1]
mu = numpy.mean(x_train, axis=0).reshape(1, d)
sig = numpy.std(x_train, axis=0).reshape(1, d)

# transform the training features
x_train = (x_train - mu) / (sig + 1E-6)

# transform the test features
x_test = (x_test - mu) / (sig + 1E-6)
```

Centering images

# Rescaling images

```
# calculate mu and sig using the training set
d = x_train.shape[1]
mu = numpy.mean(x_train, axis=0).reshape(1, d)
sig = numpy.std(x_train, axis=0).reshape(1, d)

# transform the training features
x_train = (x_train - mu) / (sig + 1E-6)

# transform the test features
x_test = (x_test - mu) / (sig + 1E-6)
```

Centering images      Standardize images

# Rescaling images

```
# calculate mu and sig using the training set
d = x_train.shape[1]
mu = numpy.mean(x_train, axis=0).reshape(1, d)
sig = numpy.std(x_train, axis=0).reshape(1, d)


# transform the training features
x_train = (x_train - mu) / (sig + 1E-6)

# transform the test features
x_test = (x_test - mu) / (sig + 1E-6)
```


Centering images      Standardize images

To bound the values of data

# Local response normalization




a) Inter-Channel LRN ( $n=2$ )




b) Intra-Channel LRN ( $n=2$ )



# Local response normalization




# Local response normalization



$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

# Local response normalization




$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

Output value

Original value

# Local response normalization



a) Inter-Channel LRN (n=2)

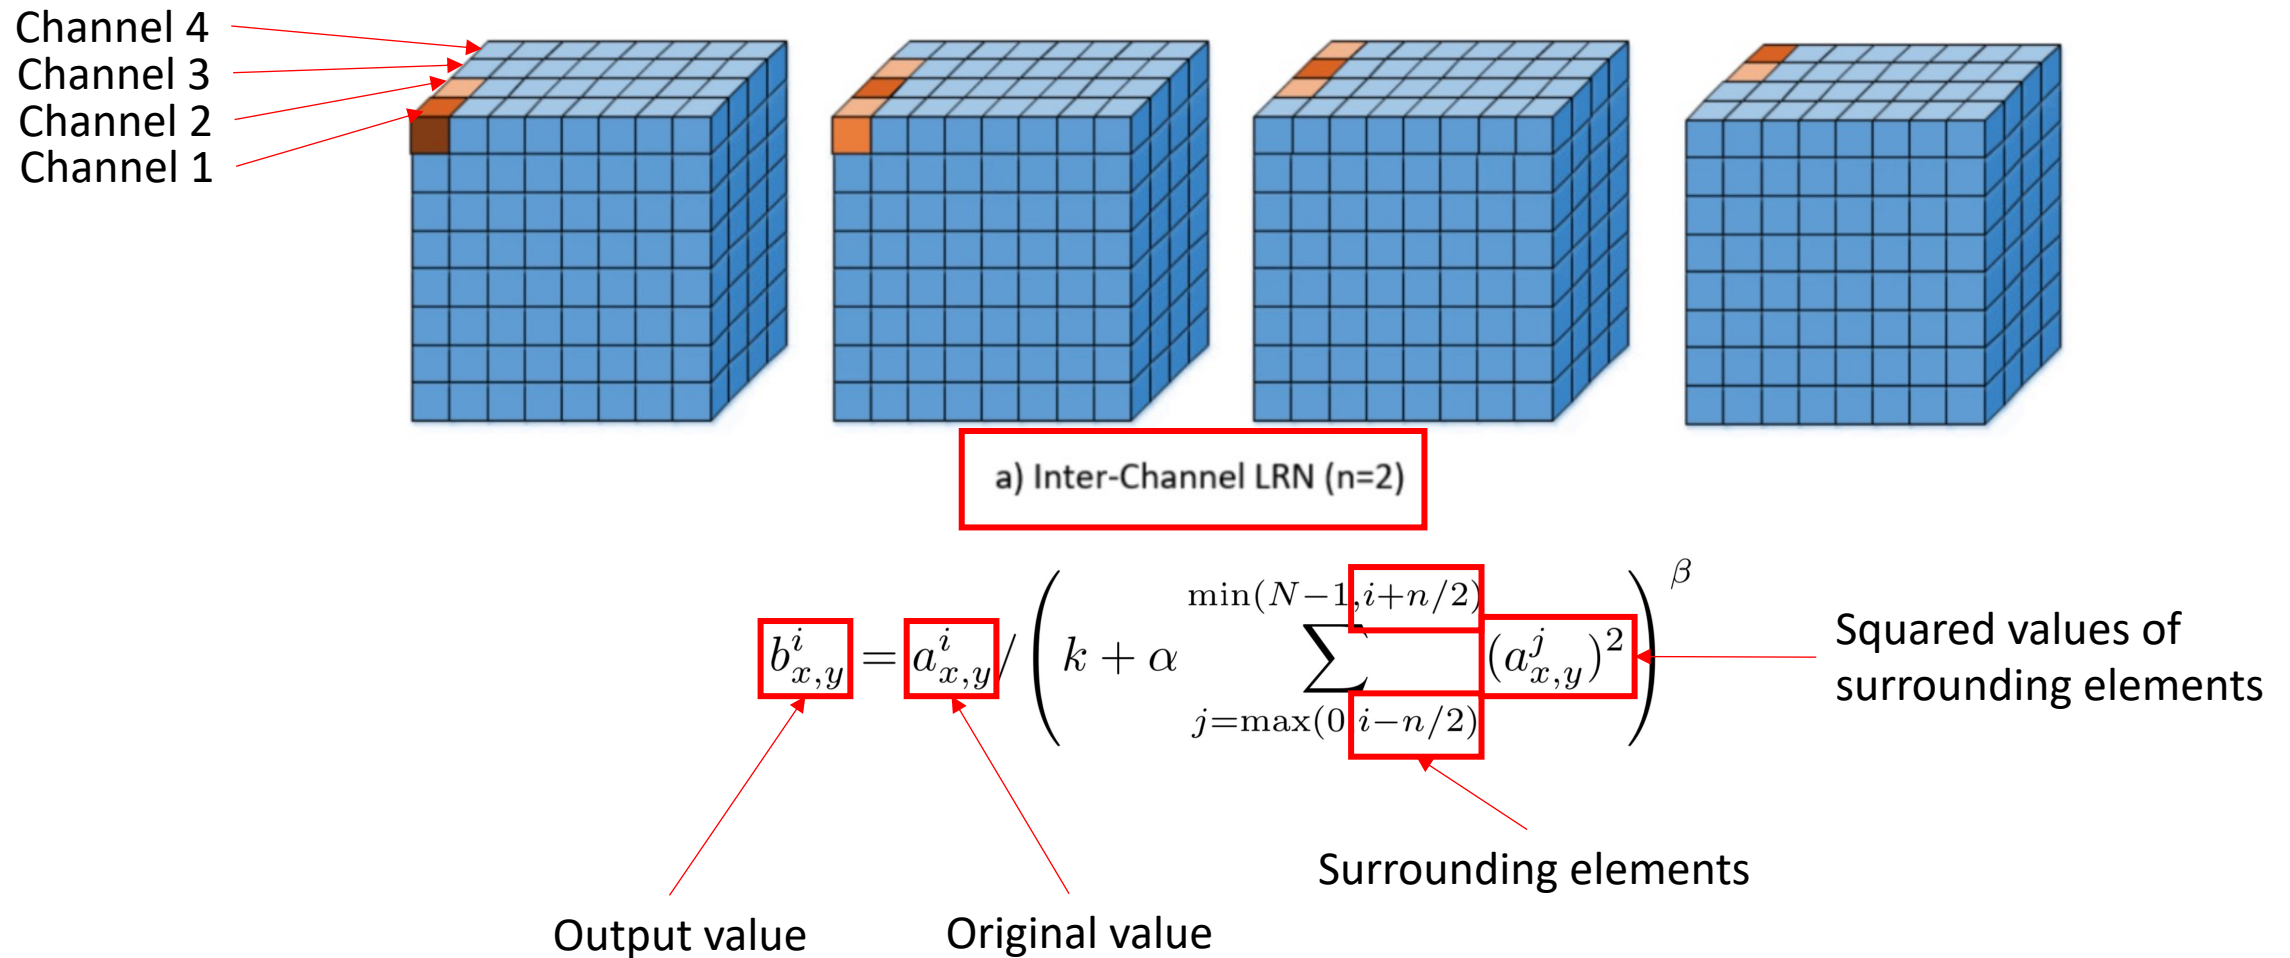
$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

Output value

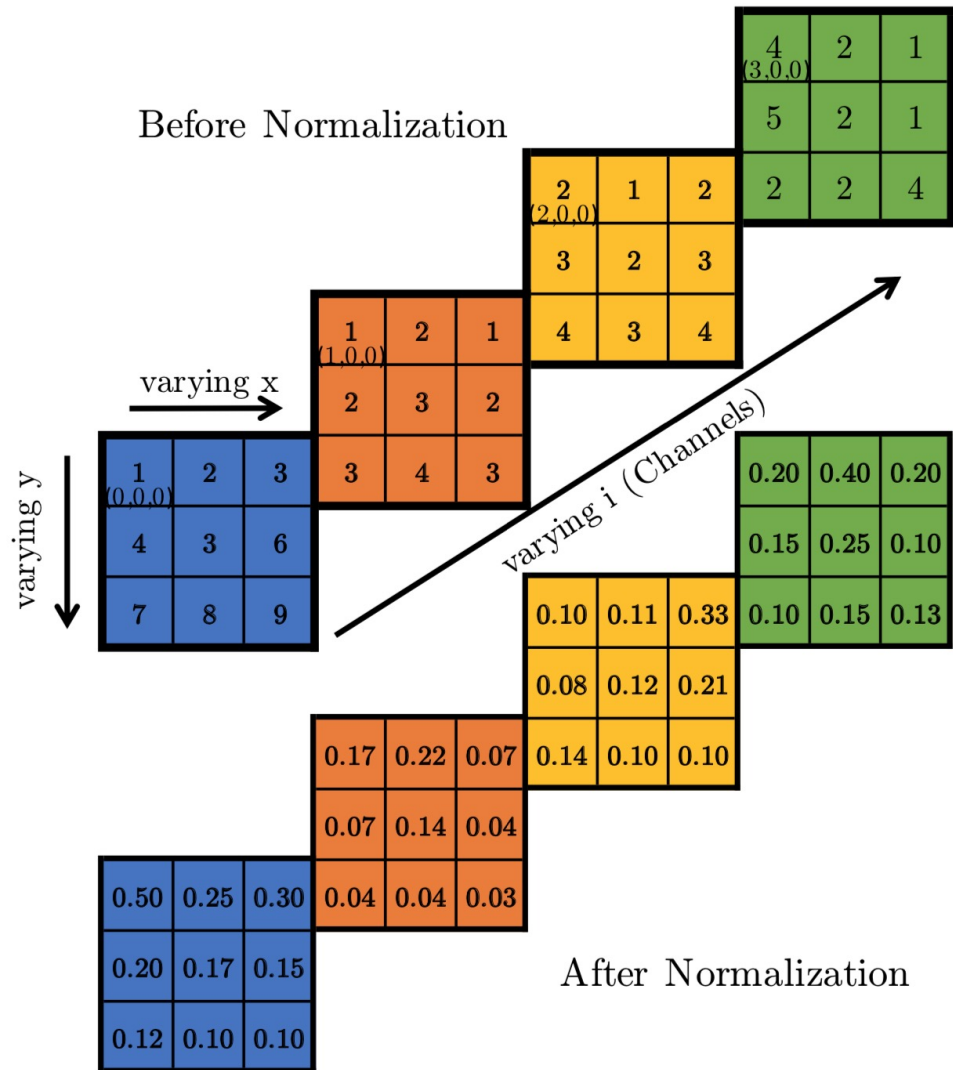
Original value

Surrounding elements

# Local response normalization



# Local response normalization

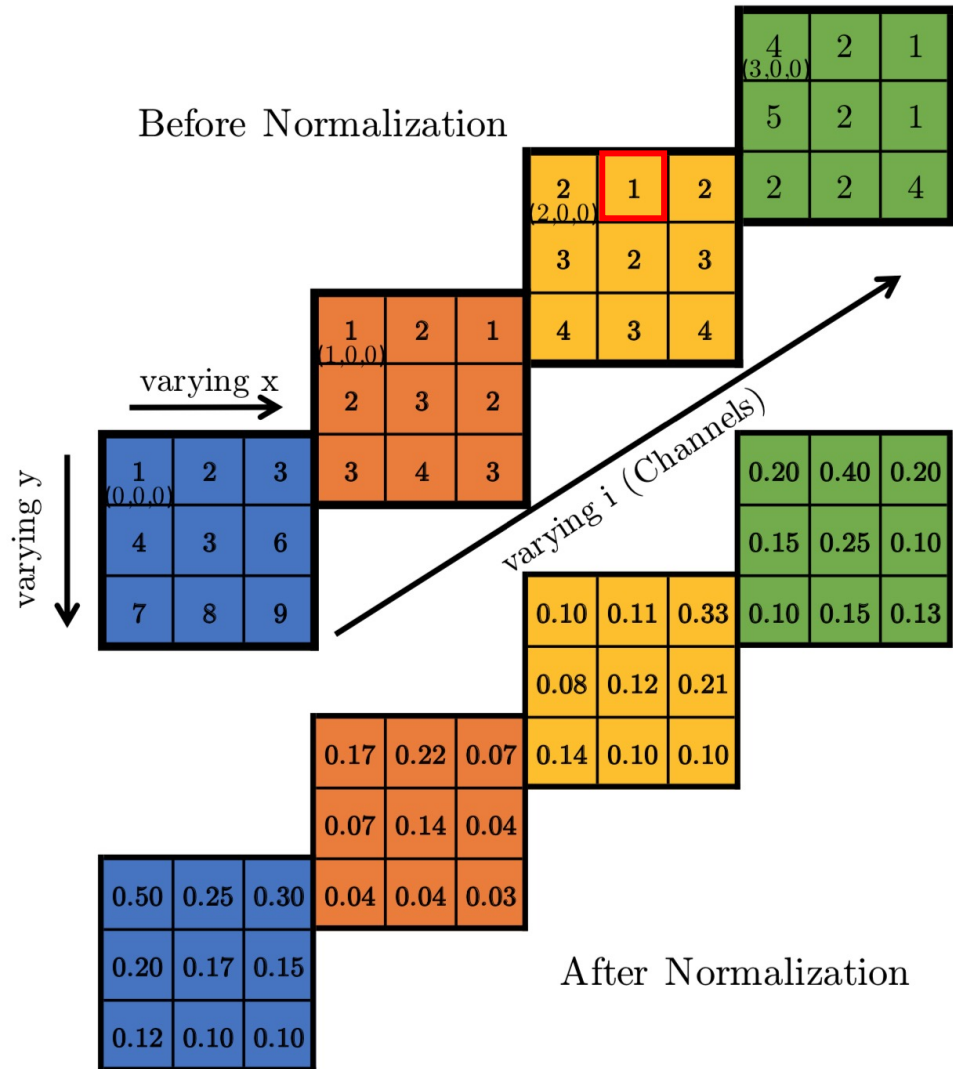


$N=2, K=0, \alpha = 1, \beta = 1$

$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0,i-n/2)}^{\min(N-1,i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$



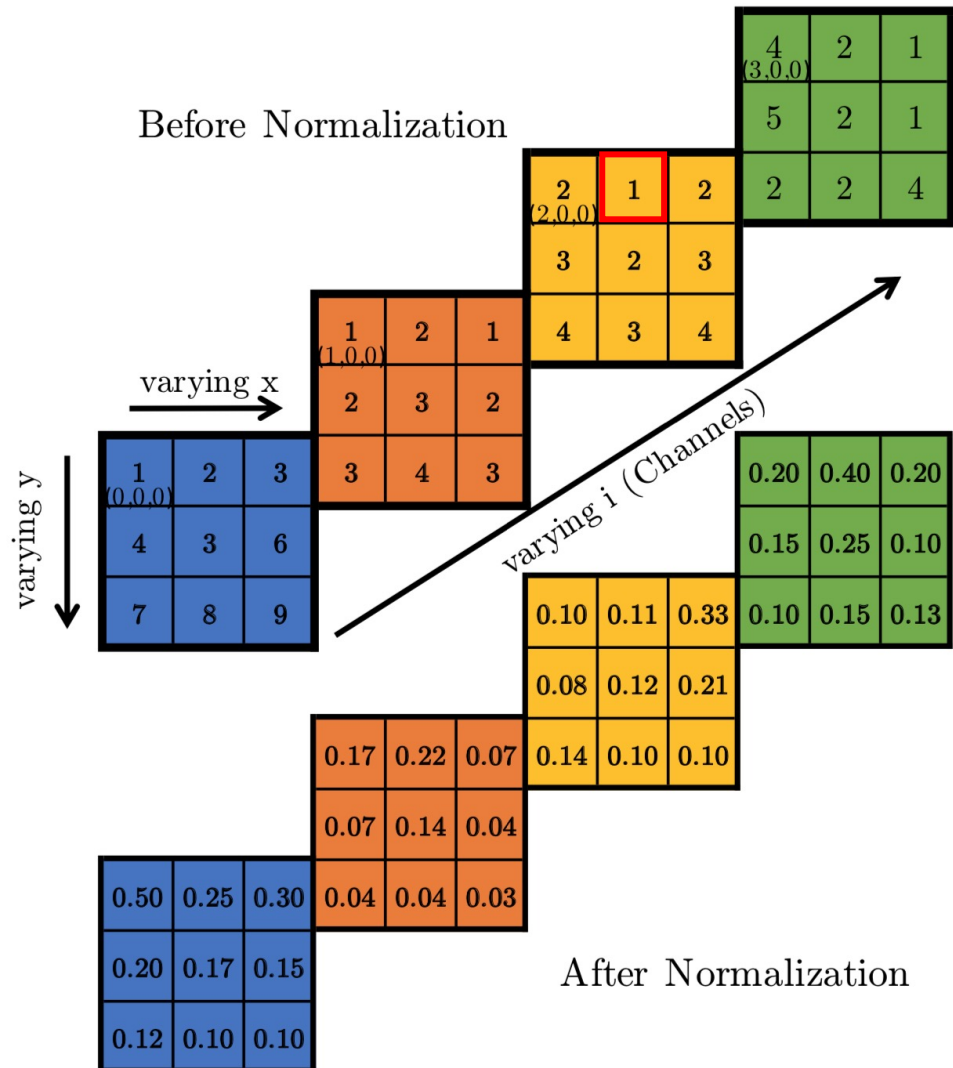
# Local response normalization



$$N=2, K=0, \alpha = 1, \beta = 1$$

$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

# Local response normalization



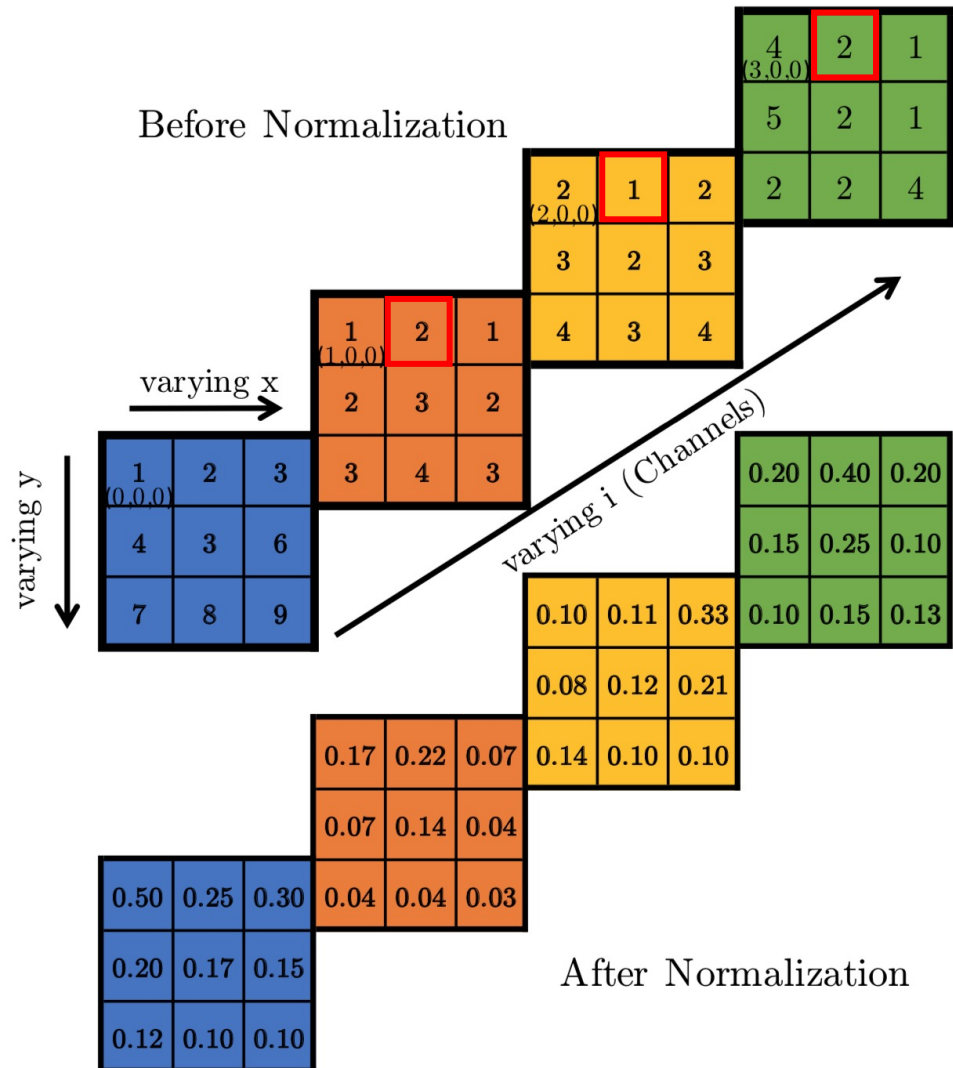
$$N=2, K=0, \alpha = 1, \beta = 1$$

$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

1/( ? ) = ?



# Local response normalization

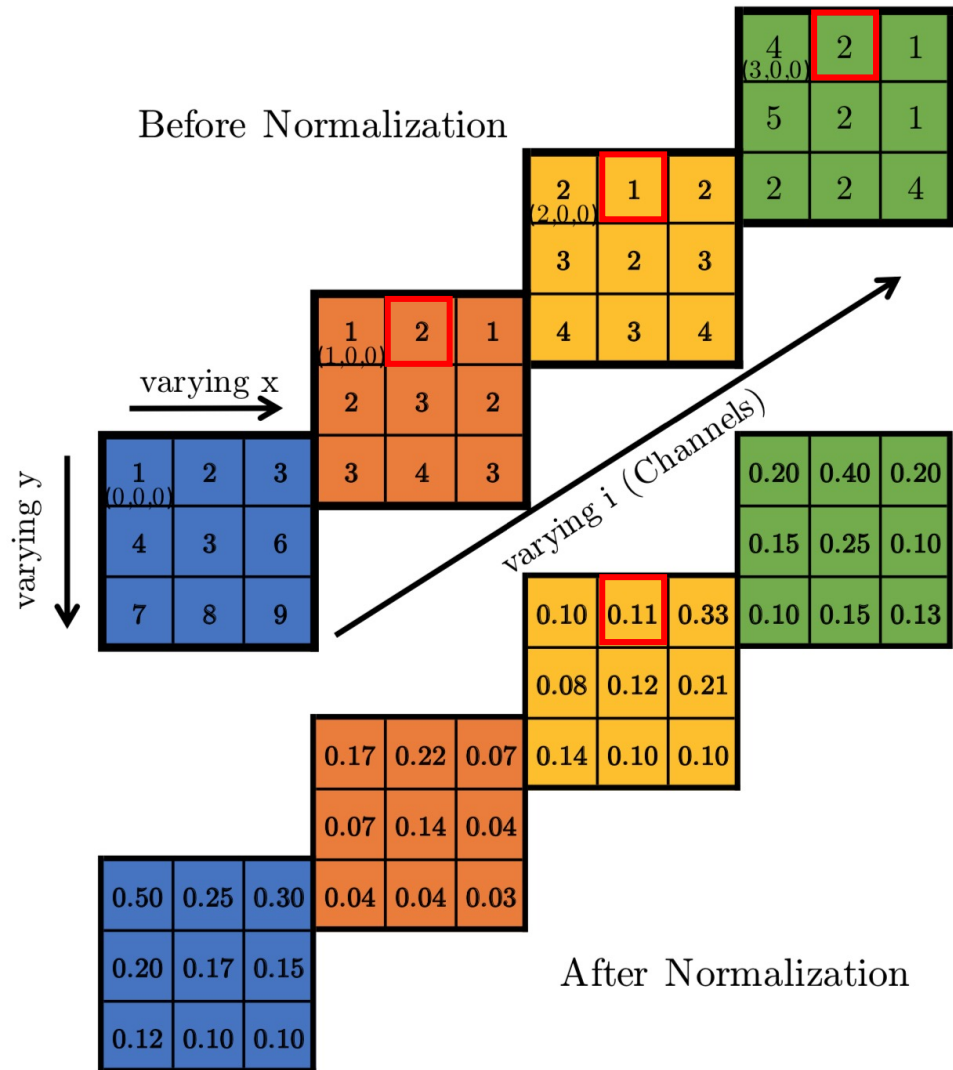


$$N=2, K=0, \alpha = 1, \beta = 1$$

$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

1/(0 + 4 + 1 + 4)=?

# Local response normalization




$N=2, K=0, \alpha = 1, \beta = 1$

$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

$1/(0 + 4 + 1 + 4) = 1/9$


# Local response normalization



b) Intra-Channel LRN (n=2)

$$b_{x,y}^k = a_{x,y}^k / \left( k + \alpha \sum_{i=\max(0,x-n/2)}^{\min(W,x+n/2)} \sum_{j=\max(0,y-n/2)}^{\min(H,y+n/2)} (a_{i,j}^k)^2 \right)^\beta$$

# Local response normalization




b) Intra-Channel LRN (n=2)

$$b_{x,y}^k = a_{x,y}^k / \left( k + \alpha \sum_{i=\max(0,x-n/2)}^{\min(W,x+n/2)} \sum_{j=\max(0,y-n/2)}^{\min(H,y+n/2)} (a_{i,j}^k)^2 \right)^\beta$$

Inter-channel: 
$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0,i-n/2)}^{\min(N-1,i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

# Local response normalization




b) Intra-Channel LRN (n=2)

$$b_{x,y}^k = a_{x,y}^k / \left( k + \alpha \sum_{i=\max(0,x-n/2)}^{\min(W,x+n/2)} \sum_{j=\max(0,y-n/2)}^{\min(H,y+n/2)} (a_{i,j}^k)^2 \right)^\beta$$

Inter-channel:  $b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0,i-n/2)}^{\min(N-1,i+n/2)} (a_{x,y}^j)^2 \right)^\beta$

# Local response normalization



$$b_{x,y}^k = a_{x,y}^k / \left( k + \alpha \sum_{i=\max(0, x-n/2)}^{\min(W, x+n/2)} \sum_{j=\max(0, y-n/2)}^{\min(H, y+n/2)} (a_{i,j}^k)^2 \right)^\beta$$

# Batch normalization [BN]

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;  
Parameters to be learned:  $\gamma, \beta$   
**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$
$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.



# Batch normalization [BN]

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.



# Batch normalization [BN]

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Rescaling for a batch

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

# Batch normalization [BN]

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Rescaling for a batch

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

# Batch normalization [BN]

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Rescaling for a batch

A linear model as output

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

# Batch normalization [BN]

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$



Rescaling for a batch

A linear model as output:  
There are two learnable parameters

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

# Data augmentation



- Increase the amount of data by:
  - Adding slightly modified copies of already existing data, or
  - Newly created synthetic data from existing data



Vertical flipping

# Data augmentation



- Increase the amount of data by:
  - Adding slightly modified copies of already existing data, or
  - Newly created synthetic data from existing data



Vertical flipping

# Data augmentation

- Increase the amount of data by:
  - Adding slightly modified copies of already existing data, or
  - Newly created synthetic data from existing data




Horizontal flipping



# Data augmentation



- Increase the amount of data by:
  - Adding slightly modified copies of already existing data, or
  - Newly created synthetic data from existing data





# Data augmentation



- Increase the amount of data by:
  - Adding slightly modified copies of already existing data, or
  - Newly created synthetic data from existing data



Changing brightness

# Data augmentation

- Increase the amount of data by:
  - Adding slightly modified copies of already existing data, or
  - Newly created synthetic data from existing data



Color jitter:  
brightness, contrast, saturation, hue

# Why data augmentation



The two classes in our hypothetical dataset. The one in the left represents Brand A (Ford), and the one in the right represents Brand B (Chevrolet).

# Why data augmentation



The two classes in our hypothetical dataset. The one in the left represents Brand A (Ford), and the one in the right represents Brand B (Chevrolet).

Consider: when our images only contain Ford cars facing left and Chevrolet cars facing right...



# Why data augmentation



The two classes in our hypothetical dataset. The one in the left represents Brand A (Ford), and the one in the right represents Brand B (Chevrolet).

Consider: when our images only contain Ford cars facing left and Chevrolet cars facing right...



A Ford car (Brand A), but facing right.

# Why data augmentation



The two classes in our hypothetical dataset. The one in the left represents Brand A (Ford), and the one in the right represents Brand B (Chevrolet).

Consider: when our images only contain Ford cars facing left and **Chevrolet cars facing right**...



A Ford car (Brand A), but facing right.

Our CNN may predict this car (**facing right**) to Chevrolet...

# Why data augmentation



The two classes in our hypothetical dataset. The one in the left represents Brand A (Ford), and the one in the right represents Brand B (Chevrolet).

Consider: when our images only contain Ford cars facing left and Chevrolet cars facing right...



A Ford car (Brand A), but facing right.

Our CNN may predict this car (**facing right**) to Chevrolet...

Data augmentation:  
Gives more variations for data



# Why data augmentation



The two classes in our hypothetical dataset. The one in the left represents Brand A (Ford), and the one in the right represents Brand B (Chevrolet).

Consider: when our images only contain Ford cars facing left and Chevrolet cars facing right...



A Ford car (Brand A), but facing right.

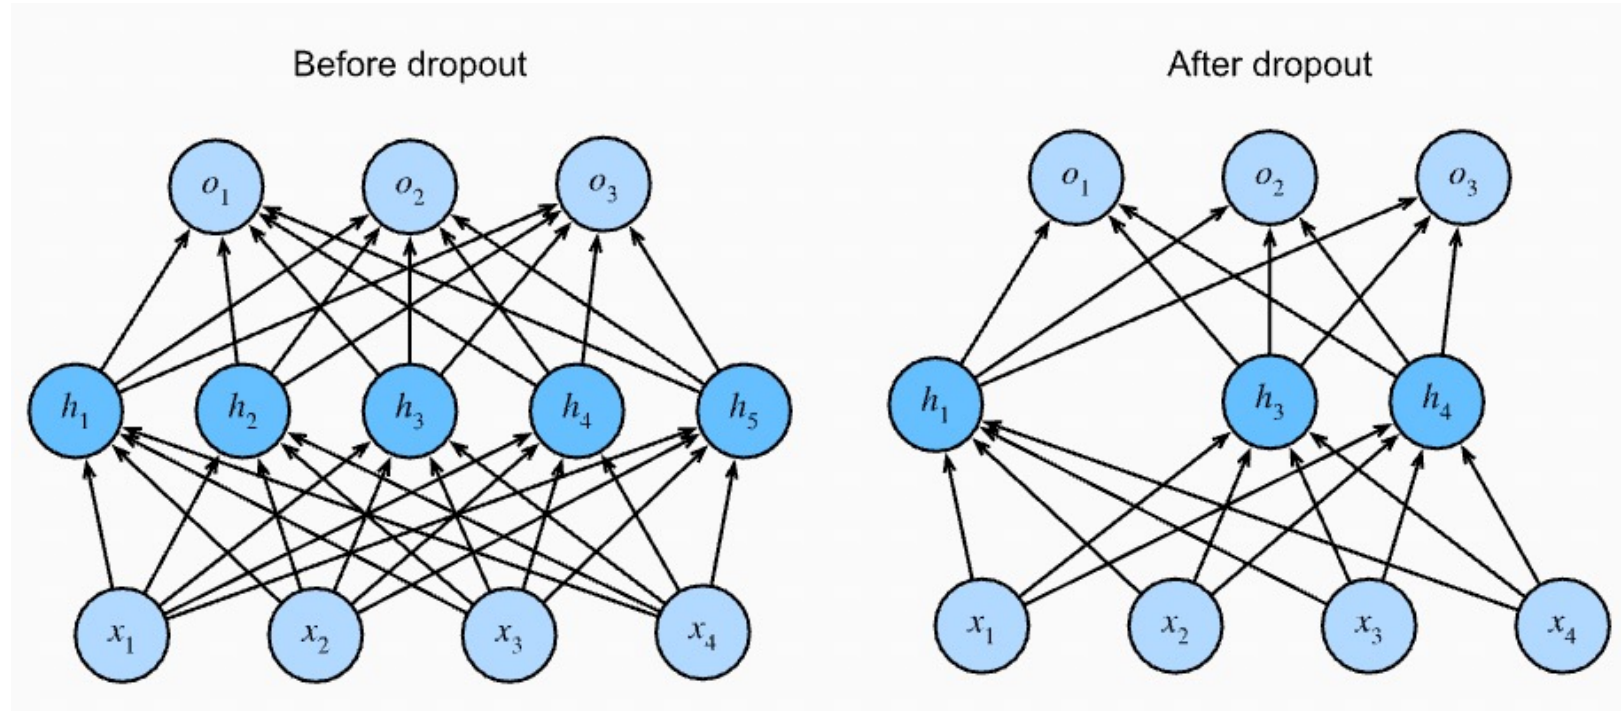
Our CNN may predict this car (**facing right**) to Chevrolet...

Data augmentation:


Gives more variations for data → better generalization




# Dropout



# Dropout

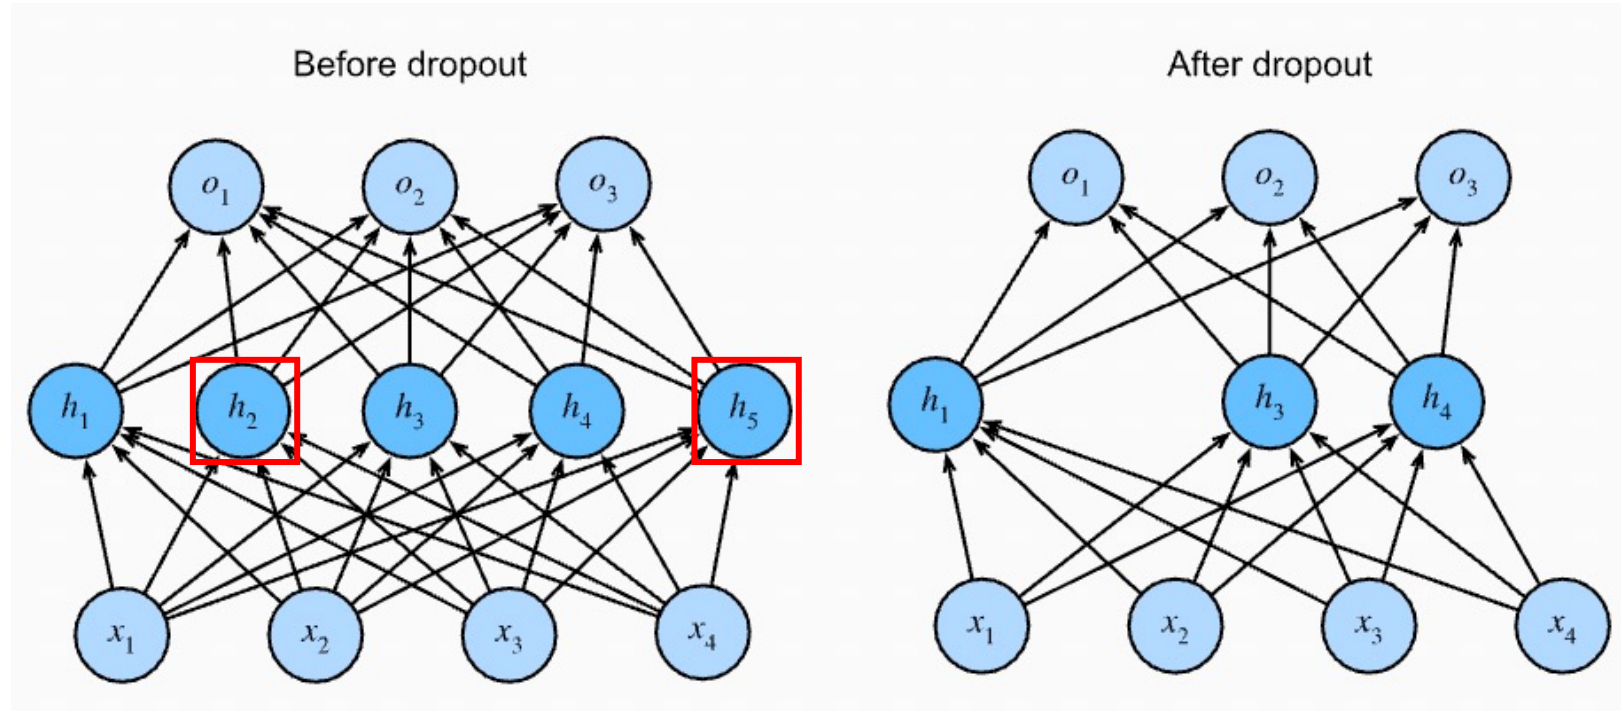


# Dropout



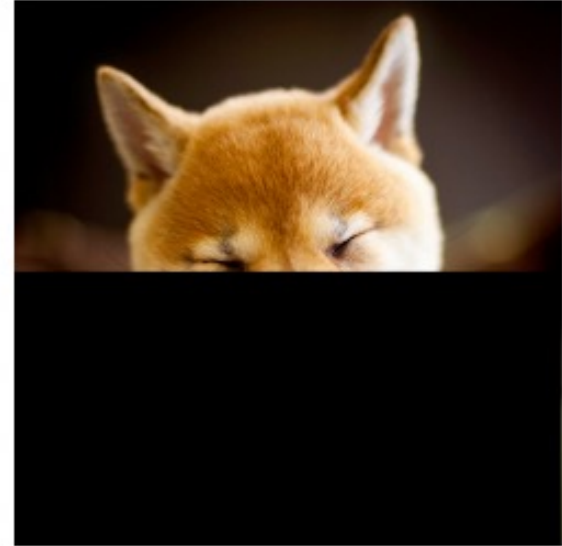
Dropout in training: select an arbitrary percentage of neurons (weights) and mask them

# Dropout



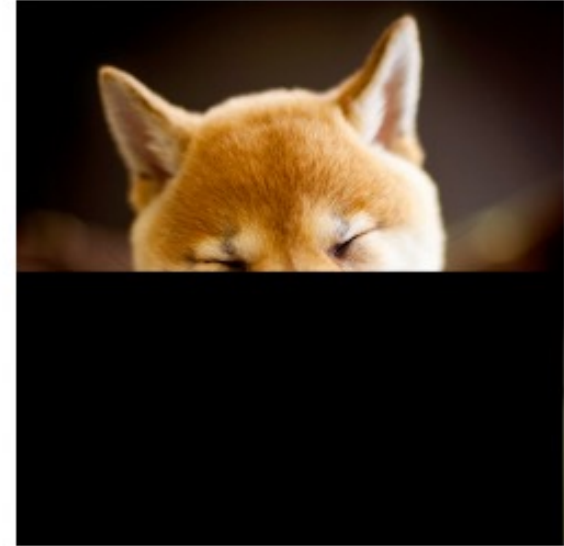
Dropout in training: select an arbitrary percentage of neurons (weights) and mask them  
Dropout in testing: use all parameters, no dropout

# Dropout



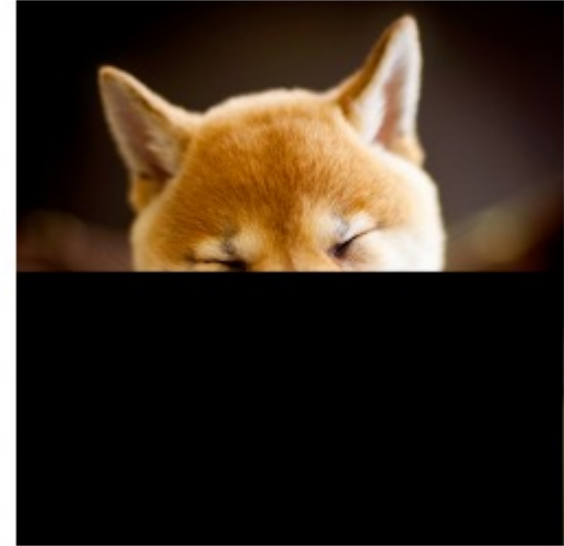


# Dropout



Why dropout?

# Dropout




Why dropout? → alleviate overfitting



# Regularization/weight decay


$$\tilde{J}(\boldsymbol{\theta}; \mathbf{X}, \mathbf{y}) = J(\boldsymbol{\theta}; \mathbf{X}, \mathbf{y}) + \alpha\Omega(\boldsymbol{\theta}),$$

$$\Omega(\boldsymbol{\theta}) = \frac{1}{2}\|\mathbf{w}\|_2^2$$


# Regularization/weight decay

$$\tilde{J}(\theta; \mathbf{X}, \mathbf{y}) = J(\theta; \mathbf{X}, \mathbf{y}) + \alpha\Omega(\theta),$$

$$\Omega(\theta) = \frac{1}{2}\|\mathbf{w}\|_2^2$$





Origin


# Regularization/weight decay

$$\tilde{J}(\theta; \mathbf{X}, \mathbf{y}) = J(\theta; \mathbf{X}, \mathbf{y}) + \alpha\Omega(\theta),$$


$$\Omega(\theta) = \frac{1}{2}\|\mathbf{w}\|_2^2$$



Linear  
model



Quadratic  
model




Polynomial model  
(9 degree)

# Regularization/weight decay


$$\tilde{J}(\theta; \mathbf{X}, \mathbf{y}) = J(\theta; \mathbf{X}, \mathbf{y}) + \alpha\Omega(\theta),$$

$$\Omega(\theta) = \frac{1}{2}\|\mathbf{w}\|_2^2$$




Origin


Improve generalization performance



Linear  
model




Quadratic  
model



Polynomial model  
(9 degree)


# Pre-train



Suppose we have a learned model → weight parameters are determined and fixed

# Pre-train

AlexNet for ImageNet: 1000 classes



Suppose we have a learned model → weight parameters are determined and fixed

# Pre-train


Now: **Two** classes



VS



## AlexNet for ImageNet: 1000 classes




Suppose we have a learned model  $\rightarrow$  weight parameters are determined and fixed



# Pre-train


Now: **Two** classes



Suppose we have a learned model → weight parameters are determined and fixed


# Pre-train

Now: **Two** classes




Suppose we have a learned model → weight parameters are determined and fixed

# Stagewise/restart training



# Stagewise/restart training



## One-loop SGD


For  $t=1 \rightarrow T$

Compute stochastic gradients  $G_t$  for  $w_t$

Update  $w_{t+1} = w_t - \eta G_t$

Endfor


# Stagewise/restart training



## Two-loop SGD

```
For s=1→S
  For t=1→T
    Compute stochastic gradients  $G_t$  for  $w_t$ 
    Update  $w_{t+1} = w_t - \eta_s G_t$ 
  Endfor
Endfor
```

# Stagewise/restart training




## Two-loop SGD

```
For s=1→S
  For t=1→T
    Compute stochastic gradients  $G_t$  for  $w_t$ 
    Update  $w_{t+1} = w_t - \eta_s G_t$ 
  Endfor
Endfor
```



# Stagewise/restart training



## Two-loop SGD

```
For s=1→S
  For t=1→T
    Compute stochastic gradients  $G_t$  for  $w_t$ 
    Update  $w_{t+1} = w_t - \eta_s G_t$ 
  Endfor
Endfor
```

# Many other tricks for CNNs

- Large mini-batch in stochastic gradient descent
- Learning rate warmup [warmup]
- Mixup augmentation [mixup]
- Others, e.g., [BagOfTricks]

# References

- [BN] Ioffe, Sergey, and Christian Szegedy. "Batch normalization: Accelerating deep network training by reducing internal covariate shift." *arXiv preprint arXiv:1502.03167* (2015).
- [warmup] Goyal, Priya, Piotr Dollár, Ross Girshick, Pieter Noordhuis, Lukasz Wesolowski, Aapo Kyrola, Andrew Tulloch, Yangqing Jia, and Kaiming He. "Accurate, large minibatch sgd: Training imagenet in 1 hour." *arXiv preprint arXiv:1706.02677* (2017).
- [mixup] Zhang, Hongyi, Moustapha Cisse, Yann N. Dauphin, and David Lopez-Paz. "mixup: Beyond empirical risk minimization." *arXiv preprint arXiv:1710.09412* (2017).
- [BagOfTricks] He, Tong, Zhi Zhang, Hang Zhang, Zhongyue Zhang, Junyuan Xie, and Mu Li. "Bag of tricks for image classification with convolutional neural networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 558-567. 2019.