Practical GAN Training

Neural Networks Design And Application

Tricks

Architecture guidelines for stable Deep Convolutional GANs

- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.

Radford, Alec, Luke Metz, and Soumith Chintala. "Unsupervised representation learning with deep convolutional generative adversarial networks." *arXiv preprint arXiv:1511.06434* (2015). https://arxiv.org/pdf/1511.06434.pdf

https://machinelearningmastery.com/how-to-train-stable-generative-adversarial-networks/

Replace pooling with strided conv layer

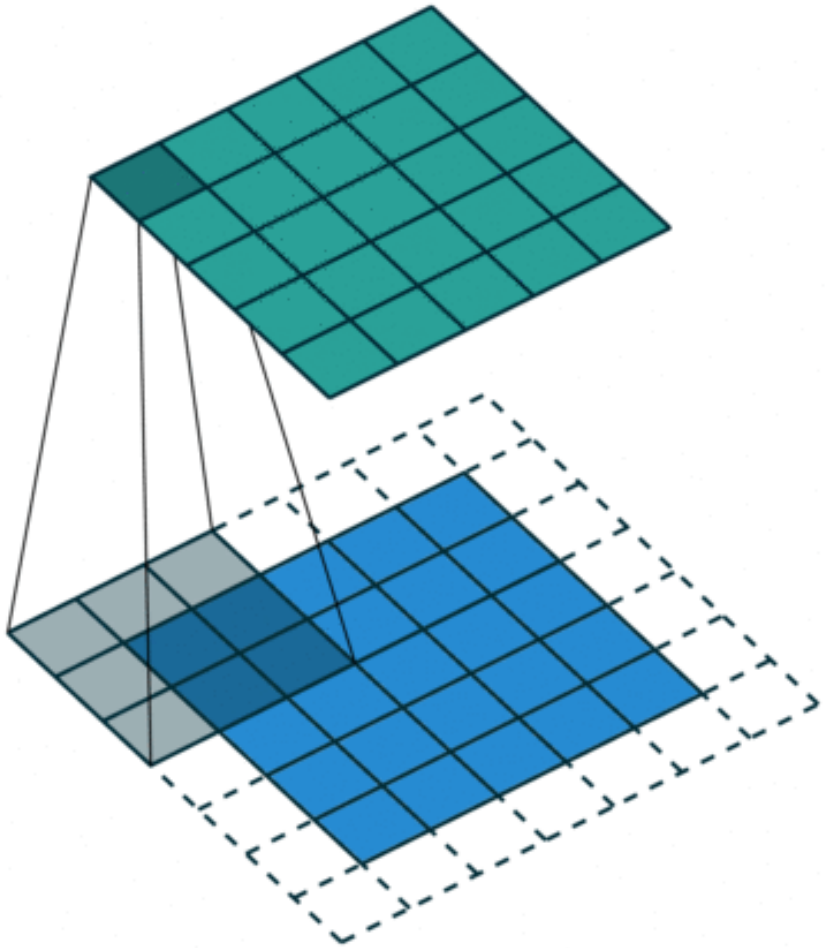

Replace pooling with strided conv layer

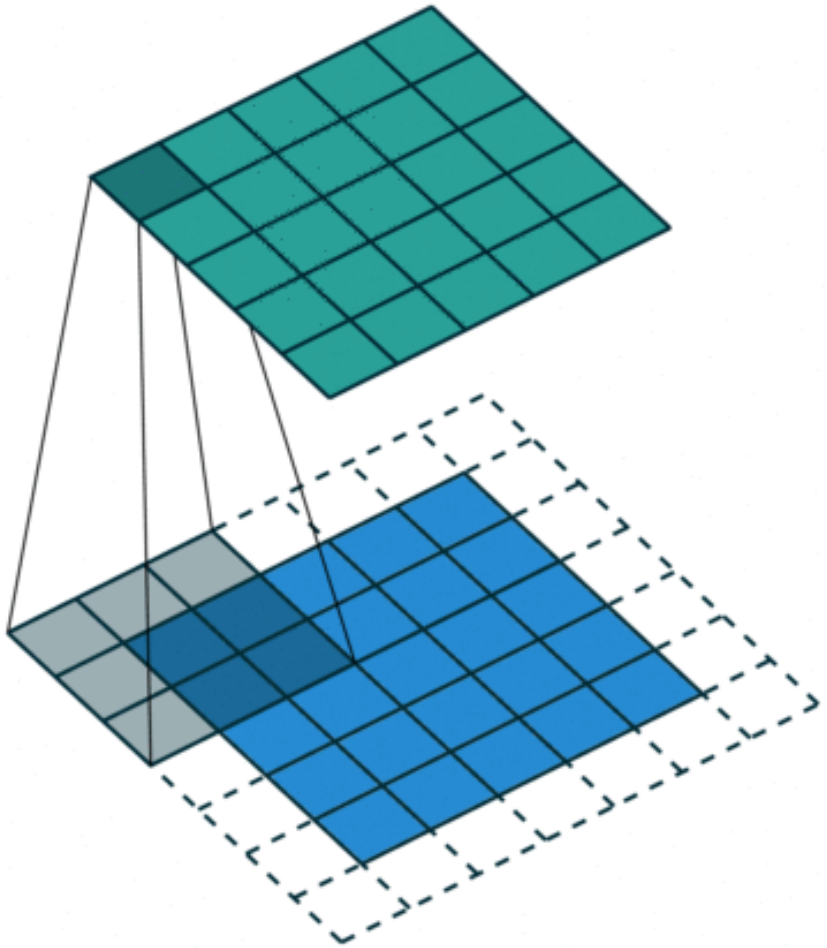

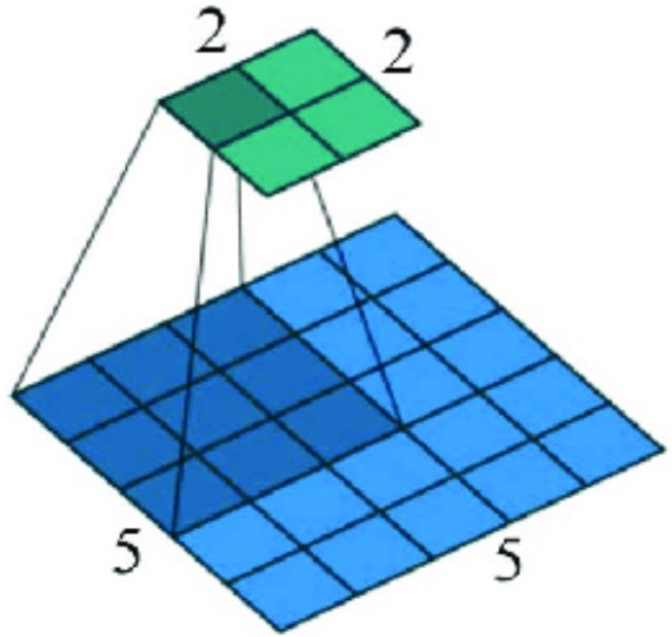

Stride 2 Padding 0

Strided Convolution

Batch normalization [BN]

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1...m}\}$;
Parameters to be learned: γ, β
Output: $\{y_i = \mathrm{BN}_{\gamma,\beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m}\sum_{i=1}^{m} x_i \qquad \text{// mini-batch mean}$$

$$\sigma^2_{\mathcal{B}} \leftarrow \frac{1}{m}\sum_{i=1}^{m}(x_i - \mu_{\mathcal{B}})^2 \qquad \text{// mini-batch variance}$$

$$\widehat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma^2_{\mathcal{B}} + \epsilon}} \qquad \text{// normalize}$$

$$y_i \leftarrow \gamma \widehat{x}_i + \beta \equiv \mathrm{BN}_{\gamma,\beta}(x_i) \qquad \text{// scale and shift}$$

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.

Rescaling for a batch

A linear model as output:
There are two learnable parameters

LeNet-5 in 1999

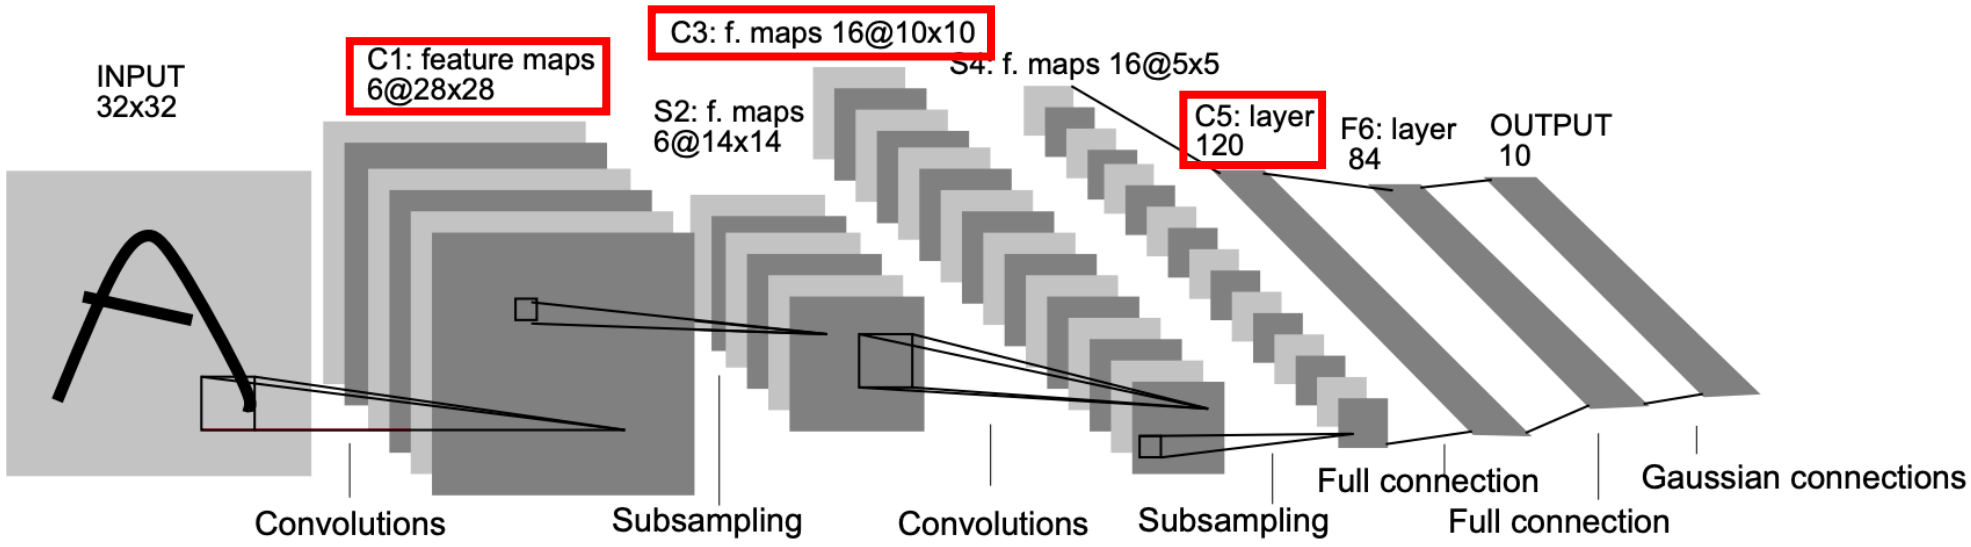

Fig. 1. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeCun, Yann, Patrick Haffner, Léon Bottou, and Yoshua Bengio. "Object recognition with gradient-based learning." In *Shape, contour and grouping in computer vision*, pp. 319-345. Springer, Berlin, Heidelberg, 1999.

Remove full connected layer

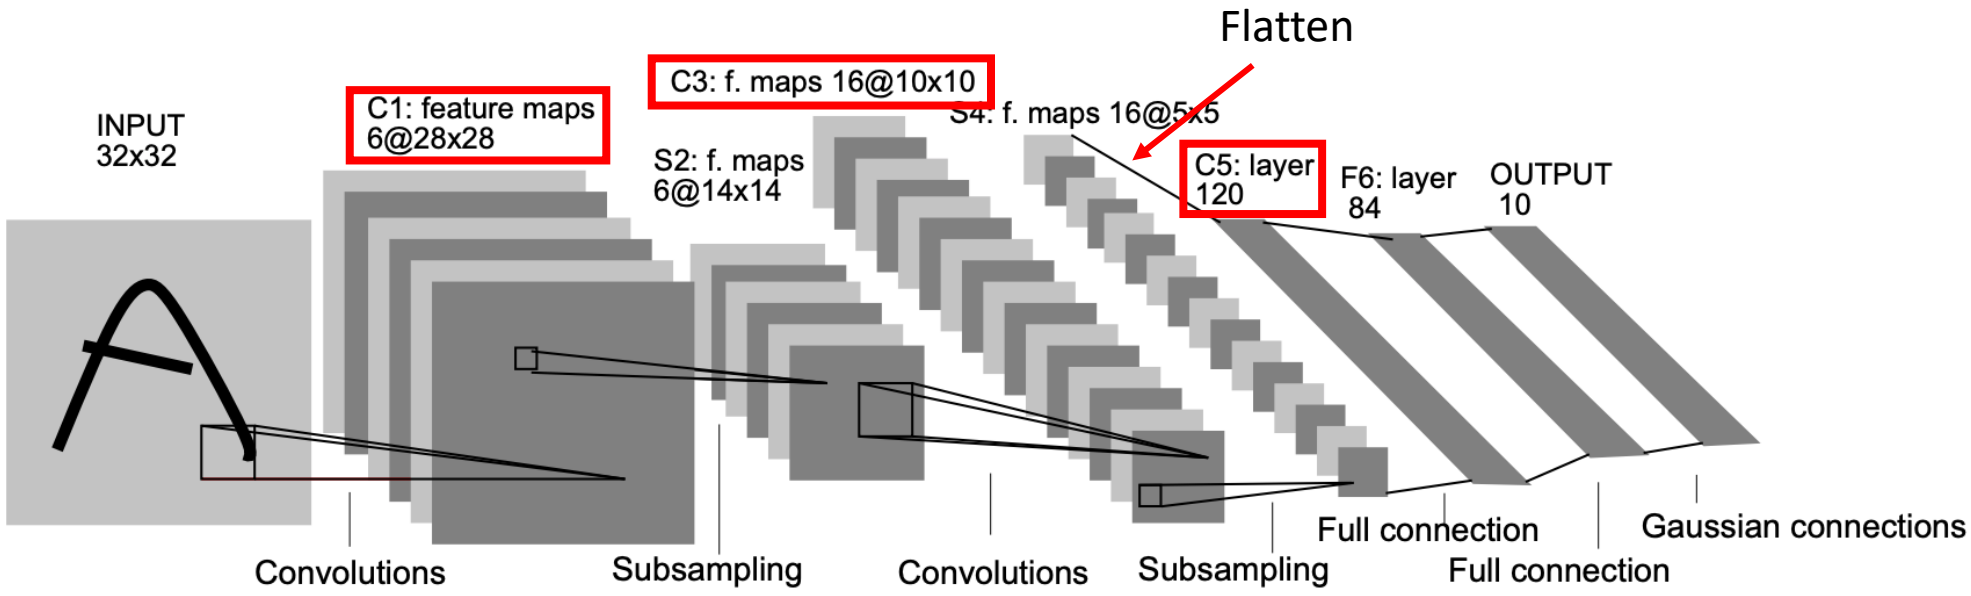

Fig. 1. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeCun, Yann, Patrick Haffner, Léon Bottou, and Yoshua Bengio. "Object recognition with gradient-based learning." In *Shape, contour and grouping in computer vision*, pp. 319-345. Springer, Berlin, Heidelberg, 1999.

Remove full connected layer

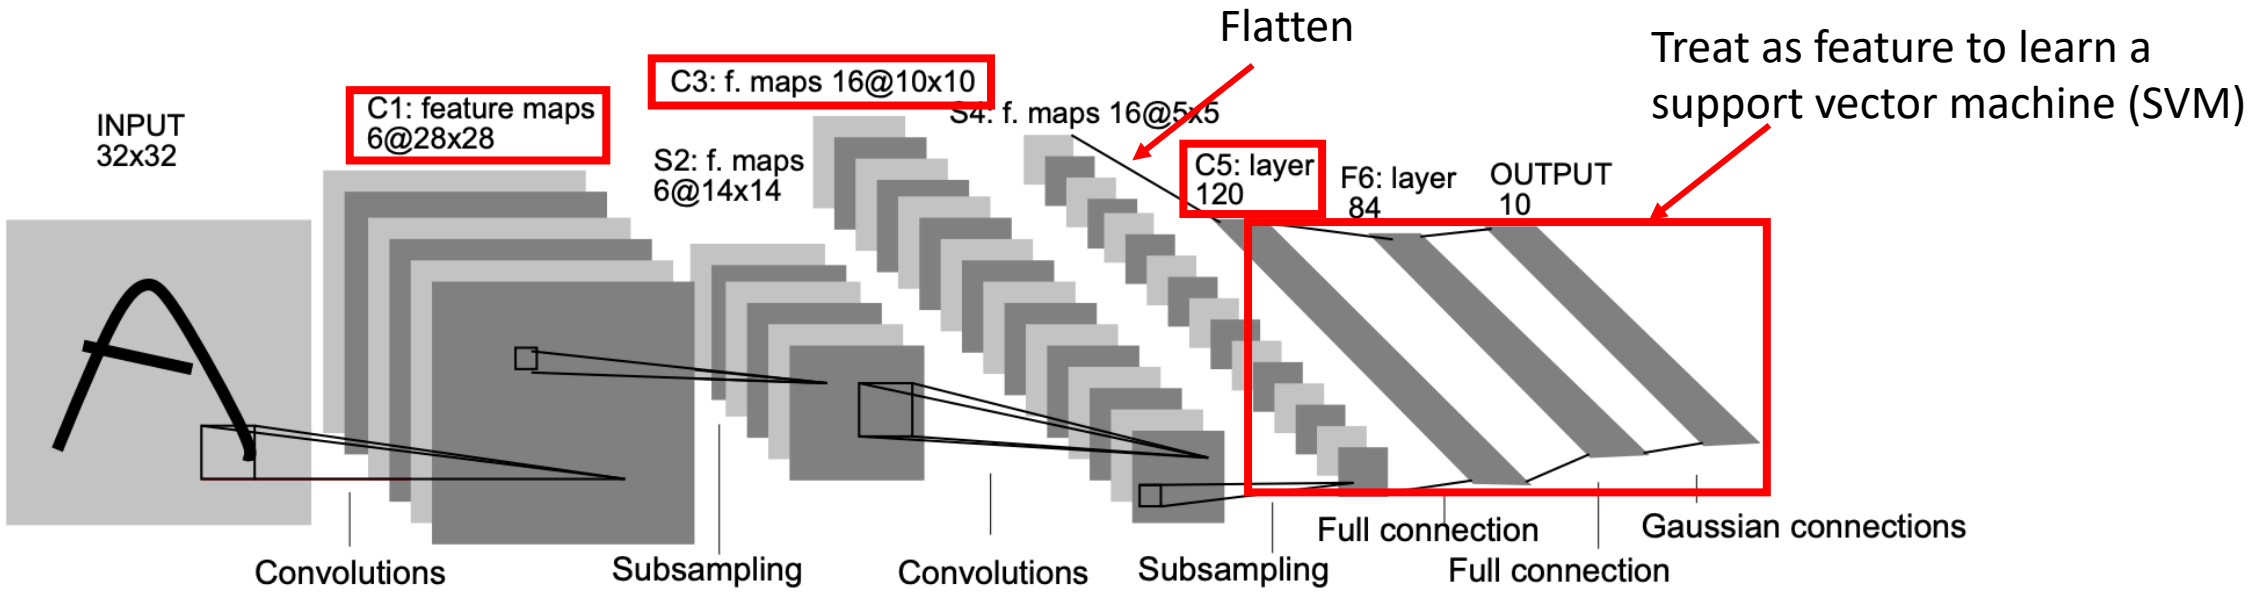

Fig. 1. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeCun, Yann, Patrick Haffner, Léon Bottou, and Yoshua Bengio. "Object recognition with gradient-based learning." In *Shape, contour and grouping in computer vision*, pp. 319-345. Springer, Berlin, Heidelberg, 1999.

Architectures

- CycleGAN
- StyleGAN
- …

https://machinelearningmastery.com/tour-of-generative-adversarial-network-models/
https://neptune.ai/blog/6-gan-architectures

GAN

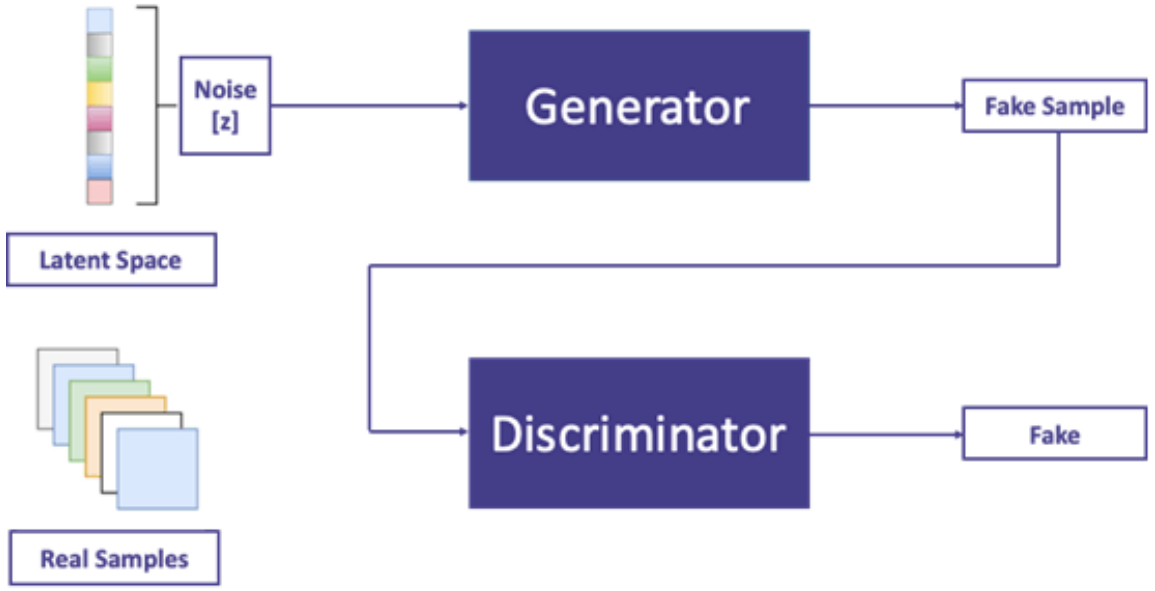

$$\min_G \max_D V(D,G) = \mathbb{E}_{\boldsymbol{x}\sim p_{\text{data}}(\boldsymbol{x})}[\log D(\boldsymbol{x})] + \mathbb{E}_{\boldsymbol{z}\sim p_{\boldsymbol{z}}(\boldsymbol{z})}[\log(1 - D(G(\boldsymbol{z})))].$$

Image credit https://neptune.ai/blog/6-gan-architectures

CycleGAN

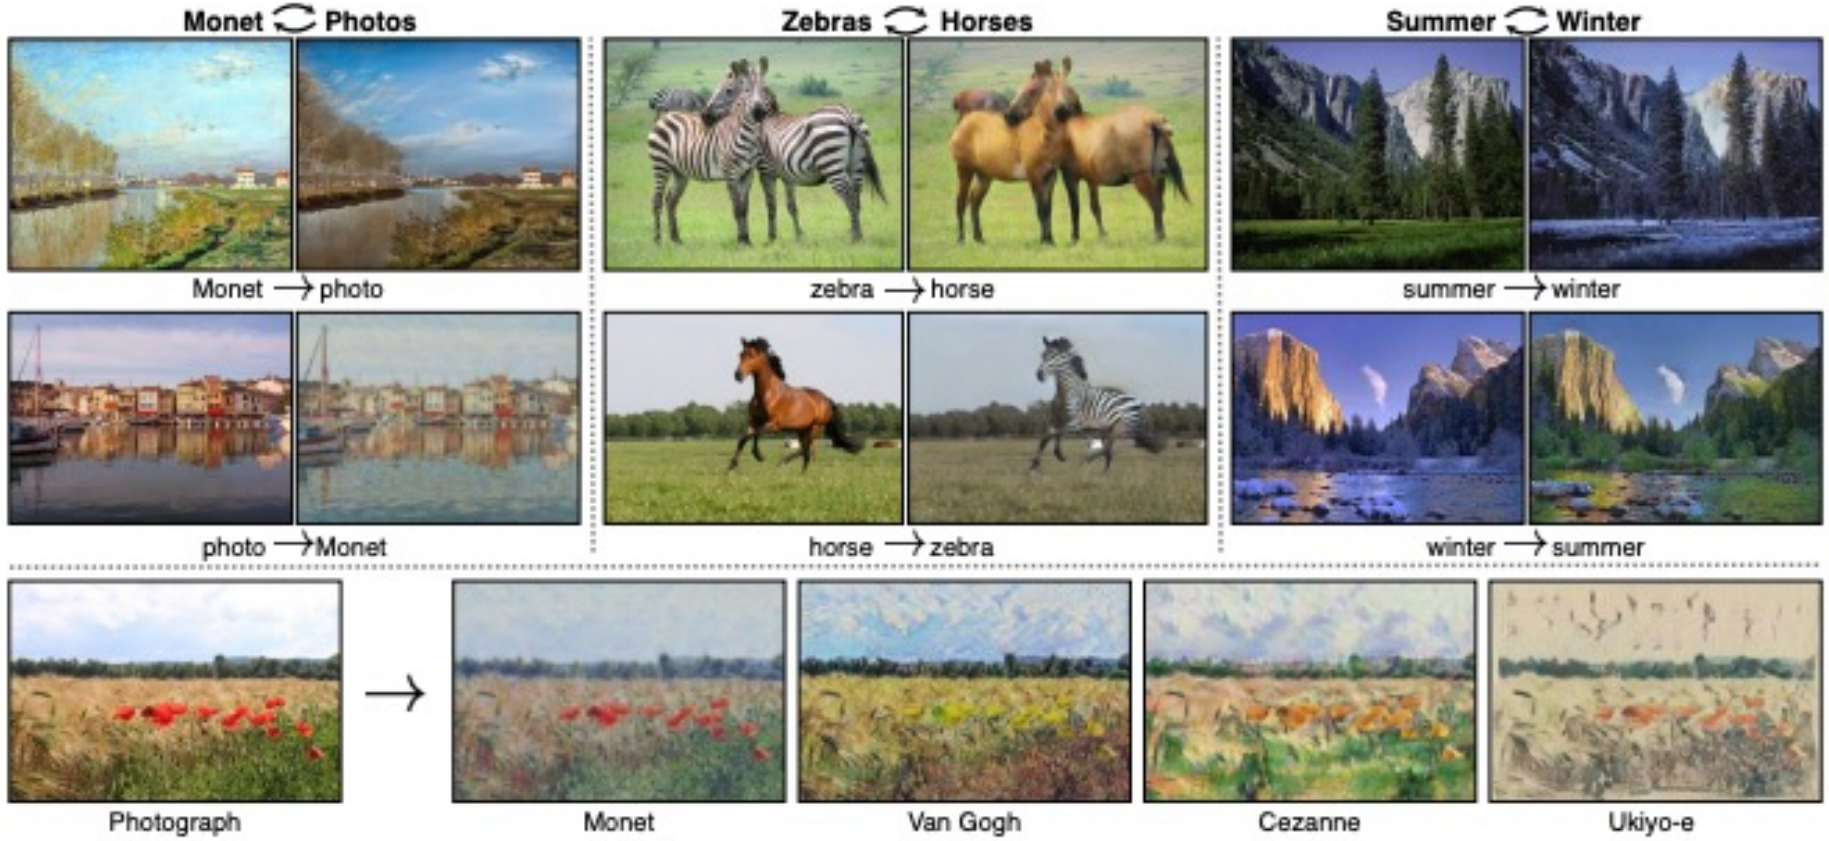
Monet ⇄ Photos
Zebras ⇄ Horses
Summer ⇄ Winter
Monet → photo
zebra → horse
summer → winter
photo → Monet
horse → zebra
winter → summer
Photograph
Monet
Van Gogh
Cezanne
Ukiyo-e

CycleGAN

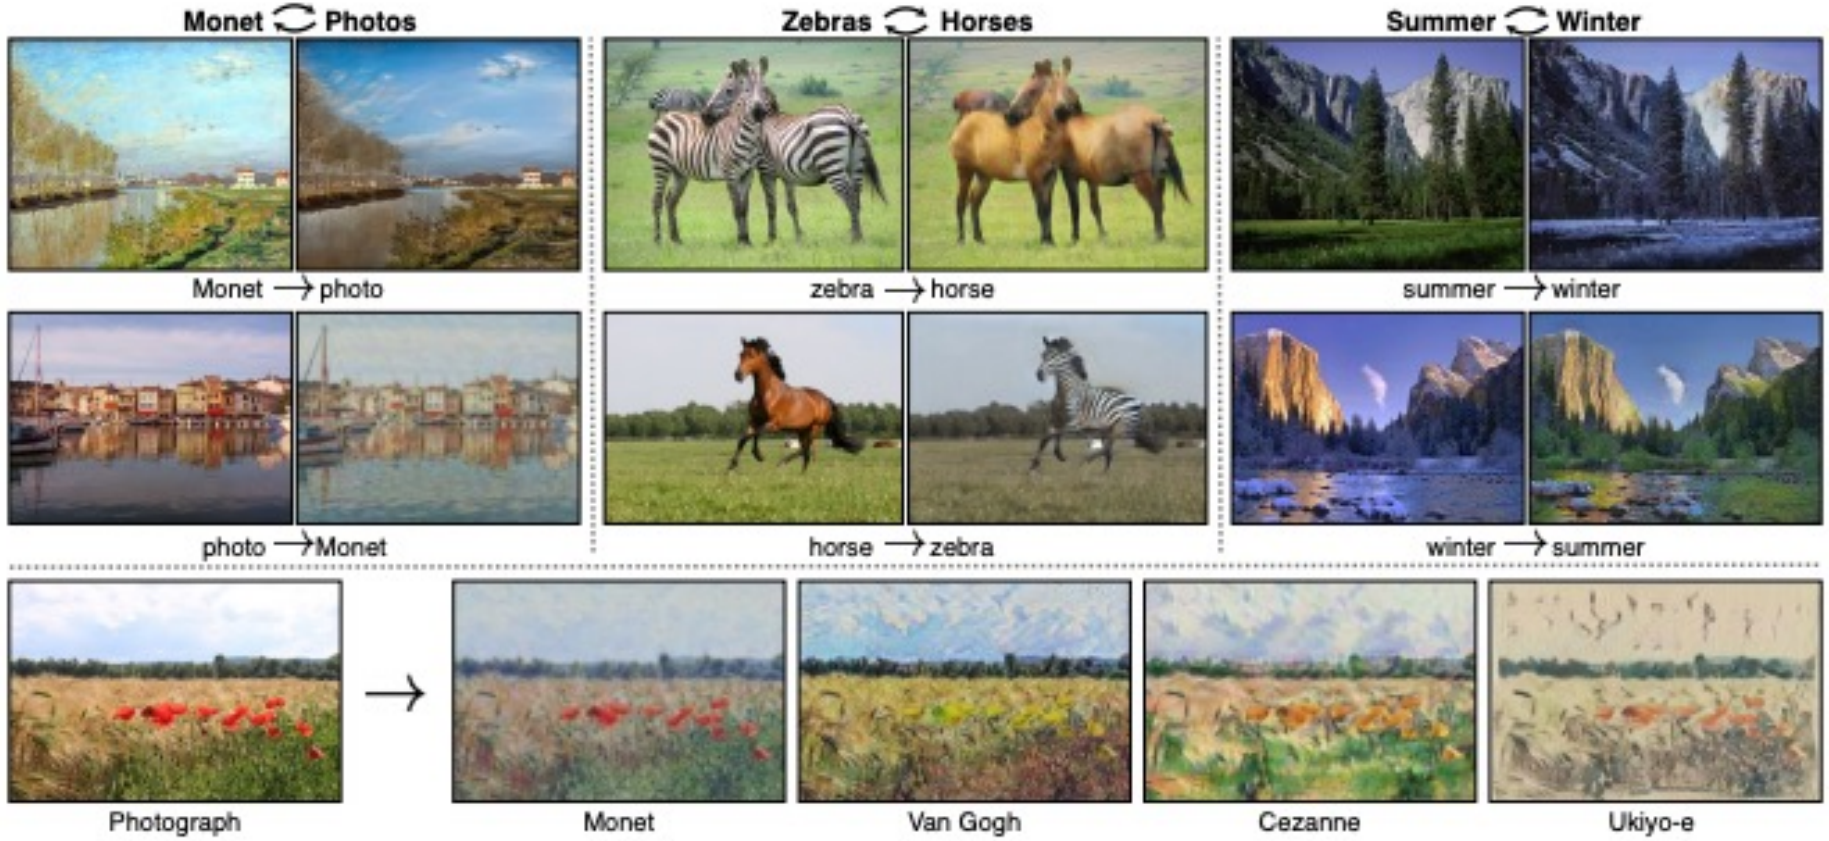

Q1: content changed?

CycleGAN

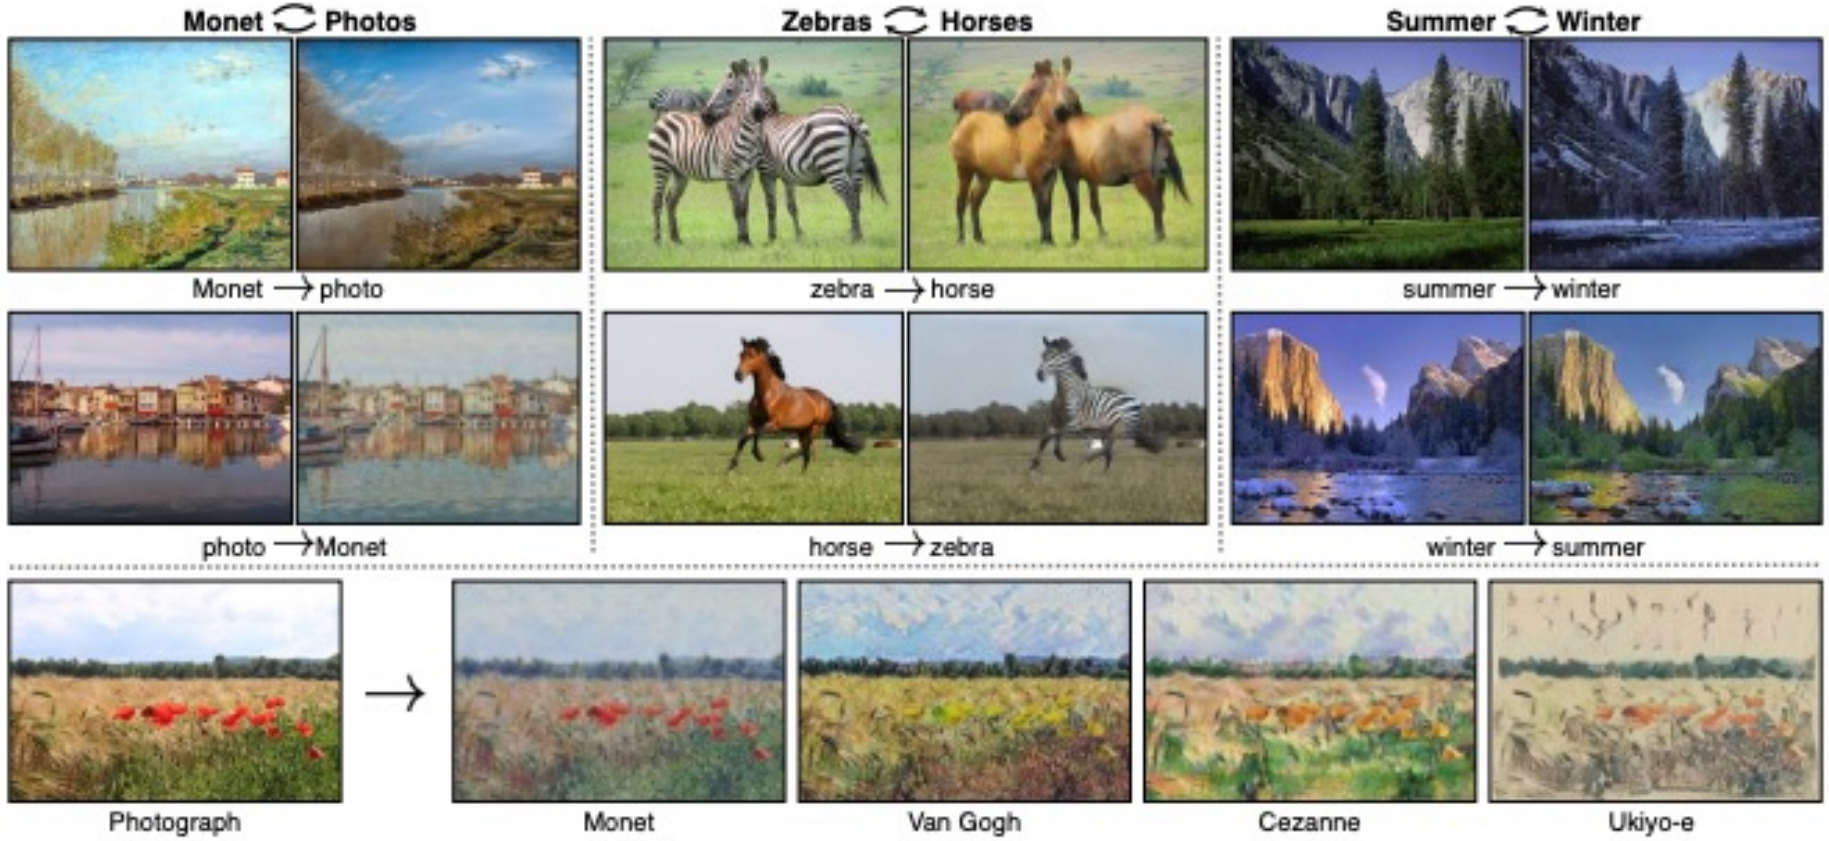

Q1: content changed?
Q2: what changed?

CycleGAN

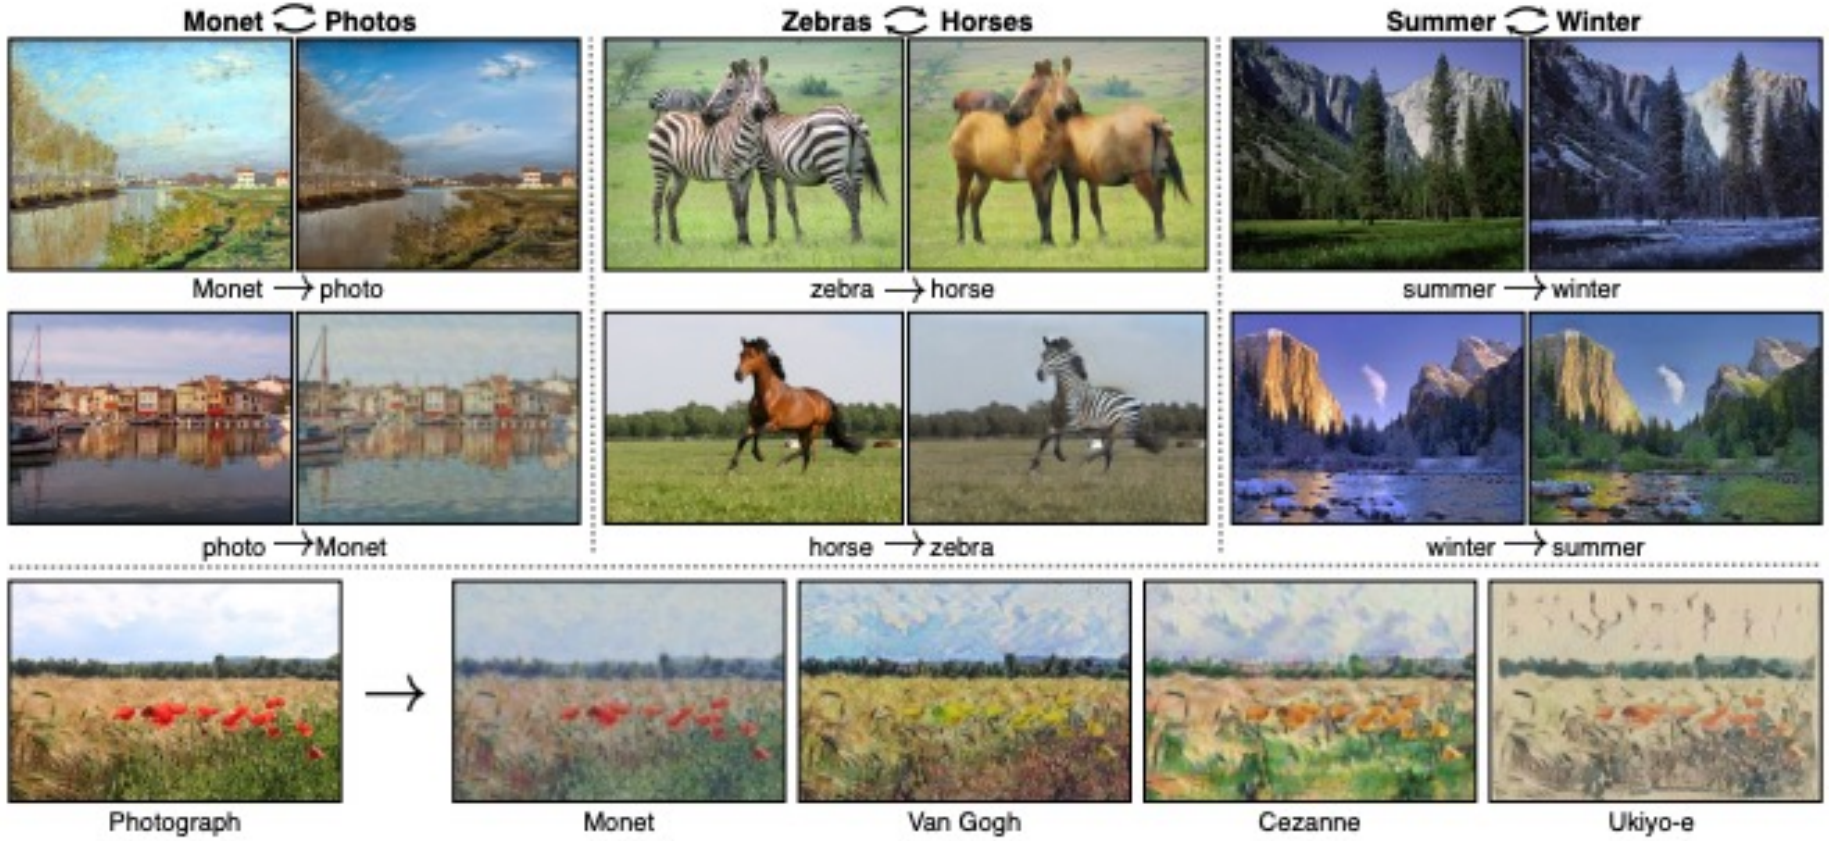

Q1: content changed?

Q2: what changed?→image-to-image translation

CycleGAN

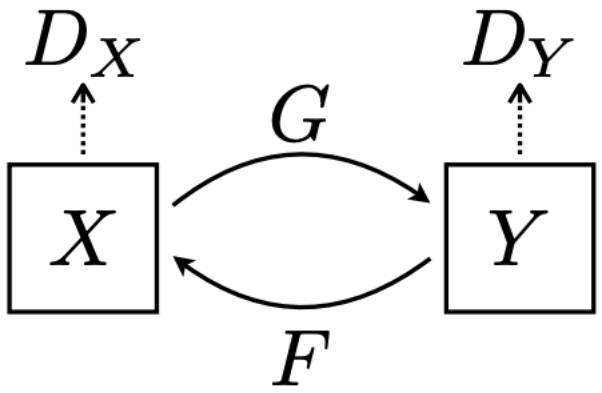

CycleGAN

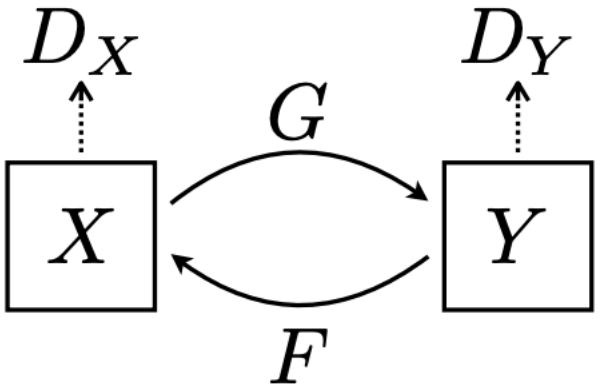

CycleGAN

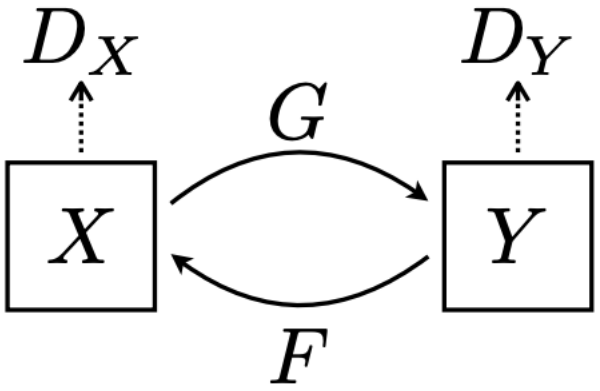

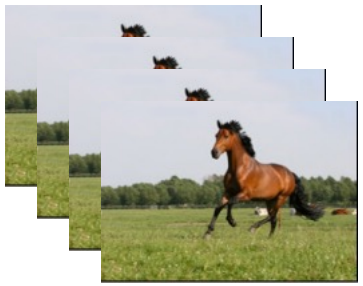

CycleGAN

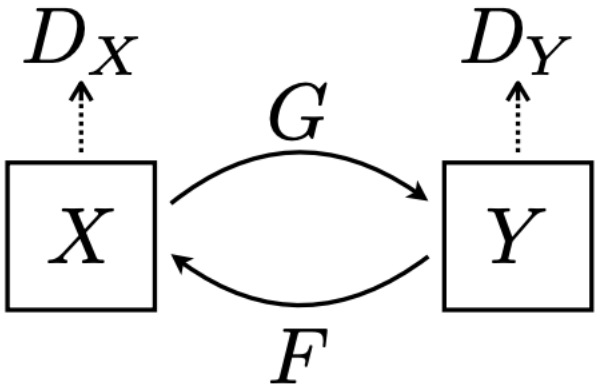

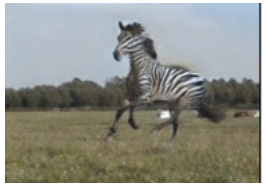

CycleGAN

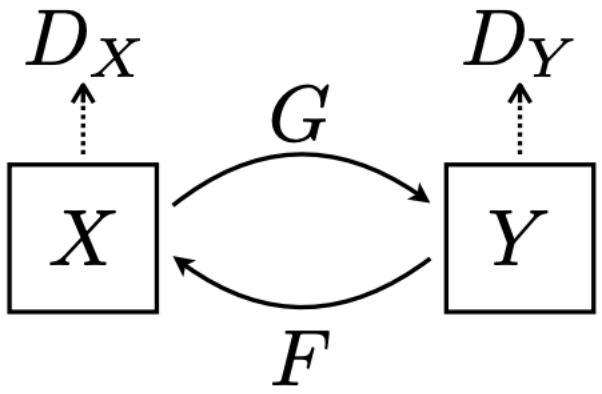

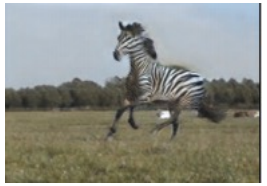

G

CycleGAN

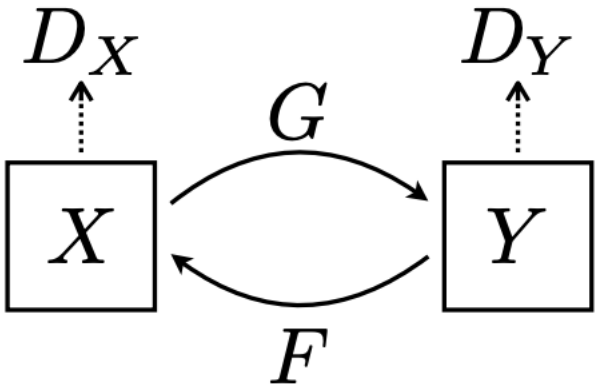

G

F

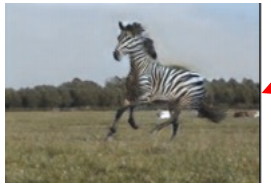

CycleGAN

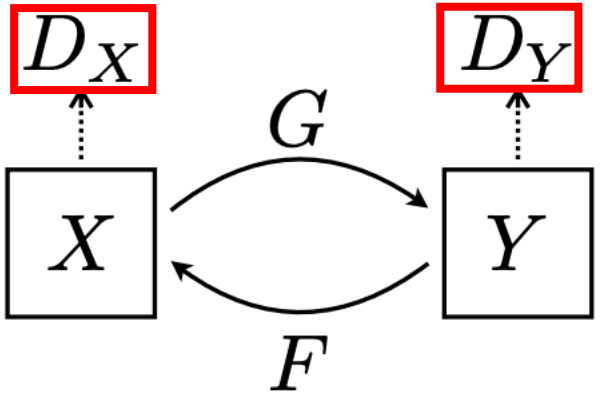

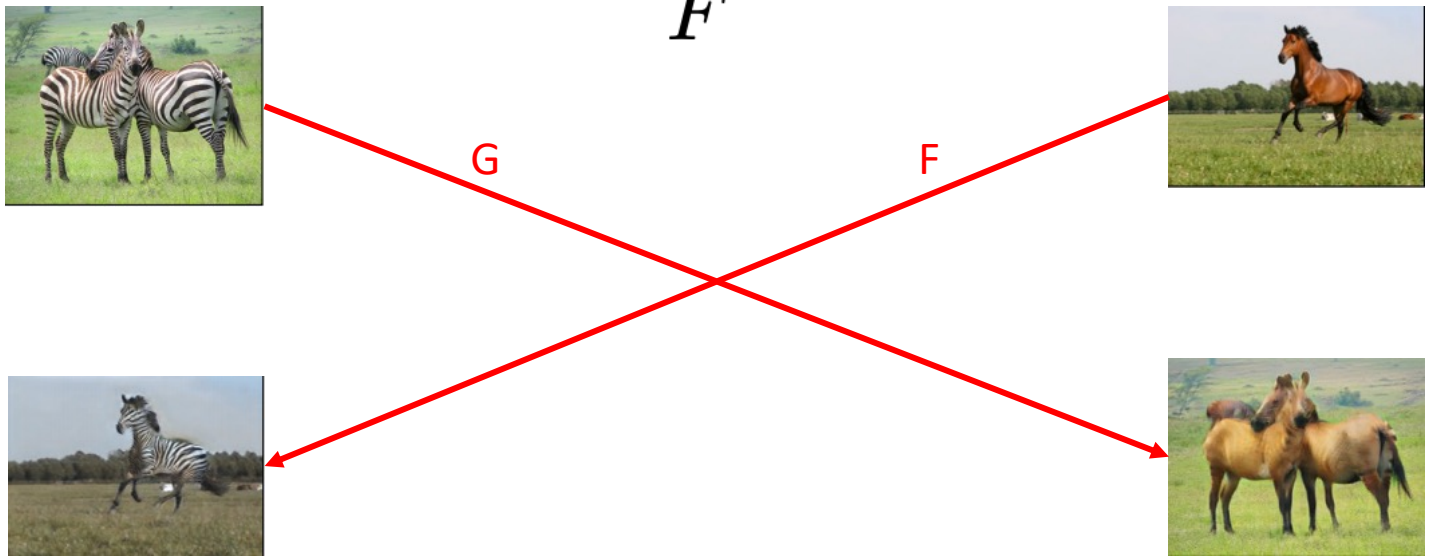

CycleGAN

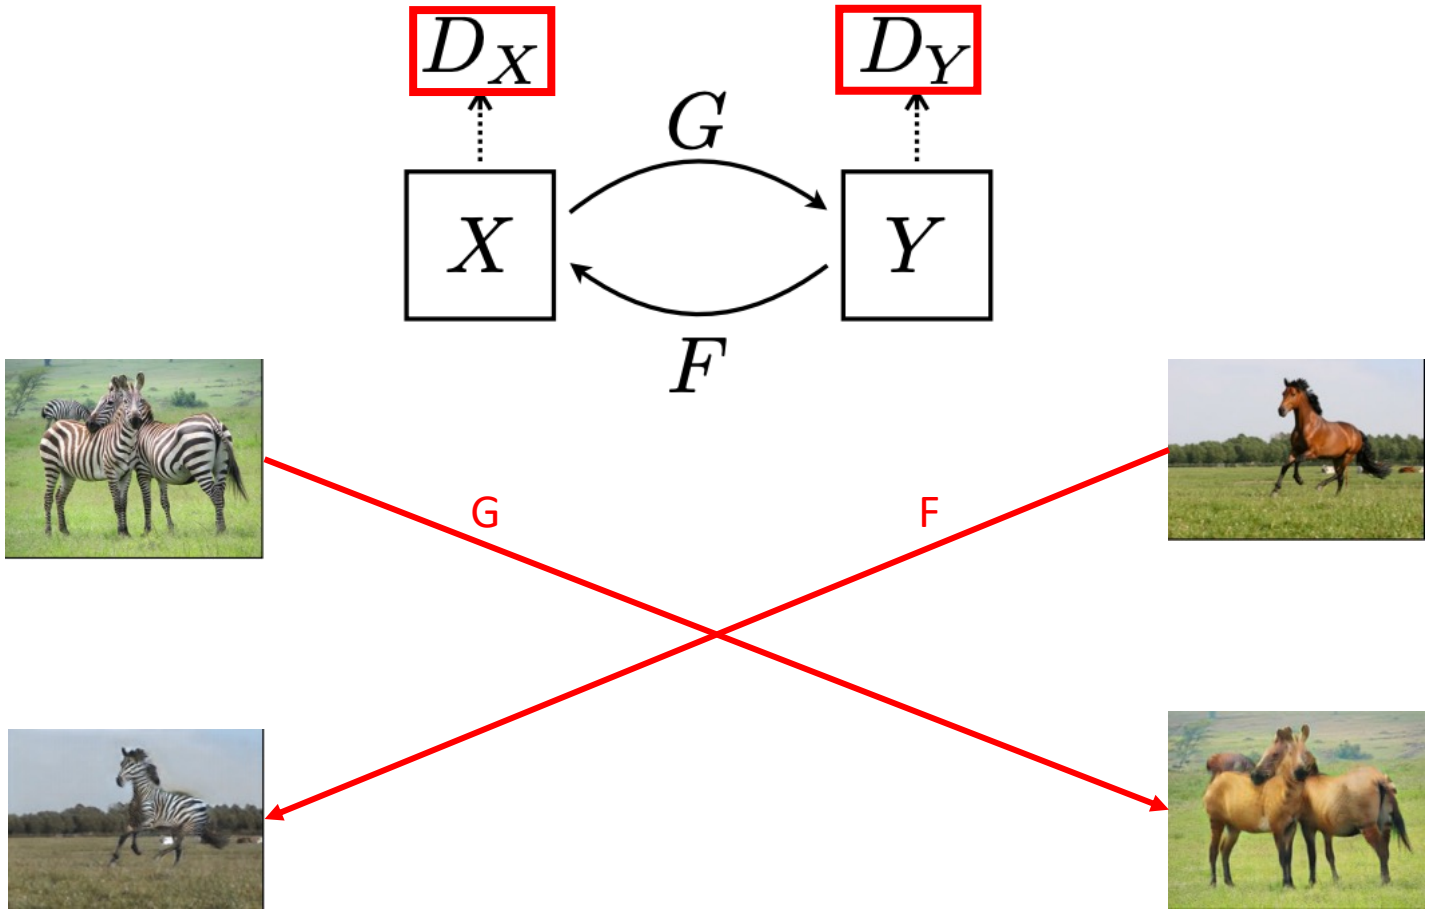

Q: how to ensure the content is not changed?

CycleGAN

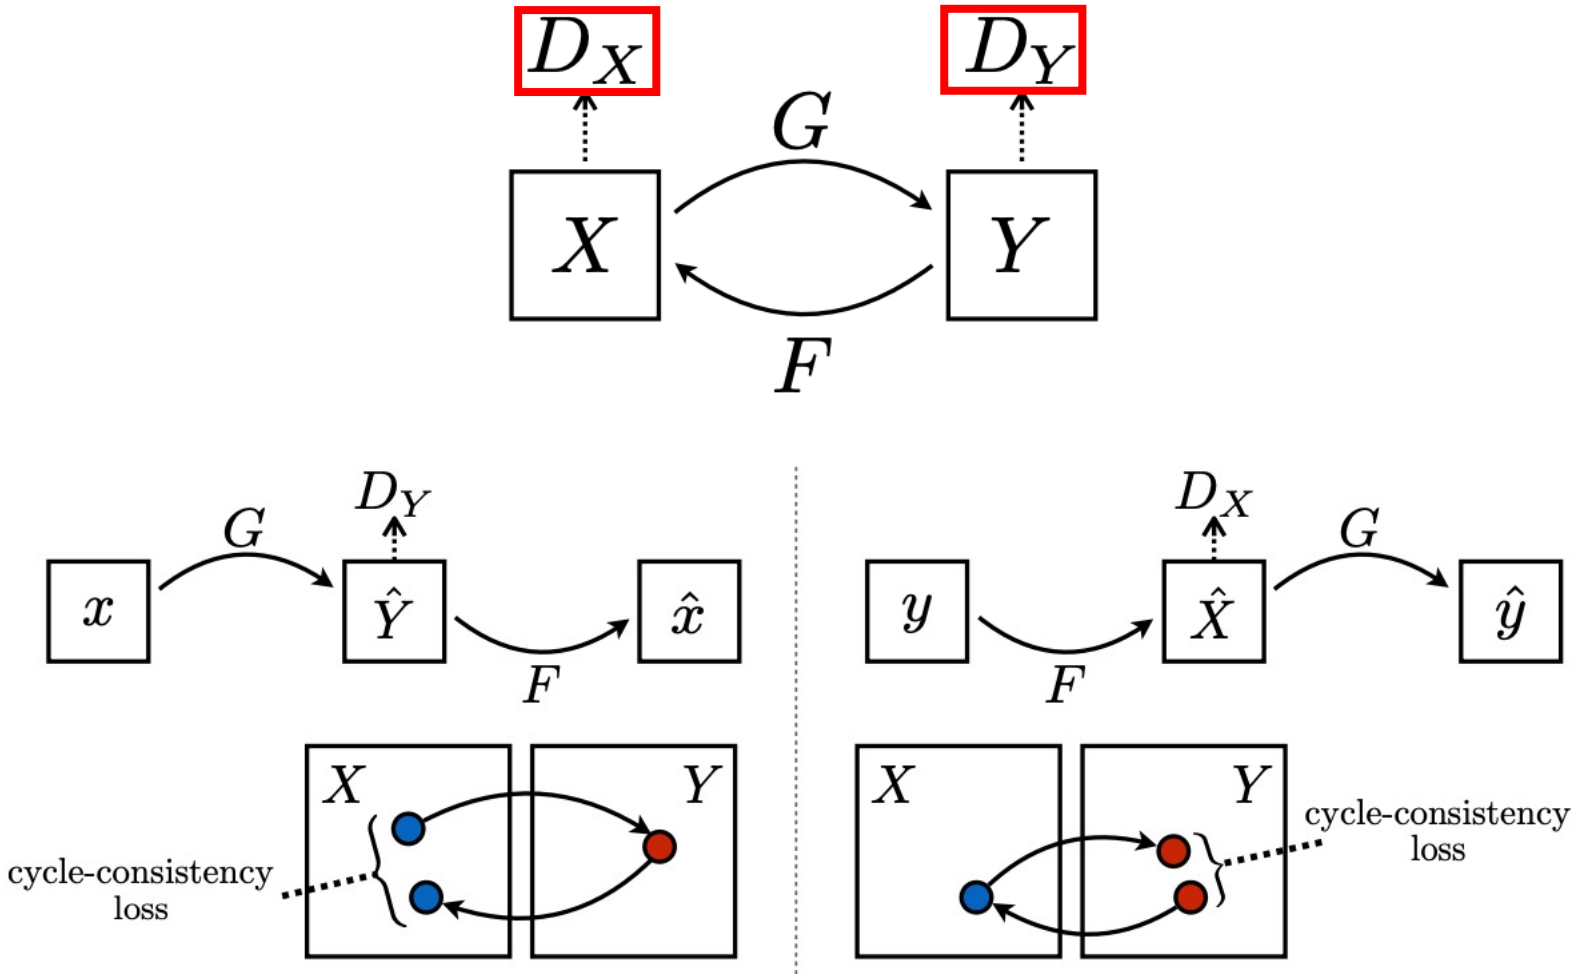
D_X
D_Y
G
X
Y
F
D_Y
G
x
$\hat{Y}$
$\hat{x}$
F
D_X
G
y
$\hat{X}$
$\hat{y}$
F
X
Y
cycle-consistency loss
X
Y
cycle-consistency loss

CycleGAN

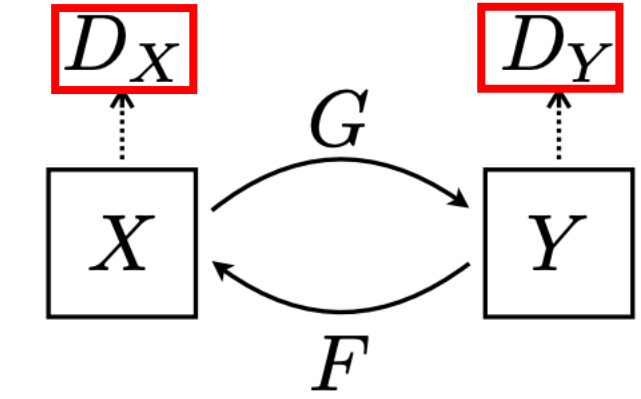
D_X
D_Y
G
X
Y
F

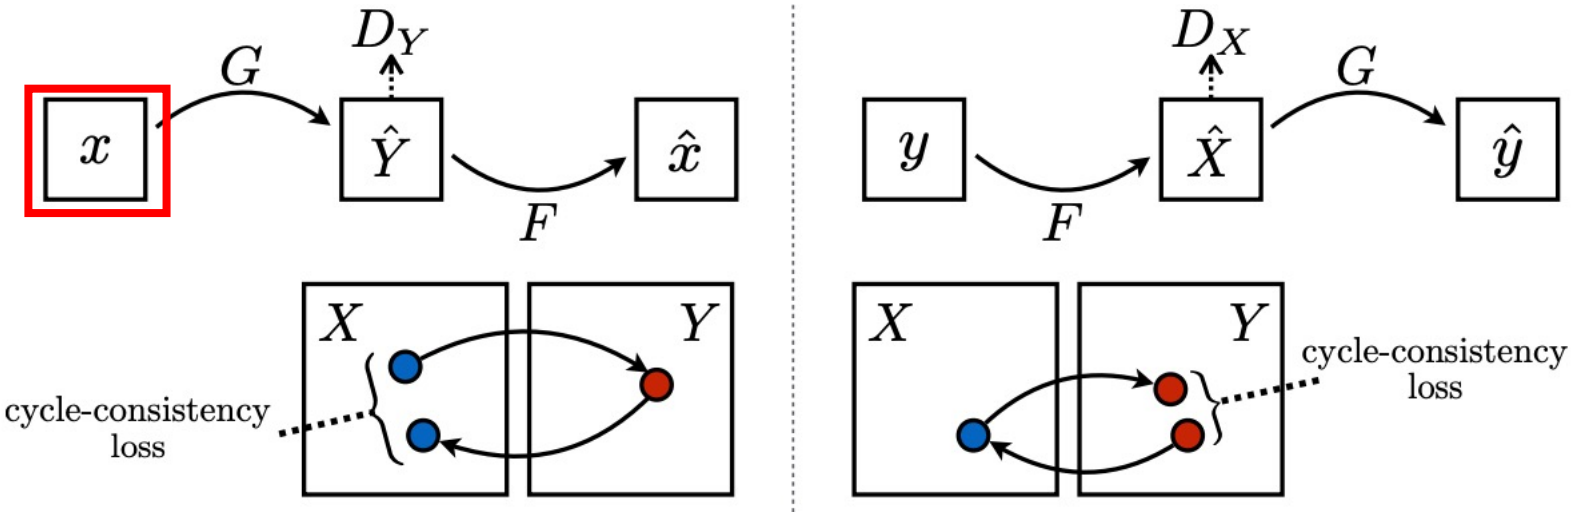
D_Y
G
x
$\hat{Y}$
$\hat{x}$
F
X
Y
cycle-consistency loss
D_X
G
y
$\hat{X}$
$\hat{y}$
F
X
Y
cycle-consistency loss

CycleGAN

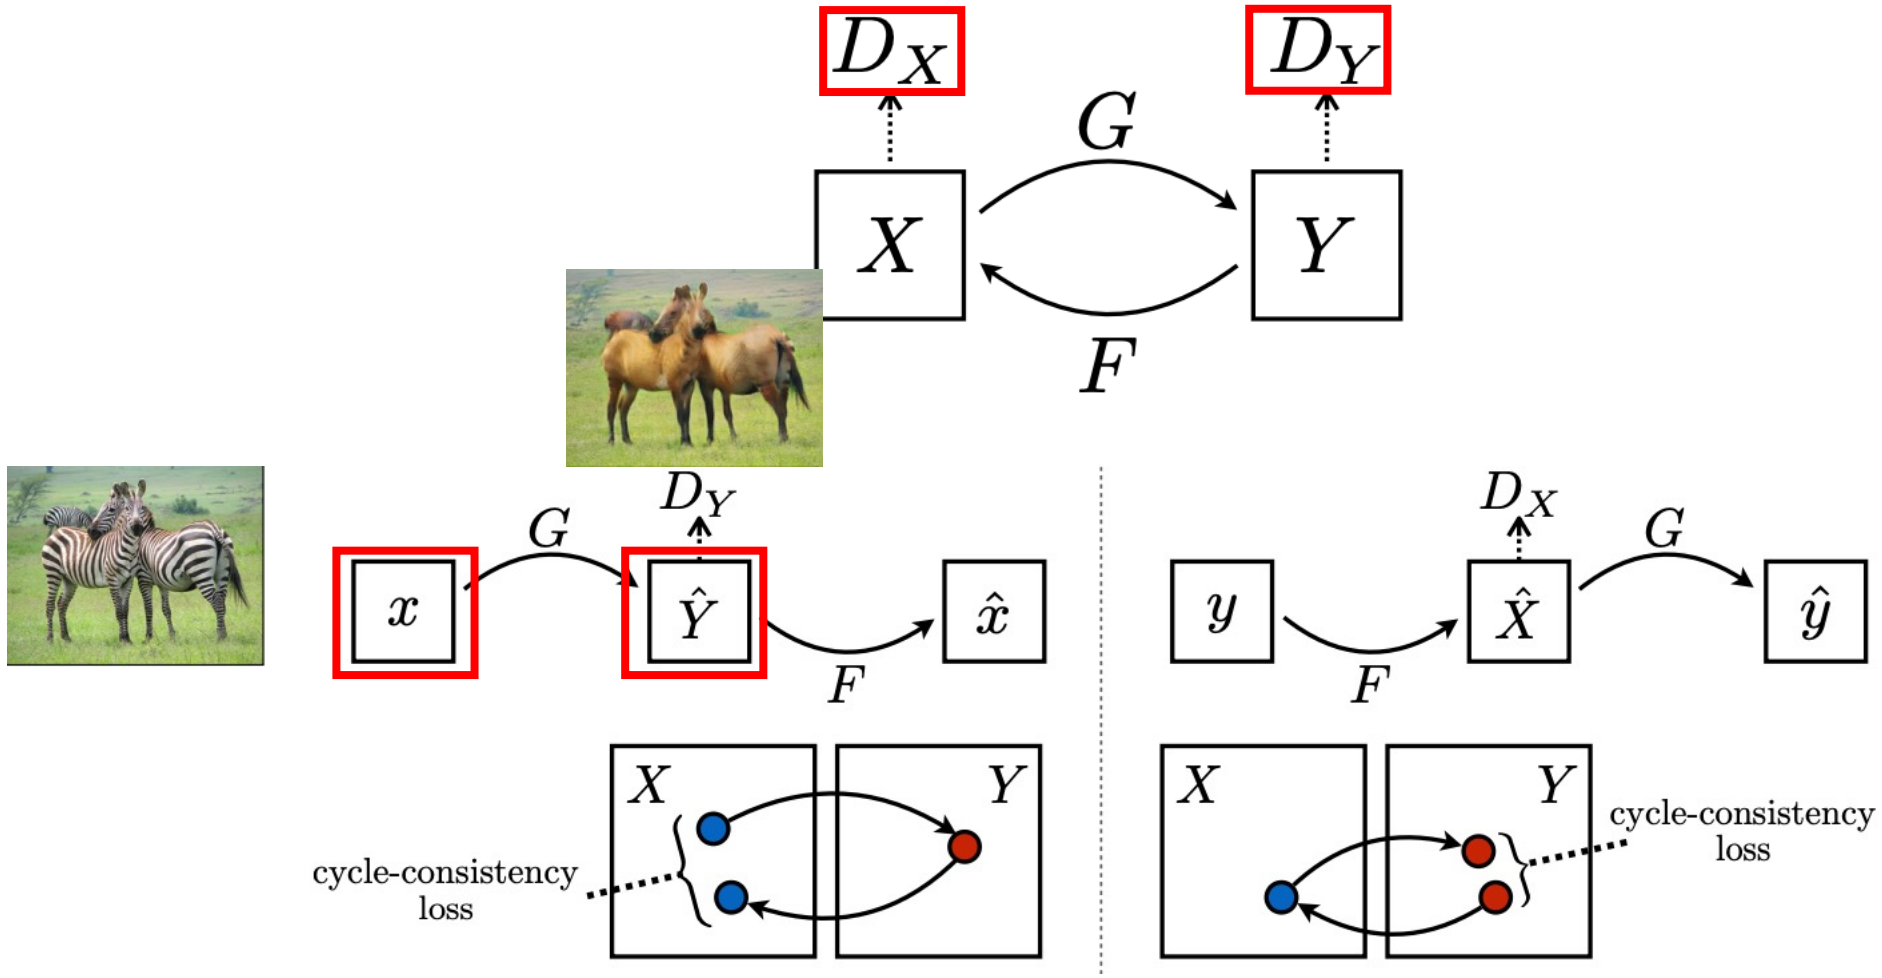
D_X
D_Y
G
X
Y
F
D_Y
G
x
$\hat{Y}$
$\hat{x}$
F
D_X
G
y
$\hat{X}$
$\hat{y}$
F
X
Y
cycle-consistency loss
X
Y
cycle-consistency loss

CycleGAN

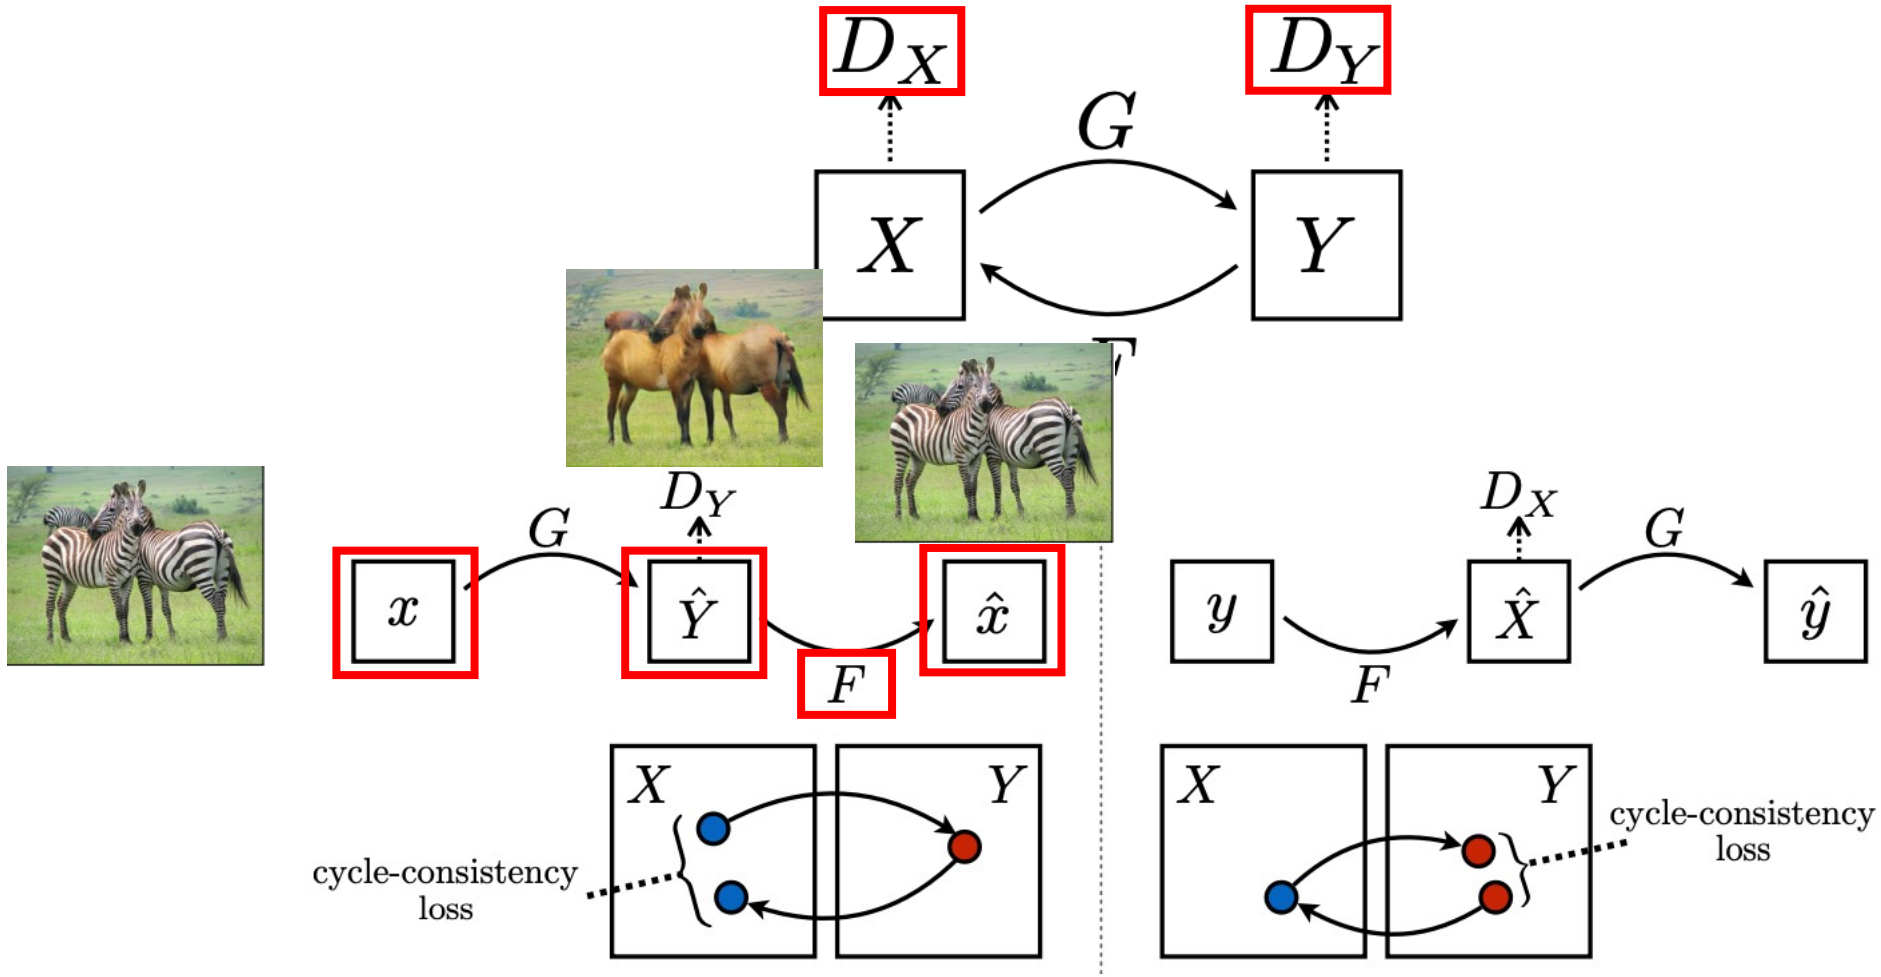
D_X
D_Y
G
X
Y
F
D_Y
G
x
$\hat{Y}$
$\hat{x}$
F
D_X
G
y
$\hat{X}$
$\hat{y}$
F
X
Y
cycle-consistency loss
X
Y
cycle-consistency loss

CycleGAN

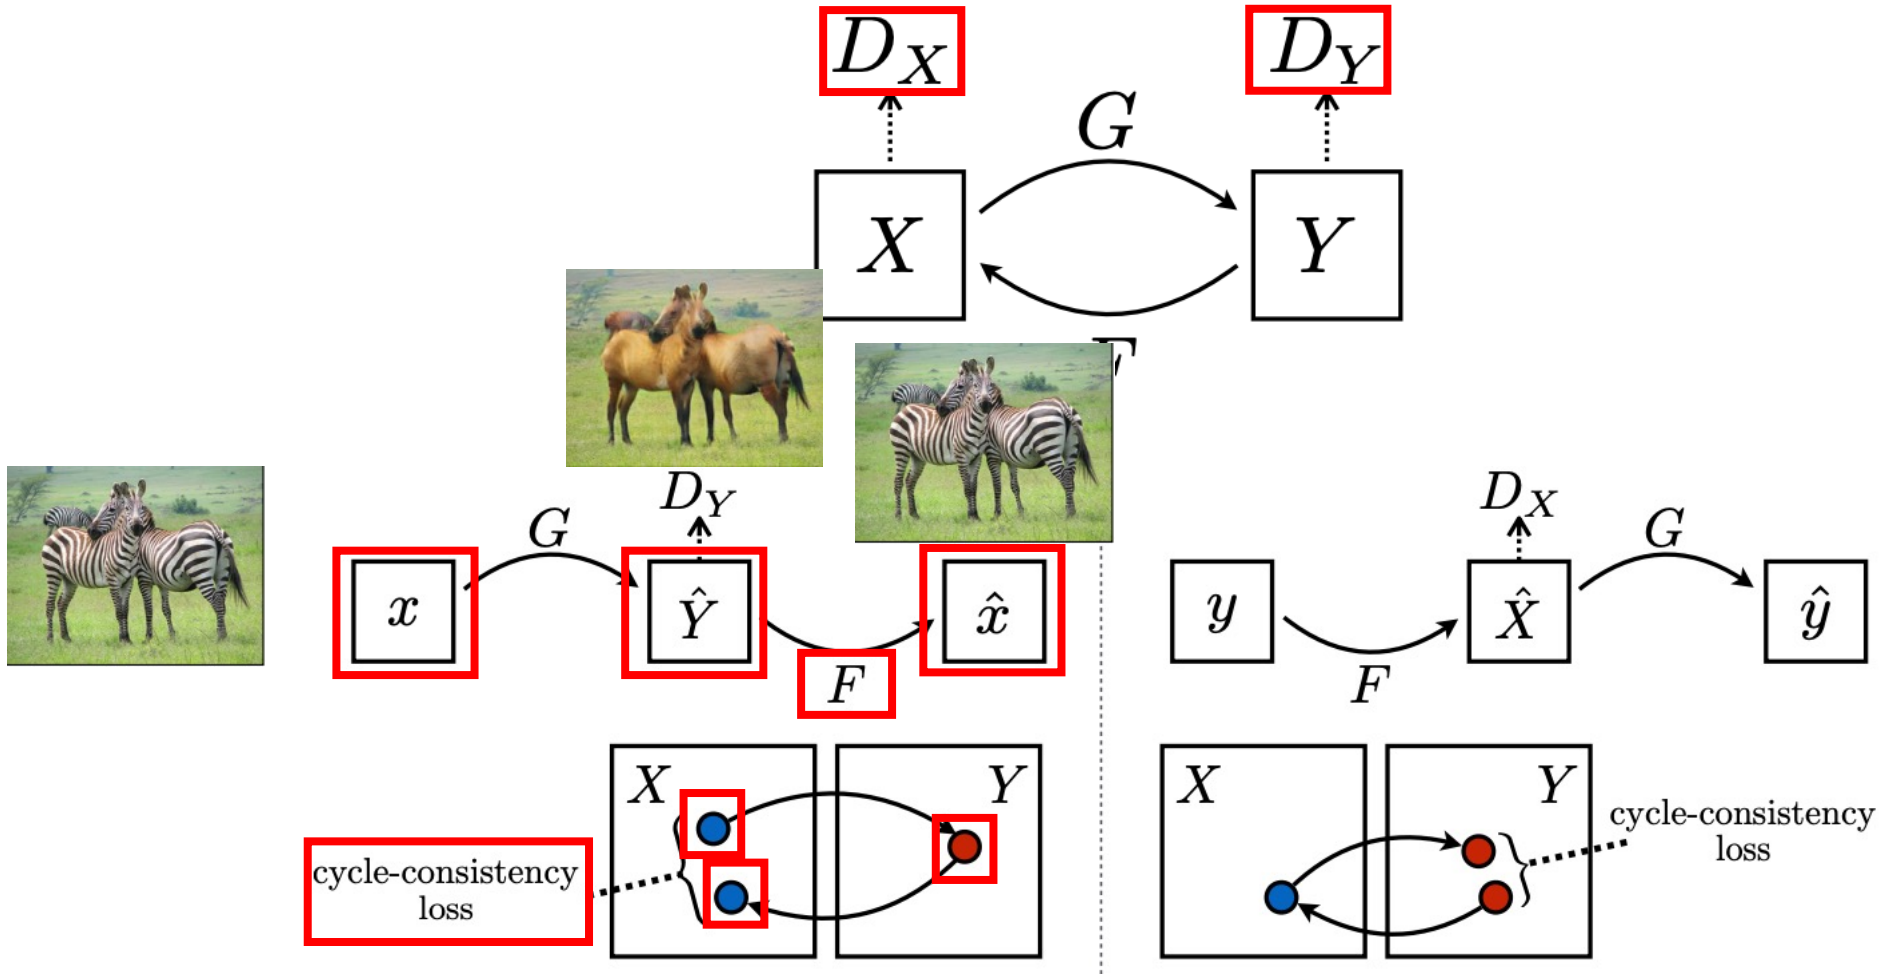
D_X
D_Y
G
X
Y
F
D_Y
G
x
$\hat{Y}$
$\hat{x}$
F
D_X
G
y
$\hat{X}$
$\hat{y}$
F
X
Y
cycle-consistency loss
X
Y
cycle-consistency loss

CycleGAN

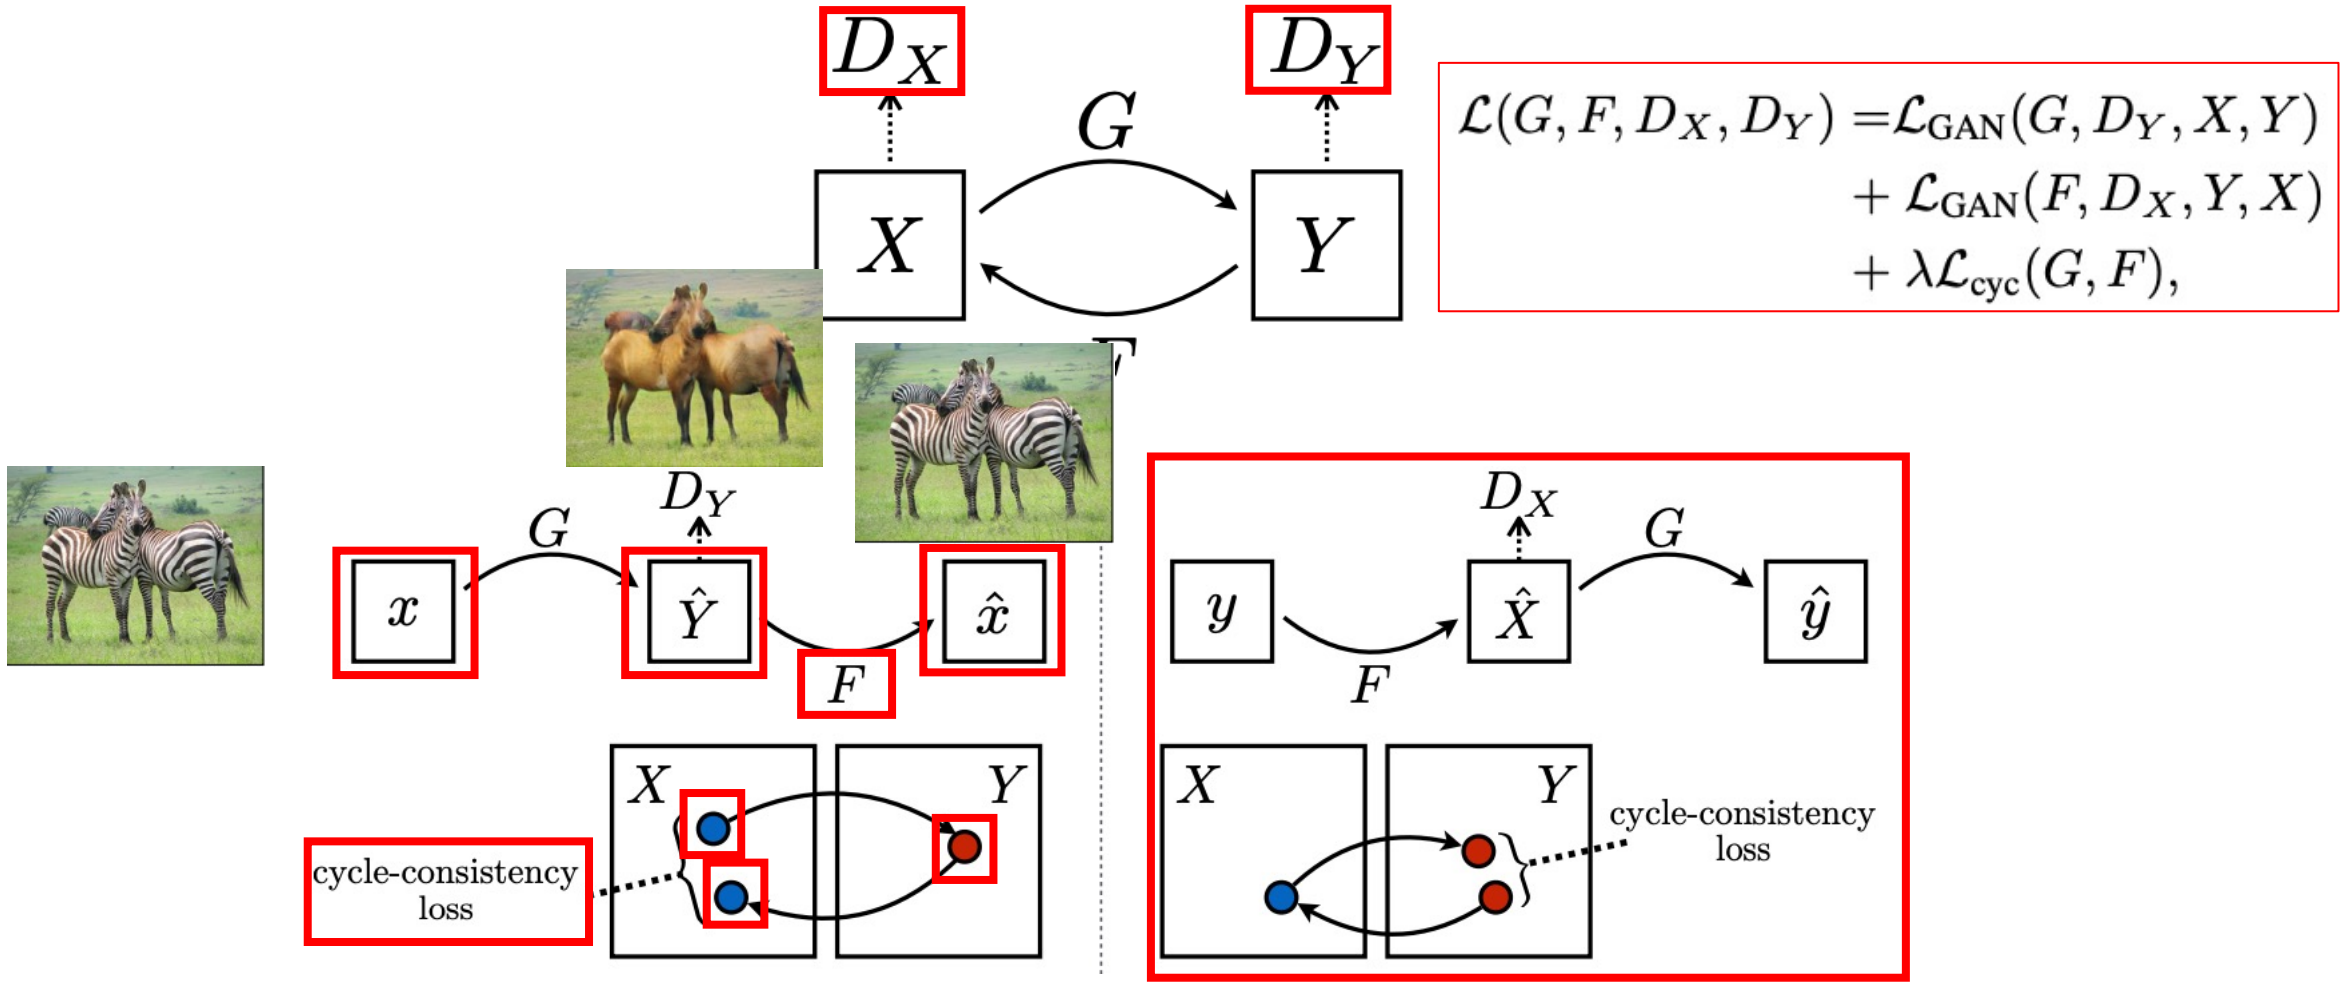
D_X
D_Y
G
X
Y
F
$\mathcal{L}(G, F, D_X, D_Y) = \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y)$
$+ \mathcal{L}_{\text{GAN}}(F, D_X, Y, X)$
$+ \lambda \mathcal{L}_{\text{cyc}}(G, F),$
x
$\hat{Y}$
$\hat{x}$
y
$\hat{X}$
$\hat{y}$
cycle-consistency loss
cycle-consistency loss

StyleGAN

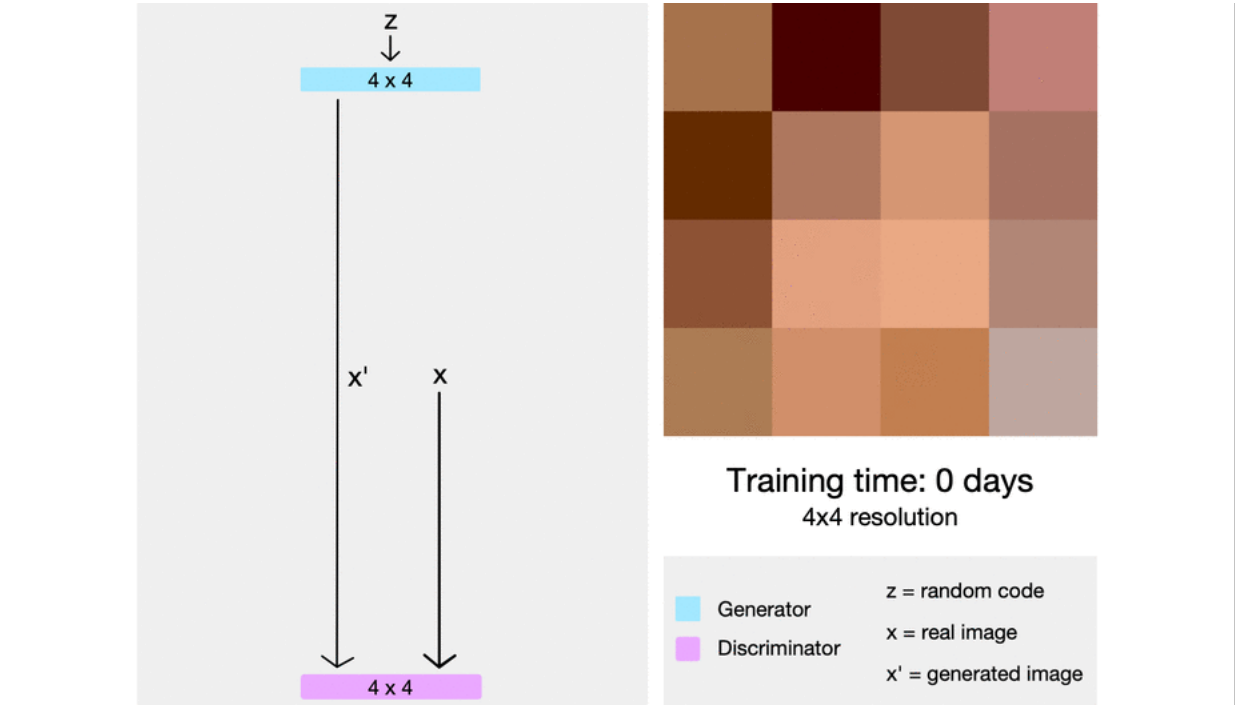

Progressive GAN https://arxiv.org/pdf/1710.10196.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

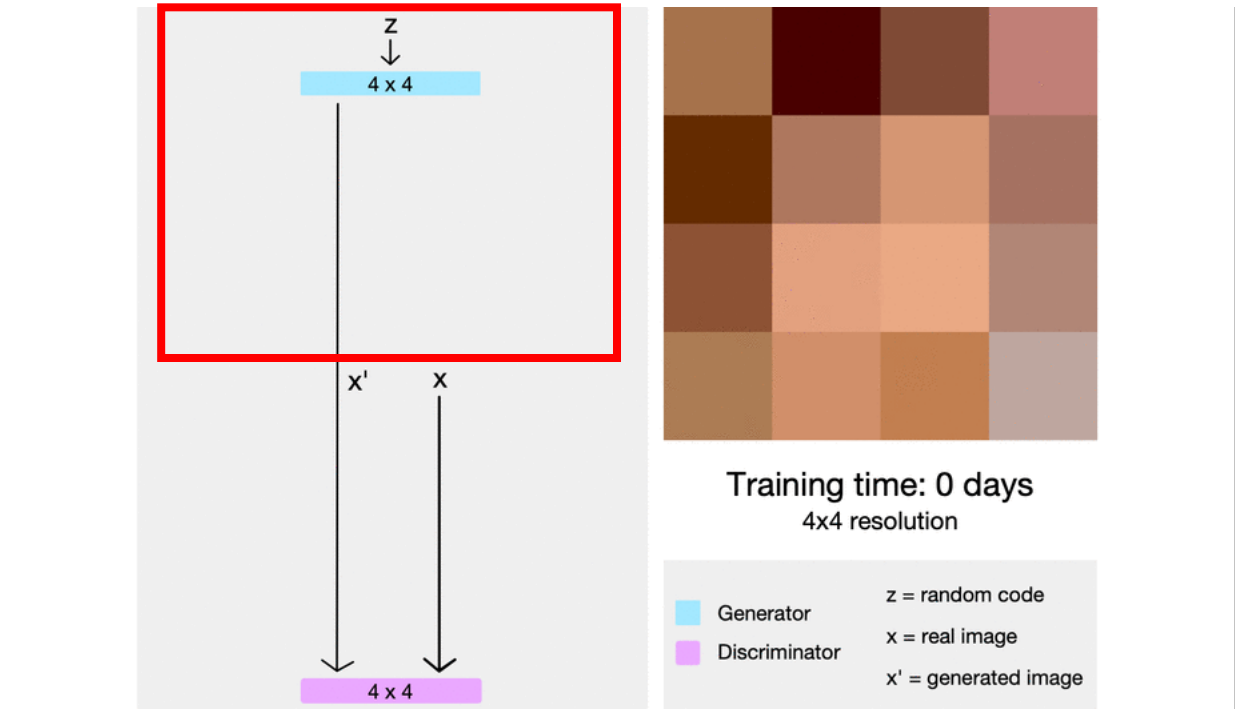

Progressive GAN https://arxiv.org/pdf/1710.10196.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

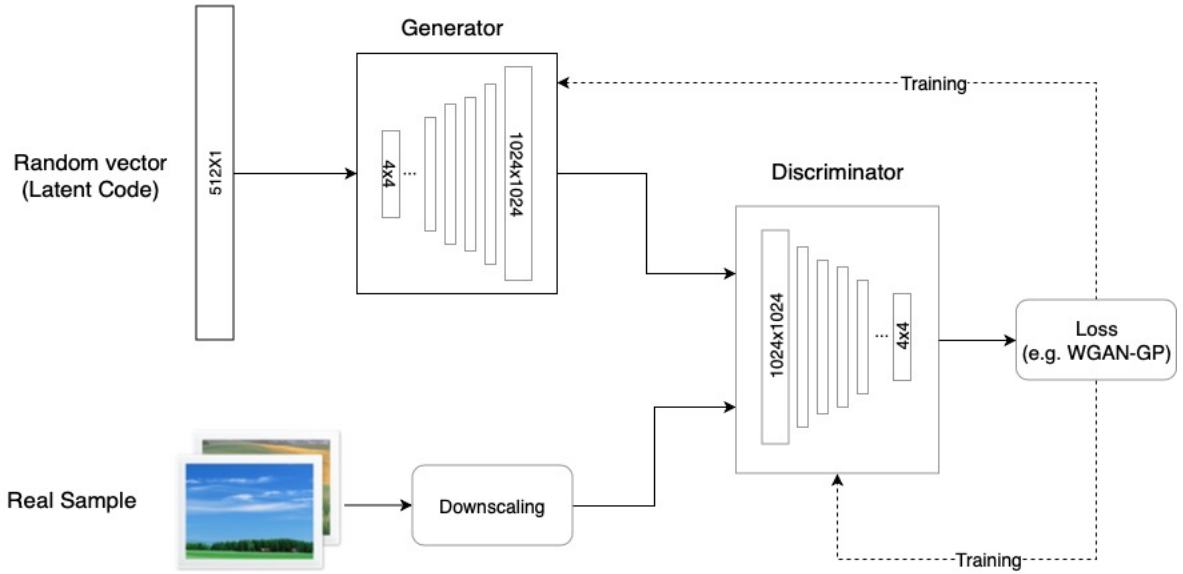

Progressive GAN https://arxiv.org/pdf/1710.10196.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

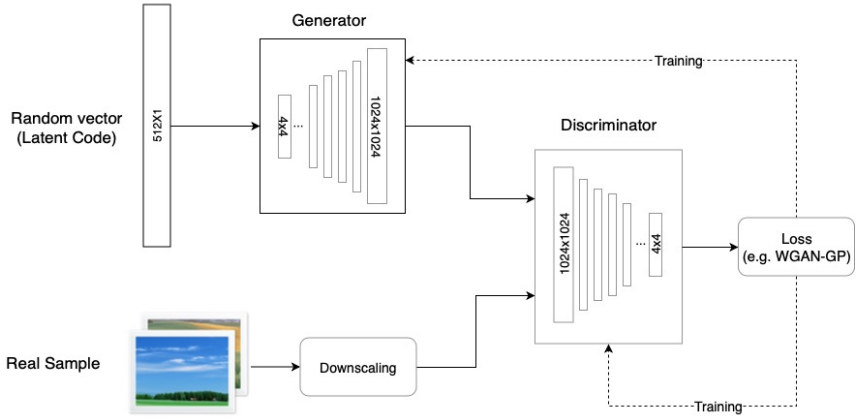

1. Coarse - resolution of up to 82 - affects pose, general hair style, face shape, etc
2. Middle - resolution of 162 to 322 - affects finer facial features, hair style, eyes open/closed, etc.
3. Fine - resolution of 642 to 10242 - affects color scheme (eye, hair and skin) and micro features.

Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019. https://arxiv.org/pdf/1812.04948.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

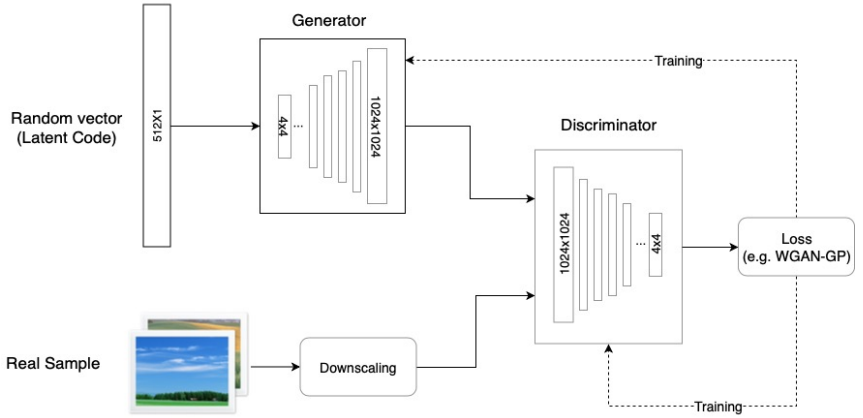

1. Coarse - resolution of up to 82 - affects pose, general hair style, face shape, etc
2. Middle - resolution of 162 to 322 - affects finer facial features, hair style, eyes open/closed, etc.
3. Fine - resolution of 642 to 10242 - affects color scheme (eye, hair and skin) and micro features.

Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019. https://arxiv.org/pdf/1812.04948.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

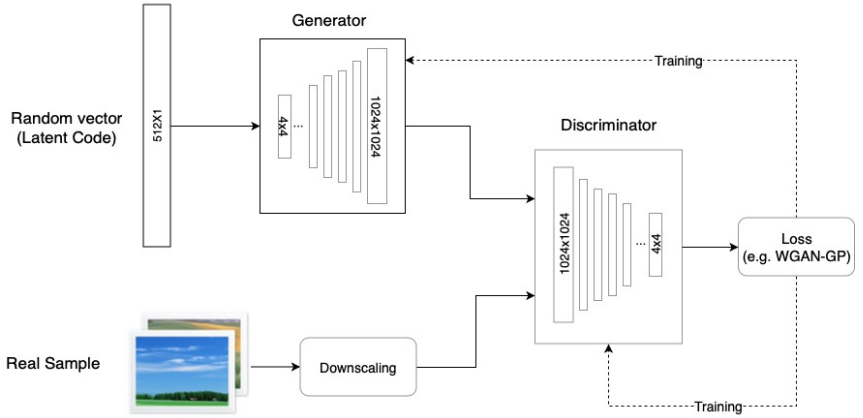

1. Coarse - resolution of up to 82 - affects pose, general hair style, face shape, etc
2. Middle - resolution of 162 to 322 - affects finer facial features, hair style, eyes open/closed, etc.
3. Fine - resolution of 642 to 10242 - affects color scheme (eye, hair and skin) and micro features.

Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019. https://arxiv.org/pdf/1812.04948.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

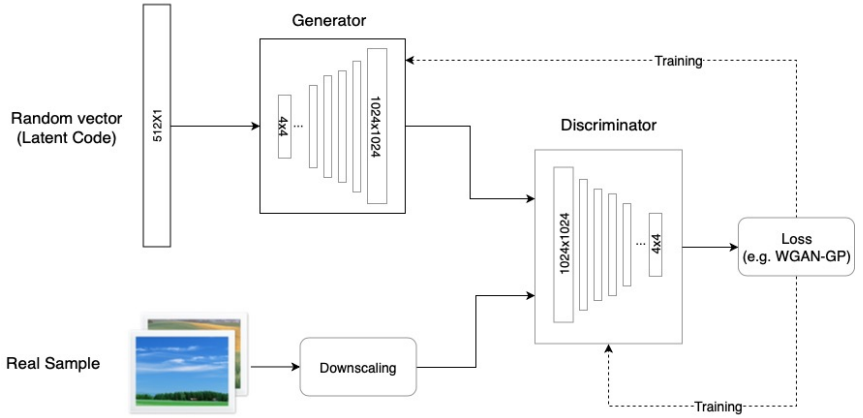

Q: can we control these styles?

1. Coarse - resolution of up to 82 - affects pose, general hair style, face shape, etc
2. Middle - resolution of 162 to 322 - affects finer facial features, hair style, eyes open/closed, etc.
3. Fine - resolution of 642 to 10242 - affects color scheme (eye, hair and skin) and micro features.

Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019. https://arxiv.org/pdf/1812.04948.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

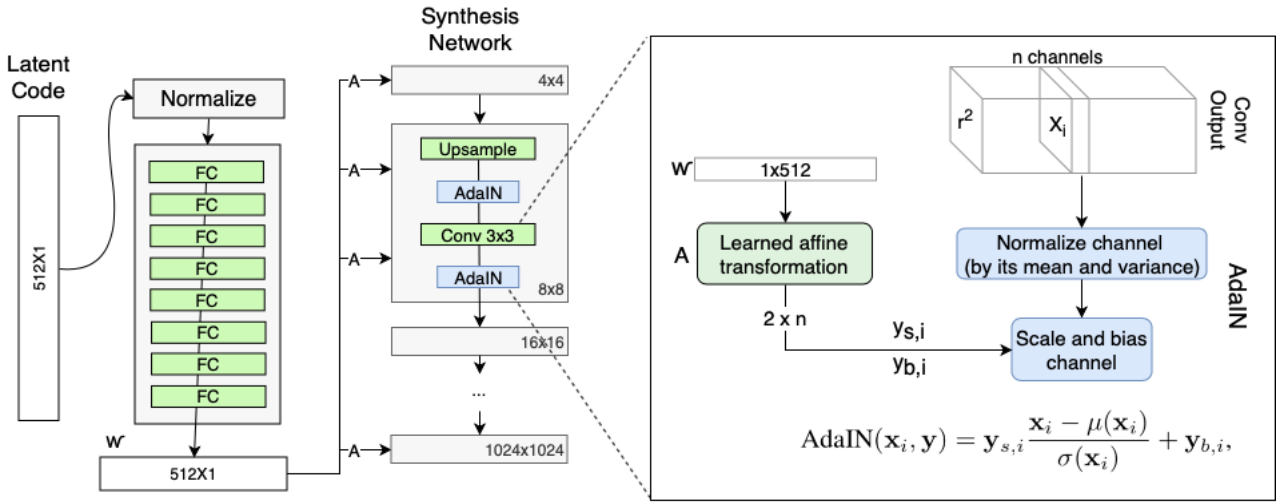

Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019. https://arxiv.org/pdf/1812.04948.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

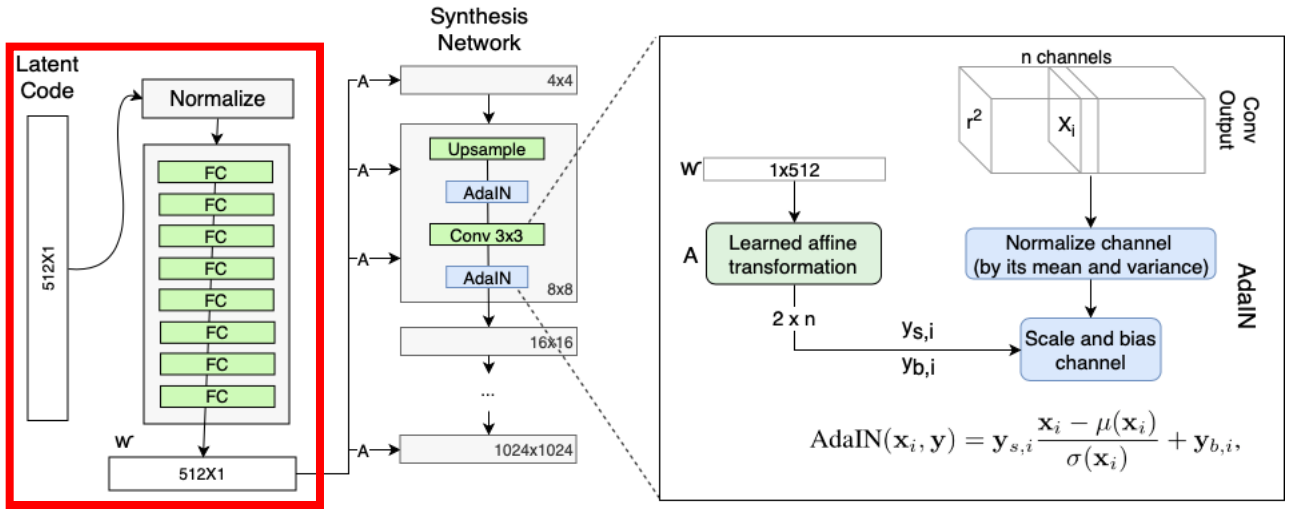

Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019. https://arxiv.org/pdf/1812.04948.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

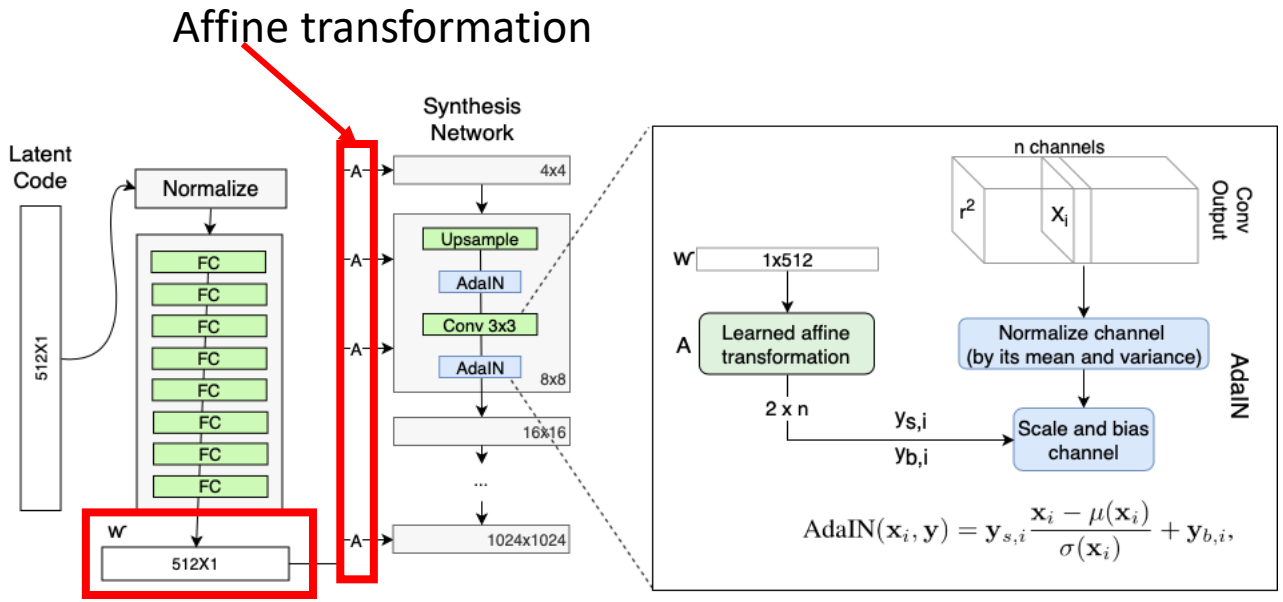

Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019. https://arxiv.org/pdf/1812.04948.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

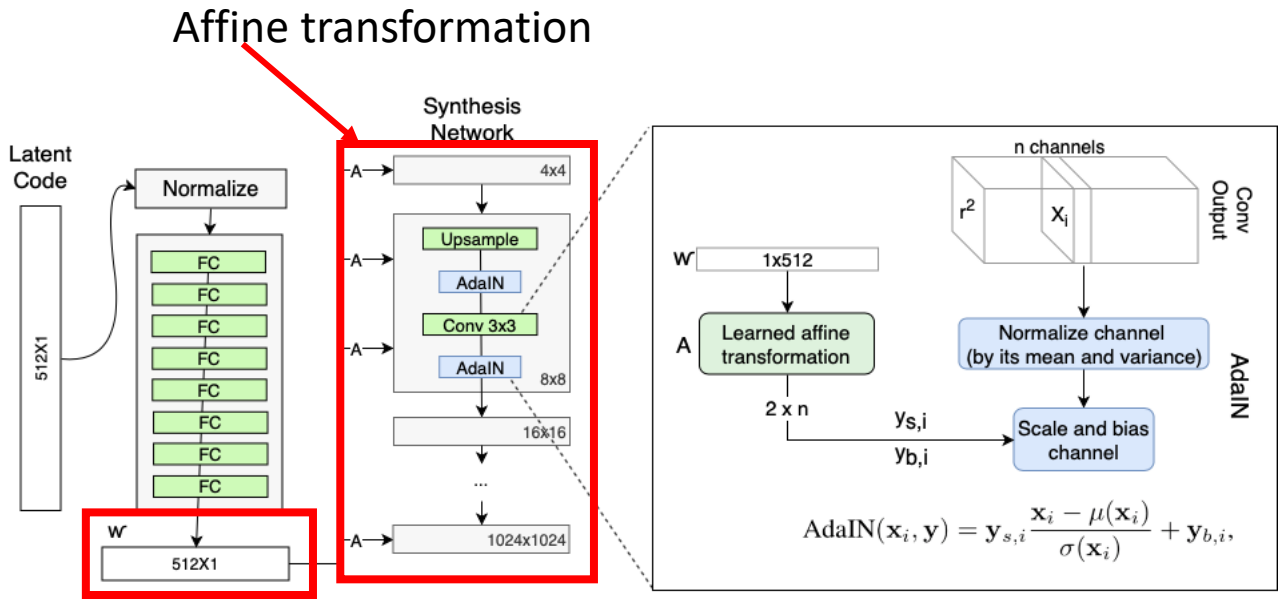

Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019. https://arxiv.org/pdf/1812.04948.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

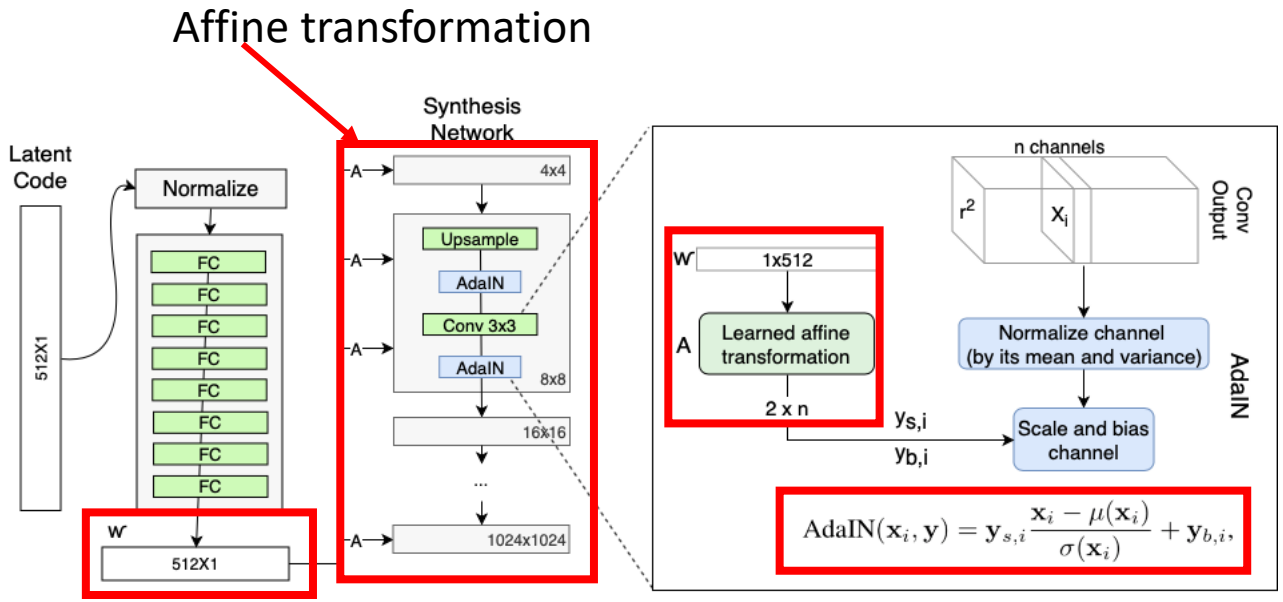

Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019. https://arxiv.org/pdf/1812.04948.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

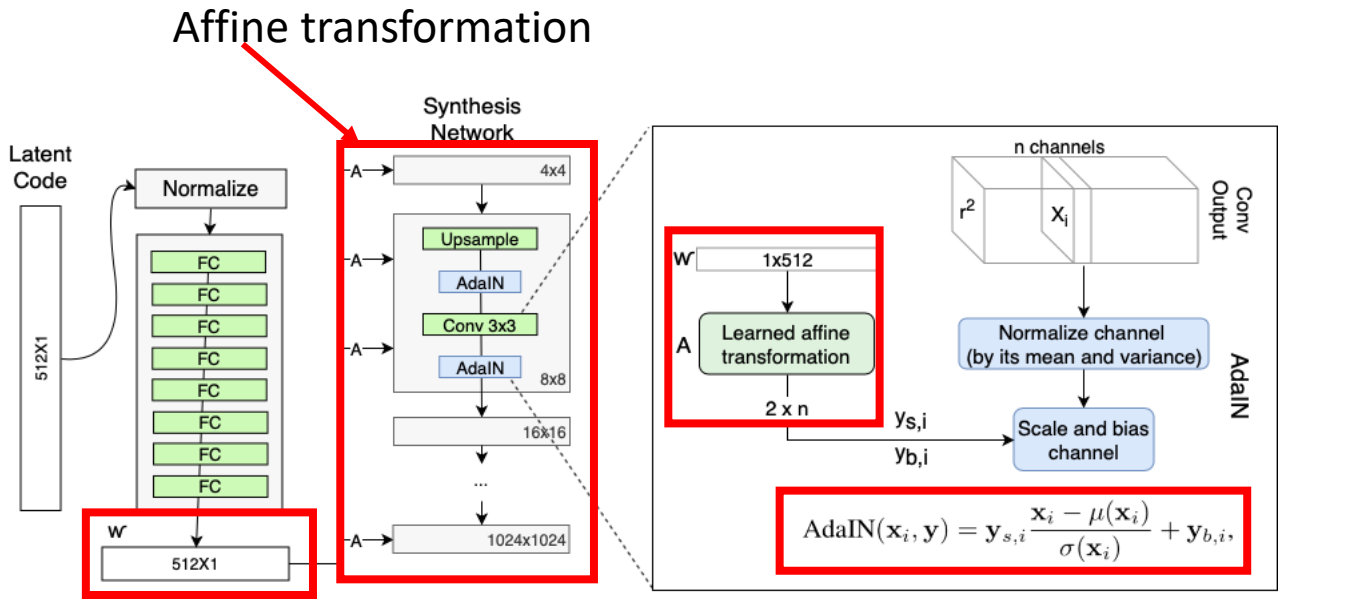

ADAptive Instance Normalization

Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019. https://arxiv.org/pdf/1812.04948.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

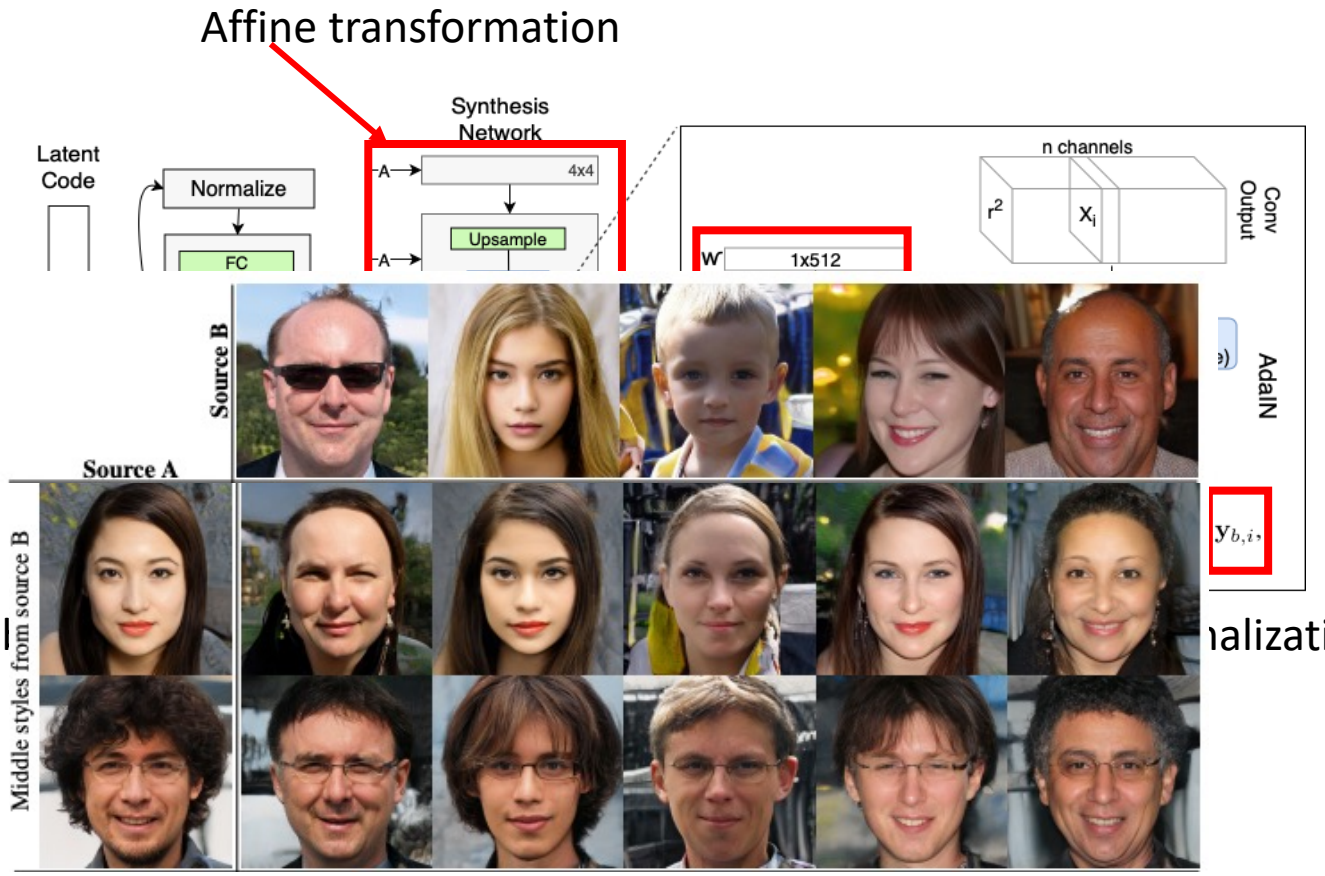

Style mixing: what al… …nalization recourses at various …

Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 4401-4410. 2019. https://arxiv.org/pdf/1812.04948.pdf

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

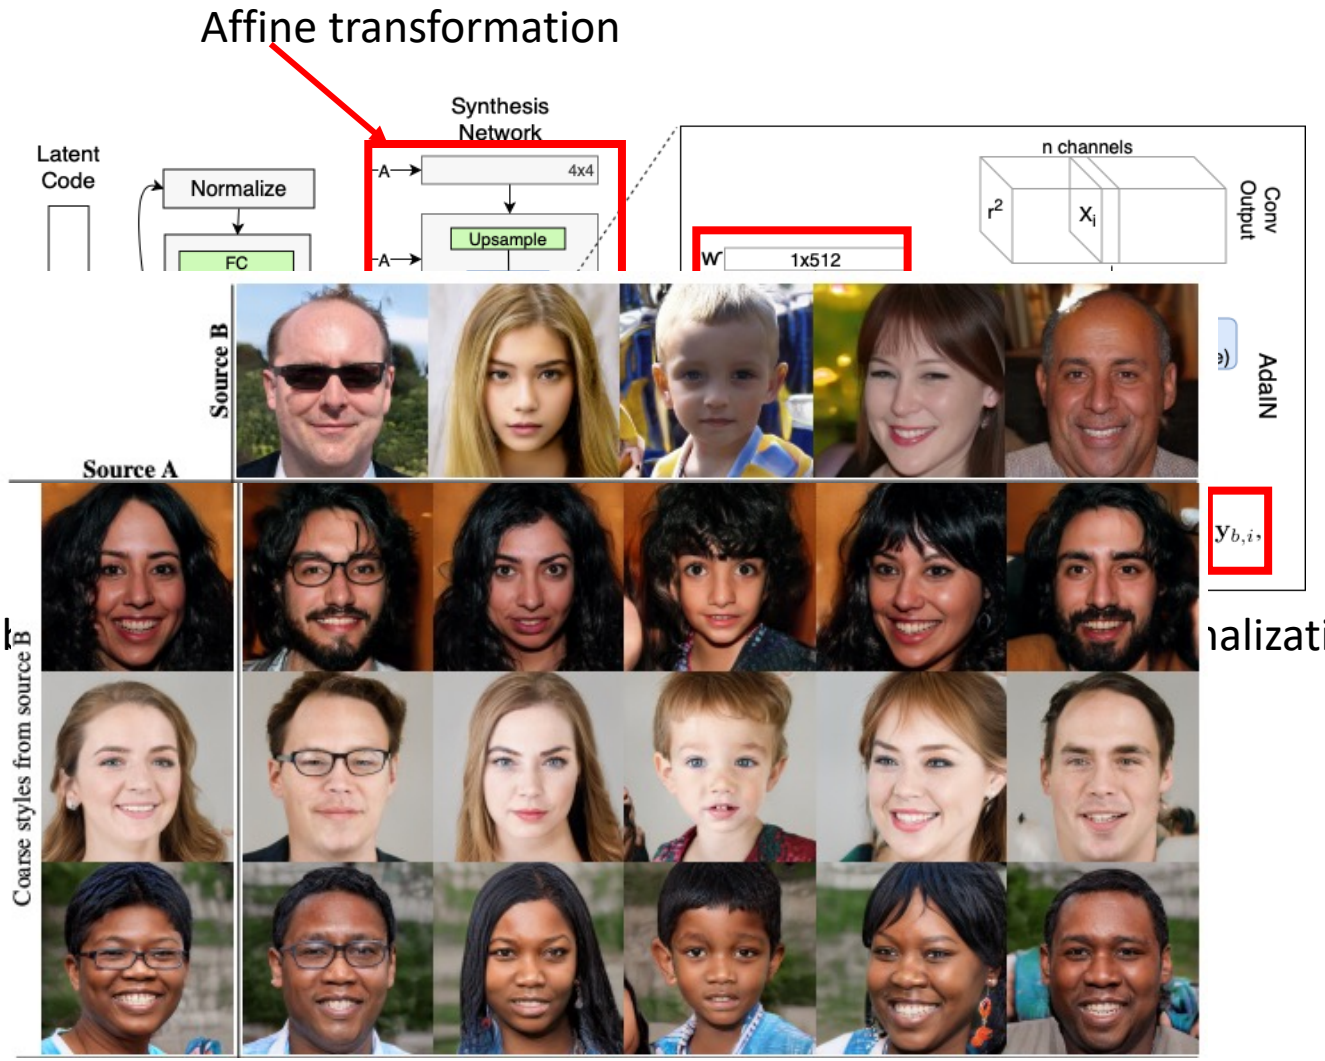

Style mixing: what al… …nalization recourses at various …

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

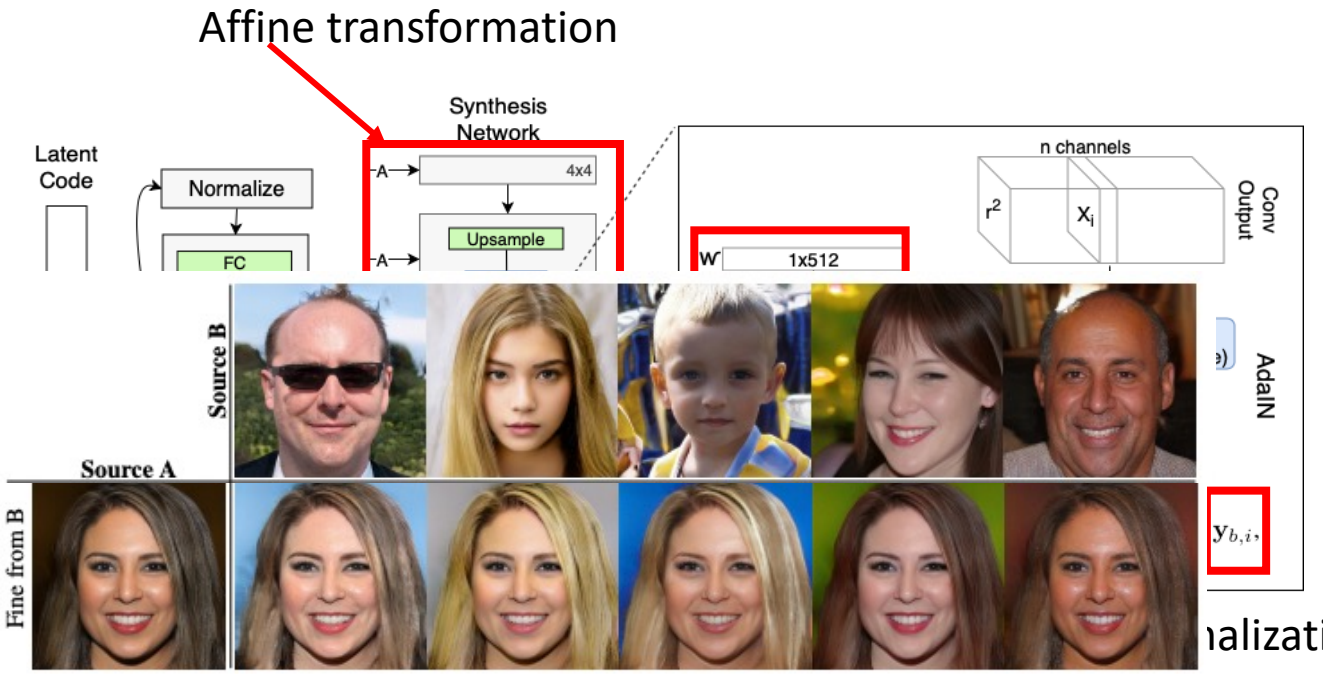

Style mixing: what a ...alization recourses at various layers?

https://towardsdatascience.com/explained-a-style-based-generator-architecture-for-gans-generating-and-tuning-realistic-6cb2be0f431

StyleGAN

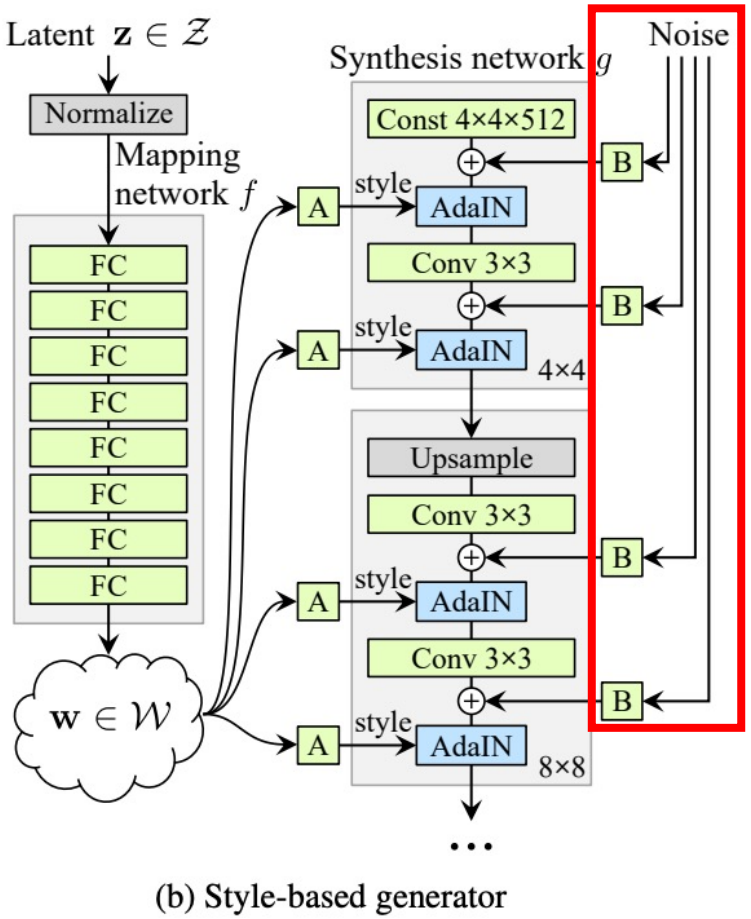

(b) Style-based generator

Stochastic variation

StyleGAN

: variation